

Universal Stochastic Predictor

Phase 1: API Foundations

Implementation Team

February 19, 2026

Índice

1 Phase 1: API Foundations Overview	2
1.1 Scope	2
1.2 Design Principles	2
2 Type System (types.py)	3
2.1 Overview	3
2.2 PredictorConfig Class	3
2.2.1 Purpose	3
2.2.2 Core Configuration Fields	3
2.3 Data Structures for Prediction API	5
2.3.1 ProcessState	5
2.3.2 PredictionResult	5
2.4 Immutability Guarantees	5
3 Configuration Management (config.py)	6
3.1 Architecture	6
3.2 ConfigManager Class	6
3.3 FIELD_TO_SECTION_MAP (Single Source of Truth)	6
3.4 config.toml Structure	7
4 Validation Framework (validation.py)	9
4.1 Purpose	9
4.2 Key Validators	9
4.2.1 validate_finite()	9
4.2.2 validate_simplex()	9
4.2.3 validate_holder_exponent()	9
4.2.4 validate_alpha_stable()	9
4.2.5 sanitize_array()	10
4.3 Zero-Heuristics Policy	10
5 Random Number Generation (random.py)	11
5.1 JAX PRNG Infrastructure	11
5.2 Reproducibility Verification	11
6 Schema Definitions (schemas.py)	12
6.1 Pydantic v2 Models	12
6.1.1 ProcessStateSchema	12
6.1.2 PredictionResultSchema	12
6.1.3 TelemetryDataSchema	12
6.2 Validation Features	12

7	Code Quality Metrics	14
7.1	Lines of Code	14
7.2	Compliance Checklist	14
8	Phase 1 Summary	15

Capítulo 1

Phase 1: API Foundations Overview

Phase 1 implements the foundational API layer for the Universal Stochastic Predictor (USP). This phase establishes core data structures, configuration management, validation framework, random number generation, and schema definitions required for kernel implementations (Phase 2).

1.1 Scope

Phase 1 covers:

- **Type System** (`types.py`): Core immutable dataclasses for configuration and predictions
- **Configuration Management** (`config.py`): Singleton ConfigManager with TOML-based parameter injection
- **Validation Framework** (`validation.py`): Domain-specific validation and sanitization logic
- **Random Number Generation** (`random.py`): JAX-based PRNG utilities
- **Schema Definitions** (`schemas.py`): Pydantic models for API contracts

1.2 Design Principles

- **Zero-Heuristics Policy:** All hyperparameters must reside in configuration, never hardcoded in code
- **100% English:** All code, comments, docstrings, and identifiers in English only
- **Immutability:** Data structures use frozen dataclasses for thread-safety and JAX compatibility
- **Type Safety:** Dimension checking via jaxtyping; strict validation boundaries

Capítulo 2

Type System (types.py)

2.1 Overview

The `types.py` module defines all immutable data structures using frozen dataclasses. This ensures thread-safe configuration sharing, JAX JIT compilation cache compatibility, and proper type checking.

2.2 PredictorConfig Class

2.2.1 Purpose

`PredictorConfig` is the system hyperparameter vector (denoted Λ in the specification). It contains all configurable parameters for orchestration, kernels, validation, and I/O. Total: **47 fields**.

2.2.2 Core Configuration Fields

Schema Versioning

```
1 schema_version: str = "1.0"
```

JKO Orchestrator (Optimal Transport)

```
1 epsilon: float = 1e-3           # Entropic regularization (Sinkhorn)
2 learning_rate: float = 0.01     # Learning rate tau
3 sinkhorn_epsilon_min: float = 0.01 # Min epsilon for coupling
4 sinkhorn_epsilon_0: float = 0.1   # Base epsilon
5 sinkhorn_alpha: float = 0.5      # Volatility coupling coefficient
```

Entropy Monitoring

```
1 entropy_window: int = 100        # Sliding window size
2 entropy_threshold: float = 0.8    # Mode collapse detection
```

Kernel Parameters

Kernel D (Log-Signatures)

```
1 kernel_d_depth: int = 3          # Truncation level
2 kernel_d_alpha: float = 0.1       # Extrapolation scaling
```

Kernel A (WTMM + Fokker-Planck)

```
1 wtmm_buffer_size: int = 128          # Memory buffer
2 besov_cone_c: float = 1.5            # Cone of influence
3 kernel_ridge_lambda: float = 1e-6    # RKHS regularization
4 kernel_a_bandwidth: float = 0.1      # Gaussian kernel smoothness
5 kernel_a_embedding_dim: int = 5       # Takens embedding
```

Kernel B (PDE/DGM)

```
1 dgm_width_size: int = 64             # Network width
2 dgm_depth: int = 4                  # Network depth
3 dgm_entropy_num_bins: int = 50      # Histogram bins
4 kernel_b_r: float = 0.05            # HJB interest rate
5 kernel_b_sigma: float = 0.2          # HJB volatility
6 kernel_b_horizon: float = 1.0        # Prediction horizon
```

Kernel C (SDE Integration)

```
1 stiffness_low: int = 100            # Explicit integrator threshold
2 stiffness_high: int = 1000           # Implicit integrator threshold
3 sde_dt: float = 0.01                # Time step
4 sde_numel_integrations: int = 100   # Number of steps
5 sde_diffusion_sigma: float = 0.2     # Diffusion coefficient
6 kernel_c_mu: float = 0.0             # Drift (mean reversion)
7 kernel_c_alpha: float = 1.8           # Stability (1 < alpha <= 2)
8 kernel_c_beta: float = 0.0             # Skewness (-1 <= beta <= 1)
9 kernel_c_horizon: float = 1.0         # Integration horizon
10 kernel_c_dt0: float = 0.01            # Initial time step (adaptive)
```

Risk Detection

```
1 holder_threshold: float = 0.4        # Holder singularity threshold
2 cusum_h: float = 5.0                 # CUSUM drift
3 cusum_k: float = 0.5                 # CUSUM slack
4 grace_period_steps: int = 20          # Refractory period
5 volatility_alpha: float = 0.1          # EWMA decay
```

Validation Constraints

```
1 sigma_bound: float = 20.0            # Outlier threshold (N sigma)
2 sigma_val: float = 1.0               # Reference std dev
3 max_future_drift_ns: int = 1_000_000_000 # Clock skew tolerance
4 max_past_drift_ns: int = 86_400_000_000_000 # Stale data threshold
```

I/O Policies

```
1 data_feed_timeout: int = 30           # Timeout seconds
2 data_feed_max_retries: int = 3         # Retry attempts
3 snapshot_atomic_fsync: bool = True     # Force fsync
4 snapshot_compression: str = "none"     # Compression method
5 staleness_ttl_ns: int = 500_000_000    # TTL degraded mode
6 besov_nyquist_interval_ns: int = 100_000_000 # Nyquist sample rate
7 inference_recovery_hysteresis: float = 0.8 # Recovery factor
```

Base Parameters

```
1 base_min_signal_length: int = 32           # Minimum length
2 signal_normalization_method: str = "zscore" # Normalization
3 log_sig_depth: int = 3                     # Log-signature truncation
```

2.3 Data Structures for Prediction API

2.3.1 ProcessState

```
1 @dataclass(frozen=True)
2 class ProcessState:
3     """Single observation from stochastic process."""
4     timestamp: float      # Observation time (ns)
5     price: float          # Observation value
6     volume: float         # Associated volume/energy
7     volatility_estimate: float # Auxiliary measure
```

2.3.2 PredictionResult

```
1 @dataclass(frozen=True)
2 class PredictionResult:
3     """Output prediction with uncertainty quantification."""
4     predicted_price: float      # Point estimate
5     confidence_interval_lower: float # Lower CI bound
6     confidence_interval_upper: float # Upper CI bound
7     predicted_volatility: float    # Volatility forecast
8     kernel_consensus: float       # Ensemble weight
9     entropy_diagnostic: float    # Mode collapse indicator
10    cusum_alert: bool            # Structural break
```

2.4 Immutability Guarantees

All dataclasses use `frozen=True` to enable:

- JAX JIT cache key hashing
- Thread-safe configuration sharing
- Enforcement of zero-heuristics policy

Capítulo 3

Configuration Management (config.py)

3.1 Architecture

config.py implements:

- Lazy singleton with `ConfigManager.get_instance()`
- TOML parsing with automatic field mapping
- Environment variable override support
- Runtime validation of completeness

3.2 ConfigManager Class

```
1 class ConfigManager:
2     """Singleton configuration loader."""
3
4     def load_config(self, path: str = "config.toml") -> PredictorConfig:
5         """Parse TOML and inject into PredictorConfig."""
6
7     def get_config(self) -> PredictorConfig:
8         """Retrieve cached configuration."""
9
10    @staticmethod
11    def _apply_env_overrides() -> None:
12        """Apply USP_SECTION_KEY environment variables."""
```

3.3 FIELD_TO_SECTION_MAP (Single Source of Truth)

Automated field-to-section mapping ensures all 47 config parameters have defined placement:

```
1 FIELD_TO_SECTION_MAP = {
2     # Metadata
3     "schema_version": "meta",
4
5     # Orchestration
6     "epsilon": "orchestration",
7     "learning_rate": "orchestration",
8     "sinkhorn_epsilon_min": "orchestration",
9     "sinkhorn_epsilon_0": "orchestration",
10    "sinkhorn_alpha": "orchestration",
11    "entropy_window": "orchestration",
12    "entropy_threshold": "orchestration",
```

```

13 "sigma_bound": "orchestration",
14 "sigma_val": "orchestration",
15 "max_future_drift_ns": "orchestration",
16 "max_past_drift_ns": "orchestration",
17 "holder_threshold": "orchestration",
18 "cusum_h": "orchestration",
19 "cusum_k": "orchestration",
20 "grace_period_steps": "orchestration",
21 "volatility_alpha": "orchestration",
22 "inference_recovery_hysteresis": "orchestration",
23
24 # Kernels
25 "log_sig_depth": "kernels",
26 "wtmm_buffer_size": "kernels",
27 "besov_cone_c": "kernels",
28 "besov_nyquist_interval_ns": "kernels",
29 "stiffness_low": "kernels",
30 "stiffness_high": "kernels",
31 "sde_dt": "kernels",
32 "sde_numel_integrations": "kernels",
33 "sde_diffusion_sigma": "kernels",
34 "kernel_ridge_lambda": "kernels",
35 "kernel_a_bandwidth": "kernels",
36 "kernel_a_embedding_dim": "kernels",
37 "dgm_width_size": "kernels",
38 "dgm_depth": "kernels",
39 "dgm_entropy_num_bins": "kernels",
40 "kernel_b_r": "kernels",
41 "kernel_b_sigma": "kernels",
42 "kernel_b_horizon": "kernels",
43 "kernel_c_mu": "kernels",
44 "kernel_c_alpha": "kernels",
45 "kernel_c_beta": "kernels",
46 "kernel_c_horizon": "kernels",
47 "kernel_c_dt0": "kernels",
48 "kernel_d_depth": "kernels",
49 "kernel_d_alpha": "kernels",
50
51 # I/O
52 "data_feed_timeout": "io",
53 "data_feed_max_retries": "io",
54 "snapshot_atomic_fsync": "io",
55 "snapshot_compression": "io",
56
57 # Core
58 "staleness_ttl_ns": "core",
59
60 # Base
61 "base_min_signal_length": "base",
62 "signal_normalization_method": "base",
63 }

```

3.4 config.toml Structure

```

1 [meta]
2 schema_version = "1.0"
3
4 [orchestration]
5 epsilon = 1e-3
6 learning_rate = 0.01
7 sinkhorn_epsilon_min = 0.01

```

```

8 sinkhorn_epsilon_0 = 0.1
9 sinkhorn_alpha = 0.5
10 entropy_window = 100
11 entropy_threshold = 0.8
12 sigma_bound = 20.0
13 sigma_val = 1.0
14 max_future_drift_ns = 1_000_000_000
15 max_past_drift_ns = 86_400_000_000_000
16 holder_threshold = 0.4
17 cusum_h = 5.0
18 cusum_k = 0.5
19 grace_period_steps = 20
20 volatility_alpha = 0.1
21 inference_recovery_hysteresis = 0.8
22
23 [kernels]
24 log_sig_depth = 3
25 wtmm_buffer_size = 128
26 besov_cone_c = 1.5
27 besov_nyquist_interval_ns = 100_000_000
28 stiffness_low = 100
29 stiffness_high = 1000
30 sde_dt = 0.01
31 sde_numel_integrations = 100
32 sde_diffusion_sigma = 0.2
33 kernel_ridge_lambda = 1e-6
34 kernel_a_bandwidth = 0.1
35 kernel_a_embedding_dim = 5
36 dgm_width_size = 64
37 dgm_depth = 4
38 dgm_entropy_num_bins = 50
39 kernel_b_r = 0.05
40 kernel_b_sigma = 0.2
41 kernel_b_horizon = 1.0
42 kernel_c_mu = 0.0
43 kernel_c_alpha = 1.8
44 kernel_c_beta = 0.0
45 kernel_c_horizon = 1.0
46 kernel_c_dt0 = 0.01
47 kernel_d_depth = 3
48 kernel_d_alpha = 0.1
49
50 [io]
51 data_feed_timeout = 30
52 data_feed_max_retries = 3
53 snapshot_atomic_fsync = true
54 snapshot_compression = "none"
55
56 [core]
57 staleness_ttl_ns = 500_000_000
58
59 [base]
60 base_min_signal_length = 32
61 signal_normalization_method = "zscore"

```

Capítulo 4

Validation Framework (validation.py)

4.1 Purpose

Validation functions enforce domain-agnostic constraints on all inputs. Each validator is **configuration-driven** with zero hardcoded parameters.

4.2 Key Validators

4.2.1 validate_finite()

```
1 def validate_finite(arr: Array, *,
2                     allow_nan: bool,
3                     allow_inf: bool) -> Array:
4     """Check array for NaN/Inf violations."""
```

Parameters required from config:

- `allow_nan: config.validation_finite_allow_nan`
- `allow_inf: config.validation_finite_allow_inf`

4.2.2 validate_simplex()

```
1 def validate_simplex(weights: Array, *, atol: float) -> Array:
2     """Probability simplex: all >= 0, sum = 1.0"""
```

Parameter from config: `atol ← config.validation_simplex_atol`

4.2.3 validate_holder_exponent()

```
1 def validate_holder_exponent(val: float, *,
2                               min_val: float,
3                               max_val: float) -> float:
4     """Holder continuity: bounds enforcement."""
```

Parameters from config: `min_val, max_val`

4.2.4 validate_alpha_stable()

```
1 def validate_alpha_stable(alpha: float, beta: float, *,
2                           alpha_min: float, alpha_max: float,
3                           beta_min: float, beta_max: float,
4                           exclusive_bounds: bool = True) -> tuple:
5     """Levy alpha-stable parameter space validation."""
```

4.2.5 sanitize_array()

```
1 def sanitize_array(arr: Array, *,
2                     replace_nan: float,
3                     replace_inf: Optional[float],
4                     clip_range: Optional[tuple]) -> Array:
5     """Replace NaN/Inf; optionally clip to range."""

```

4.3 Zero-Heuristics Policy

All validation parameters must come from config. No function contains hardcoded defaults:

```
1 # CORRECT (config-driven):
2 result = validate_finite(array,
3                           allow_nan=config.validation_finite_allow_nan,
4                           allow_inf=config.validation_finite_allow_inf)
5
6 # WRONG (hardcoded):
7 result = validate_finite(array, allow_nan=False, allow_inf=False)
```

Capítulo 5

Random Number Generation (random.py)

5.1 JAX PRNG Infrastructure

The `random.py` module provides deterministic sampling via JAX's threefry2x32 PRNG:

```
1 def initialize_jax_prng(seed: int) -> PRNGKeyArray:
2     """Create root PRNGKey from seed."""
3
4 def split_key(key: PRNGKeyArray) -> tuple[PRNGKeyArray, PRNGKeyArray]:
5     """Cryptographic key splitting for parallel RNG streams."""
6
7 def uniform_samples(key: PRNGKeyArray, n: int) -> Array:
8     """Generate n uniform [0, 1) samples."""
9
10 def normal_samples(key: PRNGKeyArray, n: int,
11                     loc: float = 0.0, scale: float = 1.0) -> Array:
12     """Generate n Gaussian samples."""
13
14 def exponential_samples(key: PRNGKeyArray, n: int,
15                         rate: float = 1.0) -> Array:
16     """Generate n exponential samples."""
```

5.2 Reproducibility Verification

```
1 def verify_determinism(seed: int, n_trials: int = 10) -> bool:
2     """Verify identical output across multiple runs."""
```

Capítulo 6

Schema Definitions (schemas.py)

6.1 Pydantic v2 Models

schemas.py defines API contracts with strict type enforcement:

6.1.1 ProcessStateSchema

```
1 class ProcessStateSchema(BaseModel):
2     """API contract for process observations."""
3     timestamp_utc: datetime
4     price: float = Field(..., gt=0)
5     volume: float = Field(..., ge=0)
6     volatility_proxy: Optional[float] = None
```

6.1.2 PredictionResultSchema

```
1 class PredictionResultSchema(BaseModel):
2     """API contract for predictions."""
3     predicted_price: float = Field(..., gt=0)
4     confidence_interval_lower: float
5     confidence_interval_upper: float
6     predicted_volatility: float = Field(..., ge=0)
7     kernel_consensus: float = Field(..., ge=0, le=1)
8     entropy_diagnostic: float = Field(..., ge=0)
9     cusum_alert: bool
```

6.1.3 TelemetryDataSchema

```
1 class TelemetryDataSchema(BaseModel):
2     """Diagnostic telemetry."""
3     prediction_latency_ms: float
4     kernel_latency_ms: Dict[str, float]
5     memory_usage_mb: float
6     entropy_value: float
```

6.2 Validation Features

All schemas enforce:

- Field constraints: gt, ge, le, lt

- Type strictness: No implicit coercion
- Custom validators: `@field_validator` for domain logic

Capítulo 7

Code Quality Metrics

7.1 Lines of Code

Module	LOC
types.py	347
config.py	220
validation.py	467
random.py	301
schemas.py	330
Total API Layer	1,665

7.2 Compliance Checklist

- 100% English code (no Spanish identifiers)
- Full type hints with dimensional consistency
- No hardcoded hyperparameters (zero-heuristics policy)
- All 47 config fields mapped via FIELD_TO_SECTION_MAP
- Immutable frozen dataclasses for thread-safety
- Environment variable overrides (USP_SECTION__KEY)
- Pydantic v2 strict validation

Capítulo 8

Phase 1 Summary

Phase 1 establishes production-ready API foundations:

- **Type System:** 47-field `PredictorConfig` with frozen immutability
- **Configuration:** TOML-driven, environment-overridable, automated field mapping
- **Validation:** Domain-agnostic, config-driven, zero hardcoded defaults
- **PRNG:** JAX-native threefry2x32 with reproducibility guarantees
- **Schemas:** Pydantic v2 with strict type enforcement

Ready for Phase 2 kernel implementations.