

# **Universal Stochastic Predictor**

## **Phase 4: IO Layer Initiation**

Implementation Team

February 19, 2026

# Contents

<b>1 Phase 4: IO Layer Initiation Overview</b>	<b>2</b>
1.1 Scope . . . . .	2
1.2 Design Principles . . . . .	2
<b>2 Ingestion and Validation</b>	<b>3</b>
2.1 Catastrophic Outlier Filter . . . . .	3
2.2 Frozen Signal Alarm . . . . .	3
2.3 Staleness Policy (TTL) . . . . .	3
<b>3 Telemetry Abstraction</b>	<b>4</b>
3.1 TelemetryBuffer Emission . . . . .	4
3.2 No Compute Stalls . . . . .	4
<b>4 Deterministic Logging</b>	<b>5</b>
4.1 Hash-Based Parity Checks . . . . .	5
4.2 Audit Hashes . . . . .	5
<b>5 Snapshot Strategy</b>	<b>6</b>
5.1 Atomic Persistence . . . . .	6
5.2 Binary Serialization . . . . .	6
5.3 Integrity Verification . . . . .	6
5.4 Atomic Write Protocol . . . . .	6
<b>6 Security Policies</b>	<b>7</b>
6.1 Credential Injection . . . . .	7
6.2 Version Control Exclusion . . . . .	7
<b>7 Compliance Checklist</b>	<b>8</b>
<b>8 Phase 4 Summary</b>	<b>9</b>

# Chapter 1

## Phase 4: IO Layer Initiation Overview

Phase 4 introduces the asynchronous I/O layer for snapshots, streaming, and telemetry export. The primary design goal is to preserve JAX/XLA throughput by decoupling compute from disk or network latency.

### 1.1 Scope

Phase 4 covers:

- **Telemetry Buffering:** Non-blocking emission of telemetry snapshots
- **Deterministic Logging:** Hash-based parity checks for CPU/GPU validation
- **Snapshot Strategy:** Atomic persistence of predictor state
- **Ingestion and Validation:** Input filtering, staleness policy, frozen signal detection
- **Security Enforcement:** Credential injection and secret exclusion

### 1.2 Design Principles

- **No Compute Stalls:** JAX compute threads never block on I/O
- **Determinism:** Logs capture reproducible hashes instead of raw state dumps
- **Security:** No raw signals or secrets in logs
- **Configurability:** Logging intervals and destinations injected via config
- **Integrity:** Snapshots and parity logs are hash-verified

# Chapter 2

## Ingestion and Validation

### 2.1 Catastrophic Outlier Filter

Input validation must reject catastrophic outliers when  $|y_t| > 20\sigma$  relative to historical normalization. In this case, the system must preserve inertial state and emit a critical alert without advancing the transport update.

- Reject observation and keep current state unchanged.
- Emit a critical alert for audit visibility.
- Do not update JKO/Sinkhorn weights for the rejected step.

### 2.2 Frozen Signal Alarm

If the exact same value is observed for  $N_{freeze} \geq 5$  consecutive steps, emit a `FrozenSignalAlarmEvent`. This invalidates the multifractal spectrum and requires:

- Freeze the topological branch (Kernel D).
- Switch to degraded inference mode.
- Continue monitoring until signal variation resumes.

### 2.3 Staleness Policy (TTL)

Every observation must carry a timestamp for TTL evaluation. If the target delay exceeds  $\Delta_{max}$ , the JKO update must be suspended immediately.

- Compute staleness as  $\Delta_t = t_{now} - t_{obs}$ .
- If  $\Delta_t > \Delta_{max}$ , skip the transport update.
- Preserve state and record a staleness warning event.

# Chapter 3

# Telemetry Abstraction

## 3.1 TelemetryBuffer Emission

The JKO orchestrator should emit a `TelemetryBuffer` at the end of each step. This buffer is consumed by a dedicated process outside the JAX execution thread.

- The buffer contains summary metrics (CUSUM, entropy, regime flags, OT cost).
- The compute path only enqueues the buffer and continues.
- The consumer is responsible for serialization and persistence.

## 3.2 No Compute Stalls

JAX compute threads must never block on I/O. Telemetry buffers must be non-blocking and consumed by a separate process or thread outside the JAX execution path.

# Chapter 4

## Deterministic Logging

### 4.1 Hash-Based Parity Checks

For hardware parity audits, the logger records SHA-256 hashes of the weight vector  $\rho$  and the OT cost at configurable intervals. This permits CPU/GPU parity validation without dumping VRAM data.

- Hash interval configured per deployment.
- Hashes derived from canonical float64 serialization.
- Logs are append-only and immutable.

### 4.2 Audit Hashes

Parity audits must log SHA-256 hashes of  $\rho$  and OT cost at configured intervals. Hash input must be derived from canonical float64 serialization to ensure reproducibility across CPU and GPU.

# Chapter 5

## Snapshot Strategy

### 5.1 Atomic Persistence

Snapshots must be persisted atomically to prevent partial writes. The IO layer is responsible for:

- Writing to temporary files and renaming atomically.
- Optional compression configured by policy.
- Coordinating snapshot cadence with telemetry output.

### 5.2 Binary Serialization

Text formats (JSON, XML) are prohibited for critical snapshots due to latency and ambiguity. Use dense binary formats such as Protocol Buffers or MessagePack.

- Encode all fields deterministically.
- Preserve float64 for numerical fidelity.

### 5.3 Integrity Verification

Each snapshot  $\Sigma_t$  must include a hash footer (SHA-256 or CRC32c). The load routine must verify the hash before injecting state into memory.

- Fail closed if hash verification fails.
- Log integrity failures at critical severity.

### 5.4 Atomic Write Protocol

To avoid partial writes, persist snapshots to a temporary file and then atomically rename to the target path. The rename step must be the only visible operation to consumers.

- Use a unique temporary filename per snapshot.
- Ensure the target file is replaced atomically.

# Chapter 6

# Security Policies

## 6.1 Credential Injection

Tokens and API keys must not appear in source code. Credentials must be injected at runtime via environment variables or .env files.

## 6.2 Version Control Exclusion

The repository must exclude .env files and credential directories via .gitignore. Secrets must never be committed.

## Chapter 7

# Compliance Checklist

- **No Compute Stalls:** All logging is asynchronous
- **Binary Format:** Protocol Buffers or MessagePack for snapshots
- **Atomic Snapshots:** Write-then-rename protocol
- **Deterministic Hashing:** SHA-256 on  $\rho$  and OT cost
- **Security:** No raw signals, VRAM dumps, or secrets
- **Integrity:** Snapshot hashes verified before load
- **Config-Driven:** Intervals and destinations are injected

## **Chapter 8**

# **Phase 4 Summary**

Phase 4 introduces a non-blocking I/O architecture that preserves deterministic compute while enabling telemetry, logging, and atomic snapshot persistence.