

Python API Specification - Universal Predictor

Software Engineering

February 19, 2026

Contents

1	Introduction	2
2	Data Structures (Typing)	2
2.1	Configuration (Λ)	2
2.2	Operational Input ($y_t, y_{reference}, \tau$)	2
2.3	System Output	2
3	Multi-Tenant Architecture (Stateless Functional Pattern)	3
3.1	Throughput Maximization (Vectorized Batching)	3
4	Main Class: <code>UniversalPredictor</code> (Stateful Wrapper)	3
4.1	Initialization	4
4.2	Execution Method ($t \rightarrow t + 1$)	4
5	Preventing VRAM Fragmentation (JAX Memory Management)	5
6	VRAM Monitoring	5
7	Recommended Deployment Configuration	5
8	Persistence (Atomic Snapshotting)	6
9	Asynchronous I/O for Snapshots (Non-Blocking)	6
10	Graceful Shutdown for Containers	7
11	Prometheus Integration	9
12	Adaptive CUSUM Threshold	9
13	Grace Period (Post-Regime Refractory Window)	9
14	Operational Flags and Recovery	9
15	Error Handling and Exceptions	10
16	Production Logging Example	10
17	Deterministic Floating-Point Reproducibility	10
18	Load Shedding (Adaptive Topological Pruning)	11
19	Jitter Telemetry	11
20	Dependency Pinning	11

1 Introduction

This document specifies the Python implementation of the abstract I/O interface defined in *Stochastic_Predictor_IO*. The API exposes the `UniversalPredictor` class for high-performance environments using JAX for numerical acceleration.

2 Data Structures (Typing)

We use `dataclasses` and `jaxtyping` to enforce immutability and strict dimensional typing for tensors.

2.1 Configuration (Λ)

```
1 from dataclasses import dataclass
2 from typing import Optional
3 from jaxtyping import Float, Array, Bool
4
5 @dataclass(frozen=True)
6 class PredictorConfig:
7     """Hyperparameter vector Lambda."""
8     schema_version: str = "1.0"      # Snapshot versioning
9     epsilon: float = 1e-3            # Entropic regularization (Sinkhorn)
10    learning_rate: float = 0.01     # JKO learning rate
11    log_sig_depth: int = 3          # Signature depth (Kernel D)
12    wtmm_buffer_size: int = 128    # WTMM buffer size
13    besov_cone_c: float = 1.5       # Besov cone influence
14    holder_threshold: float = 0.4   # Circuit breaker threshold
15    cusum_h: float = 5.0           # CUSUM threshold
16    cusum_k: float = 0.5           # CUSUM slack
17    grace_period_steps: int = 20    # Post-regime refractory period
18    volatility_alpha: float = 0.1   # EMA decay for variance
19
20    # Load shedding and anti-aliasing
21    staleness_ttl_ns: int = 500_000_000      # TTL (500ms)
22    besov_nyquist_interval_ns: int = 100_000_000 # Nyquist soft limit (100ms)
23    inference_recovery_hysteresis: float = 0.8  # Degraded mode recovery factor
```

2.2 Operational Input ($y_t, y_{reference}, \tau$)

```
1 @dataclass(frozen=True)
2 class ProcessState:
3     magnitude: Float[Array, "1"]      # y_t (normalized or absolute)
4     reference: Float[Array, "1"]       # y_reference
5     timestamp_ns: int                # Unix epoch (nanoseconds)
6
7     def validate_domain(self, sigma_bound: float = 20.0, sigma_val: float = 1.0) -> bool:
8         """Catastrophic outlier detection (> 20 sigma)."""
9         return abs(self.magnitude) <= (sigma_bound * sigma_val)
```

2.3 System Output

```
1 @dataclass(frozen=True)
2 class PredictionResult:
3     predicted_next: Float[Array, "1"]    # y_{t+1} (Z-score)
4
5     # Telemetry
6     holder_exponent: Float[Array, "1"]
7     cusum_drift: Float[Array, "1"]
8     distance_to_collapse: Float[Array, "1"]
9     free_energy: Float[Array, "1"]
10
11    # Advanced telemetry
12    kurtosis: Float[Array, "1"]
13    dgm_entropy: Float[Array, "1"]
14    adaptive_threshold: Float[Array, "1"]
15
16    # Orchestrator state
```

```

17     weights: Float[Array, "4"]
18
19     # Health flags
20     sinkhorn_converged: Bool[Array, "1"]
21     degraded_inference_mode: bool
22     emergency_mode: bool
23     regime_change_detected: bool
24     modeCollapse_warning: bool
25
26     mode: str # "Standard" | "Robust" | "Emergency"

```

3 Multi-Tenant Architecture (Stateless Functional Pattern)

To support hundreds of assets on a single server, the API exposes a purely functional mode. This allows state management in low-latency external storage (Redis) while sharing the compiled JAX graph across assets.

3.1 Throughput Maximization (Vectorized Batching)

This architecture enables `jax.vmap` to batch multiple asset states in a single hardware call, minimizing the Python GIL impact and maximizing GPU occupancy.

```

1 class FunctionalPredictor:
2     """
3         Stateless implementation for JAX core.
4         Scales to thousands of predictors sharing the same graph.
5     """
6
7     def __init__(self, config: PredictorConfig):
8         self.config = config
9         self._core_step = self._core_update_step
10        self._jit_update = jax.jit(self._core_step)
11        self._vmap_update = jax.jit(jax.vmap(self._core_step, in_axes=(0, 0, 0, 0)))
12
13    def init_state(self):
14        """Create a zeroed cold-state structure."""
15        return self._initialize_state_structure()
16
17    def step(self, state, obs: ProcessState) -> tuple[object, PredictionResult]:
18        """
19            Pure state transition: (S_t, Obs_t) -> (S_{t+1}, Pred_{t+1})
20        """
21        should_freeze = self._should_freeze(obs)
22        new_state, raw_result = self._jit_update(
23            state,
24            obs.magnitude,
25            obs.reference,
26            freeze_weights=should_freeze
27        )
28        result = PredictionResult(
29            predicted_next=raw_result.y_next,
30            # map the remaining fields
31        )
32        return new_state, result
33
34    def step_batch(self, states, obs_batch: ProcessState):
35        """Vectorized batch processing for N assets."""
36        freeze_flags = self._should_freeze_batch(obs_batch)
37        new_states, results = self._vmap_update(
38            states, obs_batch.magnitude, obs_batch.reference, freeze_flags
39        )
40        return new_states, results

```

4 Main Class: UniversalPredictor (Stateful Wrapper)

This class wraps the functional pattern for single-tenant usage with state held in local memory.

4.1 Initialization

```

1 class UniversalPredictor:
2     def __init__(self, config: PredictorConfig):
3         """
4             Initialize the JAX compute graph (XLA JIT compilation).
5             Allocate static device buffers (VRAM).
6             Internal state stores persistent rolling buffers updated with
7             functional ops to avoid CPU<->VRAM transfers.
8         """
9         self.config = config
10        self._state = self._initialize_state()
11        self._jit_update = jax.jit(self._core_update_step)
12        self._last_timestamp_ns = 0
13
14    def fit_history(self, history: list[float]) -> bool:
15        """
16            Cold-start bootstrapping. Requires at least N_buf samples.
17            Returns True if Sinkhorn and CUSUM converge.
18        """
19        if len(history) < self.config.wtmm_buffer_size:
20            raise ValueError(f"Insufficient history. Required: {self.config.wtmm_buffer_size}")
21
22        self._state, final_metrics = self._jit_scan_history(self._state, jnp.array(history))
23
24        is_converged = final_metrics.sinkhorn_converged
25        is_stable = final_metrics.cusum_drift < self.config.cusum_h
26        if not (is_converged and is_stable):
27            logger.warning("Cold start finished without stable convergence.")
28            return False
29        return True

```

4.2 Execution Method ($t \rightarrow t + 1$)

```

1     def step(self, obs: ProcessState) -> PredictionResult:
2         """
3             Execute one prediction cycle with domain and TTL validation.
4         """
5         if not obs.validate_domain():
6             logger.error("Catastrophic outlier detected. Ignoring tick.")
7             return self._last_valid_result
8
9         current_time = time.time_ns()
10        latency = current_time - obs.timestamp_ns
11        is_stale = latency > self.config.staleness_ttl_ns
12
13        dt_arrival = obs.timestamp_ns - self._last_timestamp_ns
14        is_sparse = (self._last_timestamp_ns > 0) and (
15            dt_arrival > self.config.besov_nyquist_interval_ns
16        )
17        if is_sparse:
18            logger.warning(
19                f"FrequencyWarning: interval {dt_arrival}ns > Nyquist limit. WTMM may alias."
20            )
21
22        self._last_timestamp_ns = obs.timestamp_ns
23        should_freeze = is_stale or is_sparse
24
25        new_state, result_data = self._jit_update(
26            self._state,
27            obs.magnitude,
28            obs.reference,
29            freeze_weights=should_freeze,
30        )
31        self._state = new_state
32
33        return PredictionResult(
34            predicted_next=result_data.y_next,
35            holder_exponent=result_data.H_t,
36            sinkhorn_converged=result_data.converged,
37            # map remaining fields
38        )

```

5 Preventing VRAM Fragmentation (JAX Memory Management)

Production problem: JAX preallocates 90% of GPU memory on first access. Long-running systems may fragment VRAM and hit silent OOM after weeks.

Solution: Configure environment variables **before** importing JAX:

```
1 import os
2
3 os.environ['XLA_PYTHON_CLIENT_MEM_FRACTION'] = '0.7'
4 os.environ['XLA_PYTHON_CLIENT_ALLOCATOR'] = 'platform'
5 os.environ['TF_FORCE_GPU_ALLOW_GROWTH'] = 'true'
6
7 import jax
8 import jax.numpy as jnp
```

6 VRAM Monitoring

```
1 import psutil
2 import subprocess
3
4 def monitor_vram_fragmentation(interval_seconds=60):
5     """Background thread for VRAM monitoring."""
6     import time
7     import threading
8
9     def _monitor():
10         while True:
11             try:
12                 result = subprocess.run(
13                     ['nvidia-smi', '--query-gpu=memory.used,memory.total',
14                      '--format=csv,nounits,noheader'],
15                     capture_output=True, text=True, timeout=5
16                 )
17                 if result.returncode == 0:
18                     used, total = map(float, result.stdout.strip().split(','))
19                     utilization = 100.0 * used / total
20                     if utilization > 0.95:
21                         print(f"[WARNING] VRAM near saturation: {utilization:.1f}%")
22                     elif utilization > 0.85:
23                         print(f"[INFO] VRAM utilization: {utilization:.1f}% (elevated)")
24                     time.sleep(interval_seconds)
25                 except Exception as e:
26                     print(f"[ERROR] VRAM monitoring failed: {e}")
27                     break
28
29         thread = threading.Thread(target=_monitor, daemon=True)
30         thread.start()
```

7 Recommended Deployment Configuration

```
1 #!/bin/bash
2 # deployment/run_predictor.sh
3
4 export XLA_PYTHON_CLIENT_MEM_FRACTION=0.7
5 export XLA_PYTHON_CLIENT_ALLOCATOR=platform
6 export TF_FORCE_GPU_ALLOW_GROWTH=true
7
8 echo "[INFO] XLA VRAM Fraction: 0.7 (28/40 GB on A100)"
9 echo "[INFO] Allocator: platform (dynamic)"
10 echo "[INFO] GPU growth: enabled"
11
12 python3 -u predictor_service.py \
13     --config config.yaml \
14     --device gpu \
15     --pool-size 100 \
16     --monitor-interval 300
```

8 Persistence (Atomic Snapshotting)

```
1 import hashlib
2 import msgpack
3
4 def save_snapshot(self, filepath: str):
5     """
6         Export internal state Sigma_t as MessagePack.
7         Append SHA-256 checksum at the end of the file.
8     """
9     state_dict = self._serialize_jax_state(self._state)
10    payload = {
11        "schema_version": self.config.schema_version,
12        "timestamp": time.time_ns(),
13        "config": asdict(self.config),
14        "global": state_dict["global"],
15        "telemetry": {
16            "kurtosis": float(self._state.kurtosis),
17            "dgm_entropy": float(self._state.dgm_entropy),
18            "adaptive_threshold": float(self._state.h_adaptive)
19        },
20        "flags": {
21            "degraded_inference": bool(self._state.degraded_mode),
22            "emergency": bool(self._state.emergency_mode),
23            "regime_change": bool(self._state.regime_changed),
24            "mode_collapse": bool(self._state.mode_collapse_warning)
25        },
26        "kernels": {
27            "A": state_dict["kernel_a"],
28            "B": state_dict["kernel_b"],
29            "C": state_dict["kernel_c"],
30            "D": state_dict["kernel_d"]
31        }
32    }
33    data_bytes = msgpack.packb(payload)
34    checksum = hashlib.sha256(data_bytes).hexdigest()
35
36    with open(filepath, "wb") as f:
37        f.write(data_bytes)
38        f.write(checksum.encode('utf-8'))
39
40
41 def load_snapshot(self, filepath: str):
42     """
43         Load state. Validate SHA-256 and schema_version.
44         Raise ValueError if validation fails.
45     """
46     with open(filepath, "rb") as f:
47         content = f.read()
48
49     data_bytes = content[:-64]
50     stored_checksum = content[-64:].decode('utf-8')
51
52     computed = hashlib.sha256(data_bytes).hexdigest()
53     if computed != stored_checksum:
54         raise ValueError("Snapshot corrupt: checksum mismatch.")
55
56     payload = msgpack.unpackb(data_bytes)
57     loaded_schema = payload.get('schema_version', 'unknown')
58     if loaded_schema != self.config.schema_version:
59         raise ValueError(
60             f"Schema version mismatch: snapshot={loaded_schema}, current={self.config.schema_version}."
61         )
62
63     self._state = self._deserialize_jax_state(payload)
```

9 Asynchronous I/O for Snapshots (Non-Blocking)

```
1 import concurrent.futures
```

```

2 import hashlib
3 import msgpack
4 import threading
5 import time
6
7 class UniversalPredictor_AsyncIO:
8     def __init__(self, n_worker_threads=2):
9         self.io_executor = concurrent.futures.ThreadPoolExecutor(
10             max_workers=n_worker_threads,
11             thread_name_prefix="snapshot_io_"
12         )
13         self.pending_snapshot_future = None
14         self.snapshot_lock = threading.Lock()
15
16     def _compute_and_save_async(self, filepath: str, data_bytes: bytes):
17         checksum = hashlib.sha256(data_bytes).hexdigest()
18         temp_filepath = filepath + ".tmp"
19         try:
20             with open(temp_filepath, "wb") as f:
21                 f.write(data_bytes)
22                 f.write(checksum.encode('utf-8'))
23             import os
24             os.replace(temp_filepath, filepath)
25         return {
26             'status': 'success',
27             'filepath': filepath,
28             'filesize_bytes': len(data_bytes),
29             'checksum': checksum,
30             'timestamp': time.time()
31         }
32     except Exception as e:
33         return {
34             'status': 'error',
35             'filepath': filepath,
36             'error': str(e),
37             'timestamp': time.time()
38         }
39
40     def save_snapshot_nonblocking(self, filepath: str) -> concurrent.futures.Future:
41         state_dict = self._serialize_jax_state(self._state)
42         payload = {
43             "schema_version": self.config.schema_version,
44             "timestamp": time.time_ns(),
45             "config": asdict(self.config),
46             "global": state_dict["global"],
47             "telemetry": {
48                 "kurtosis": float(self._state.kurtosis),
49                 "dgm_entropy": float(self._state.dgm_entropy),
50                 "adaptive_threshold": float(self._state.h_adaptive)
51             },
52             "flags": {
53                 "degraded_inference": bool(self._state.degraded_mode),
54                 "emergency": bool(self._state.emergency_mode),
55                 "regime_change": bool(self._state.regime_changed),
56                 "mode_collapse": bool(self._state.mode_collapse_warning)
57             },
58             "kernels": {
59                 "A": state_dict["kernel_a"],
60                 "B": state_dict["kernel_b"],
61                 "C": state_dict["kernel_c"],
62                 "D": state_dict["kernel_d"]
63             }
64         }
65         data_bytes = msgpack.packb(payload)
66         future = self.io_executor.submit(self._compute_and_save_async, filepath, data_bytes)
67         with self.snapshot_lock:
68             self.pending_snapshot_future = future
69         return future

```

10 Graceful Shutdown for Containers

```

1 import signal
2 import sys
3 import threading
4 import time
5 import logging
6 from typing import Optional
7
8 class UniversalPredictor_GracefulShutdown:
9     def __init__(self, config: PredictorConfig):
10         self.config = config
11         self.predictor = UniversalPredictor_AsyncIO(config)
12         self.shutdown_requested = threading.Event()
13         self.is_accepting_data = True
14         self.input_buffer_lock = threading.Lock()
15         self.residual_buffer = []
16
17         signal.signal(signal.SIGTERM, self._handle_sigterm)
18         signal.signal(signal.SIGINT, self._handle_sigint)
19
20         self.logger = logging.getLogger("predictor.shutdown")
21         self.logger.info("[INIT] Graceful shutdown handler registered")
22
23     def _handle_sigterm(self, signum, frame):
24         self.logger.warning(f"[SIGTERM] Received signal {signum}. Initiating graceful shutdown")
25         self.shutdown_requested.set()
26
27     def _handle_sigint(self, signum, frame):
28         self.logger.warning(f"[SIGINT] Received signal {signum}. Initiating graceful shutdown")
29         self.shutdown_requested.set()
30
31     def accept_observation(self, obs: ProcessState) -> Optional[PredictionResult]:
32         if self.shutdown_requested.is_set() or not self.is_accepting_data:
33             self.logger.warning(f"[REJECT] Observation rejected (shutdown in progress): {obs.timestamp_ns}")
34             return None
35         with self.input_buffer_lock:
36             self.residual_buffer.append(obs)
37         return self._process_observation(obs)
38
39     def _process_observation(self, obs: ProcessState) -> PredictionResult:
40         result = self.predictor.predict_next(obs)
41         with self.input_buffer_lock:
42             if obs in self.residual_buffer:
43                 self.residual_buffer.remove(obs)
44         return result
45
46     def graceful_shutdown(self, timeout_seconds: int = 25):
47         shutdown_start = time.time()
48         self.logger.info("GRACEFUL SHUTDOWN INITIATED")
49         self.is_accepting_data = False
50         time.sleep(0.1)
51
52         with self.input_buffer_lock:
53             for obs in list(self.residual_buffer):
54                 if time.time() - shutdown_start > timeout_seconds - 10:
55                     self.logger.warning("Timeout approaching, aborting residual processing")
56                     break
57             try:
58                 _ = self._process_observation(obs)
59             except Exception as e:
60                 self.logger.error(f"Error processing residual: {e}")
61
62             pending_snapshot = self.predictor.pending_snapshot_future
63             if pending_snapshot is not None and not pending_snapshot.done():
64                 try:
65                     remaining_time = max(1, timeout_seconds - (time.time() - shutdown_start))
66                     pending_snapshot.result(timeout=remaining_time)
67                 except Exception as e:
68                     self.logger.error(f"Async snapshot failed: {e}")
69
70             try:

```

```

71         final_snapshot_path = f"snapshots/shutdown_{int(time.time())}.pkl"
72         self.predictor.save_snapshot(final_snapshot_path)
73     except Exception as e:
74         self.logger.error(f"Final snapshot failed: {e}")
75
76     try:
77         if hasattr(self.predictor, 'io_executor'):
78             self.predictor.io_executor.shutdown(wait=True, cancel_futures=False)
79     except Exception as e:
80         self.logger.error(f"Error closing resources: {e}")
81
82     total_time = time.time() - shutdown_start
83     self.logger.info(f"SHUTDOWN COMPLETED ({total_time:.2f}s)")
84     sys.exit(0)

```

11 Prometheus Integration

```

1 from prometheus_client import Counter, Histogram
2
3 class UniversalPredictor_GracefulShutdown_Monitored:
4     def __init__(self, config: PredictorConfig):
5         self.shutdown_counter = Counter(
6             'predictor_graceful_shutdowns_total',
7             'Total number of graceful shutdowns executed'
8         )
9         self.shutdown_duration = Histogram(
10            'predictor_shutdown_duration_seconds',
11            'Time taken to complete graceful shutdown',
12            buckets=[0.5, 1.0, 2.0, 5.0, 10.0, 20.0, 30.0]
13        )
14         self.residual_observations = Histogram(
15            'predictor_shutdown_residual_observations',
16            'Number of observations in buffer during shutdown',
17            buckets=[0, 1, 5, 10, 50, 100, 500, 1000]
18        )

```

12 Adaptive CUSUM Threshold

The system implements the adaptive threshold lemma based on kurtosis:

$$h_t = k \cdot \sigma_t \cdot \left(1 + \ln\left(\frac{\kappa_t}{3}\right)\right)$$

13 Grace Period (Post-Regime Refractory Window)

After a regime change ($G^+ > h_t$), the system resets weights to uniform. A grace period prevents a cascade of false alarms while weights re-converge. The detector continues to compute G^+ but does not emit an alarm until the counter expires.

14 Operational Flags and Recovery

The system exposes explicit flags:

- **degraded_inference_mode**: TTL exceeded; weights frozen.
- **emergency_mode**: $H_t < H_{min}$; force Kernel D and Huber loss.
- **regime_change_detected**: CUSUM alarm; entropy reset.
- **mode_collapse_warning**: DGM entropy below threshold for > 10 steps.

15 Error Handling and Exceptions

Standard alerts:

- `DomainError`: catastrophic outlier $> 20\sigma$
- `StalenessWarning`: TTL exceeded
- `FrequencyWarning`: Nyquist limit violated
- `IntegrityError`: snapshot verification failed

16 Production Logging Example

```
1 import logging
2 import os
3 from datetime import datetime
4
5 def save_emergency_dump(predictor, result, asset_id: str):
6     dump_dir = os.path.expanduser("~/predictor_emergency_dumps")
7     os.makedirs(dump_dir, exist_ok=True)
8
9     timestamp = datetime.now().isoformat()
10    dump_file = f"{dump_dir}/{asset_id}_emergency_{timestamp}.msgpack"
11
12    debug_payload = {
13        "emergency_timestamp": timestamp,
14        "asset_id": asset_id,
15        "holder_exponent": float(result.holder_exponent),
16        "weights": [float(w) for w in result.weights],
17        "signal_buffer": predictor._state.signal_circular_buffer.tolist(),
18        "regime_history": predictor._state.cusum_history.tolist(),
19        "telemetry_snapshot": {
20            "kurtosis": float(result.kurtosis),
21            "dgm_entropy": float(result.dgm_entropy),
22            "adaptive_threshold": float(result.adaptive_threshold),
23            "distance_to_collapse": float(result.distance_to_collapse)
24        },
25        "flags_at_emergency": {
26            "degraded_inference": bool(result.degraded_inference_mode),
27            "regime_change": bool(result.regime_change_detected),
28            "mode_collapse": bool(result.mode_collapse_warning)
29        }
30    }
31
32    with open(dump_file, "wb") as f:
33        msgpack.packb(debug_payload, file=f)
34
35    logging.critical(f"Emergency dump saved to {dump_file} for forensic analysis")
```

17 Deterministic Floating-Point Reproducibility

Configure deterministic reductions and PRNG before importing JAX:

```
1 import os
2 import numpy as np
3 import jax
4
5 os.environ['XLA_FLAGS'] = '--xla_cpu_use_cross_replica_callbacks=false'
6 os.environ['JAX_DETERMINISTIC_REDUCTIONS'] = '1'
7 os.environ['JAX_TRACEBACK_FILTERING'] = 'off'
8
9 np.random.seed(42)
10
11 jax.config.update('jax_default_prng_impl', 'threefry2x32')
12 key = jax.random.PRNGKey(42)
13
14 jax.config.update('jax_enable_x64', True)
```

18 Load Shedding (Adaptive Topological Pruning)

When tick rate spikes, dynamically reduce signature depth M based on EWMA latency and jitter. Precompile multiple JIT graphs for $M \in \{2, 3, 5\}$ and switch by thresholds to prevent backlog.

19 Jitter Telemetry

Measure latency jitter using `time.perf_counter_ns()` and degrade if jitter exceeds 80% of Nyquist limit. Expose P95/P99 in telemetry and Prometheus.

20 Dependency Pinning

Strict version pinning is mandatory. Any update must be tested for bit-exact parity and documented. Use exact versions in `requirements.txt` and `environment.yml`, never open ranges.