

Python API Specification - Universal Predictor

Software Engineering

February 20, 2026

Contents

1	Introduction	3
2	Data Structures (Typing)	3
2.1	Configuration (Λ)	3
2.2	Operational Input ($y_t, y_{reference}, \tau$)	3
2.3	System Output	3
3	Multi-Tenant Architecture (Stateless Functional Pattern)	4
3.1	Throughput Maximization (Vectorized Batching)	4
4	Main Class: <code>UniversalPredictor</code> (Stateful Wrapper)	5
4.1	Initialization	5
4.2	Execution Method ($t \rightarrow t + 1$)	5
5	Preventing VRAM Fragmentation (JAX Memory Management)	6
6	VRAM Monitoring	6
7	Recommended Deployment Configuration	6
8	Persistence (Atomic Snapshotting)	7
9	Asynchronous I/O for Snapshots (Non-Blocking)	8
10	Graceful Shutdown for Containers	9
11	Prometheus Integration	10
12	Adaptive CUSUM Threshold	10
13	Grace Period (Post-Regime Refractory Window)	10
14	Operational Flags and Recovery	11
15	Error Handling and Exceptions	11
16	Production Logging Example	11
17	Deterministic Floating-Point Reproducibility	12
18	Load Shedding (Adaptive Topological Pruning)	12
19	Jitter Telemetry	12
20	Dependency Pinning	12

21	Meta-Optimization API (Bayesian Hyperparameter Tuning)	12
21.1	BayesianMetaOptimizer Class	12
21.2	OptimizationResult Schema (Pydantic)	16
21.3	Non-Blocking I/O Execution Pattern	19
21.4	Integration with Prediction Pipeline	21
22	Dependency Pinning	21

1 Introduction

This document specifies the Python implementation of the abstract I/O interface defined in *Stochastic_Predictor_IO*. The API exposes the `UniversalPredictor` class for high-performance environments using JAX for numerical acceleration.

2 Data Structures (Typing)

We use `dataclasses` and `jaxtyping` to enforce immutability and strict dimensional typing for tensors.

2.1 Configuration (Λ)

```
1 from dataclasses import dataclass
2 from typing import Optional
3 from jaxtyping import Float, Array, Bool
4
5 @dataclass(frozen=True)
6 class PredictorConfig:
7     """Hyperparameter vector Lambda."""
8     schema_version: str = "1.0" # Snapshot versioning
9     epsilon: float = 1e-3 # Entropic regularization (Sinkhorn)
10    learning_rate: float = 0.01 # JKO learning rate
11    jko_domain_length: float = 1.0 # Domain length for JKO scaling
12    entropy_window_relaxation_factor: float = 5.0 # Relaxation multiplier
13    entropy_window_bounds_min: int = 10 # Minimum entropy window
14    entropy_window_bounds_max: int = 500 # Maximum entropy window
15    learning_rate_safety_factor: float = 0.8 # Safety factor for learning rate
16    learning_rate_minimum: float = 1e-6 # Minimum learning rate
17    log_sig_depth: int = 3 # Signature depth (Kernel D)
18    wtmm_buffer_size: int = 128 # WTMM buffer size
19    besov_cone_c: float = 1.5 # Besov cone influence
20    koopman_top_k: int = 5 # Top-K Koopman spectral modes
21    koopman_min_power: float = 1e-10 # Minimum Koopman spectral power
22    paley_wiener_integral_max: float = 100.0 # Paley-Wiener threshold
23    kernel_c_jump_intensity: float = 0.05 # Levy jump intensity
24    kernel_c_jump_mean: float = 0.0 # Levy jump mean
25    kernel_c_jump_scale: float = 0.1 # Levy jump scale
26    kernel_c_jump_max_events: int = 16 # Max jump events per step
27    holder_threshold: float = 0.4 # Circuit breaker threshold
28    cusum_h: float = 5.0 # CUSUM threshold
29    cusum_k: float = 0.5 # CUSUM slack
30    grace_period_steps: int = 20 # Post-regime refractory period
31    volatility_alpha: float = 0.1 # EMA decay for variance
32
33    # Load shedding and anti-aliasing
34    staleness_ttl_ns: int = 500_000_000 # TTL (500ms)
35    besov_nyquist_interval_ns: int = 100_000_000 # Nyquist soft limit (100ms)
36    inference_recovery_hysteresis: float = 0.8 # Degraded mode recovery factor
```

2.2 Operational Input ($y_t, y_{\text{reference}}, \tau$)

```
1 @dataclass(frozen=True)
2 class ProcessState:
3     magnitude: Float[Array, "1"] # y_t (normalized or absolute)
4     reference: Float[Array, "1"] # y_reference
5     timestamp_ns: int # Unix epoch (nanoseconds)
6
7     def validate_domain(self, sigma_bound: float = 20.0, sigma_val: float = 1.0) -> bool:
8         """Catastrophic outlier detection (> 20 sigma)."""
9         return abs(self.magnitude) <= (sigma_bound * sigma_val)
```

2.3 System Output

```
1 @dataclass(frozen=True)
2 class PredictionResult:
3     predicted_next: Float[Array, "1"] # y_{t+1} (Z-score)
```

```

4
5     # Telemetry
6     holder_exponent: Float[Array, "1"]
7     cusum_drift: Float[Array, "1"]
8     distance_to_collapse: Float[Array, "1"]
9     free_energy: Float[Array, "1"]
10
11    # Advanced telemetry
12    kurtosis: Float[Array, "1"]
13    dgm_entropy: Float[Array, "1"]
14    adaptive_threshold: Float[Array, "1"]
15
16    # Orchestrator state
17    weights: Float[Array, "4"]
18
19    # Health flags
20    sinkhorn_converged: Bool[Array, "1"]
21    degraded_inference_mode: bool
22    emergency_mode: bool
23    regime_change_detected: bool
24    modeCollapse_warning: bool
25
26    mode: str # "Standard" | "Robust" | "Emergency"

```

3 Multi-Tenant Architecture (Stateless Functional Pattern)

To support hundreds of assets on a single server, the API exposes a purely functional mode. This allows state management in low-latency external storage (Redis) while sharing the compiled JAX graph across assets.

3.1 Throughput Maximization (Vectorized Batching)

This architecture enables `jax.vmap` to batch multiple asset states in a single hardware call, minimizing the Python GIL impact and maximizing GPU occupancy.

```

1 class FunctionalPredictor:
2     """
3         Stateless implementation for JAX core.
4         Scales to thousands of predictors sharing the same graph.
5     """
6     def __init__(self, config: PredictorConfig):
7         self.config = config
8         self._core_step = self._core_update_step
9         self._jit_update = jax.jit(self._core_step)
10        self._vmap_update = jax.jit(jax.vmap(self._core_step, in_axes=(0, 0, 0, 0)))
11
12    def init_state(self):
13        """Create a zeroed cold-state structure."""
14        return self._initialize_state_structure()
15
16    def step(self, state, obs: ProcessState) -> tuple[object, PredictionResult]:
17        """
18            Pure state transition: (S_t, Obs_t) -> (S_{t+1}, Pred_{t+1})
19        """
20        should_freeze = self._should_freeze(obs)
21        new_state, raw_result = self._jit_update(
22            state,
23            obs.magnitude,
24            obs.reference,
25            freeze_weights=should_freeze
26        )
27        result = PredictionResult(
28            predicted_next=raw_result.y_next,
29            # map the remaining fields
30        )
31        return new_state, result
32
33    def step_batch(self, states, obs_batch: ProcessState):
34        """Vectorized batch processing for N assets."""
35        freeze_flags = self._should_freeze_batch(obs_batch)

```

```

36     new_states, results = self._vmap_update(
37         states, obs_batch.magnitude, obs_batch.reference, freeze_flags
38     )
39     return new_states, results

```

4 Main Class: UniversalPredictor (Stateful Wrapper)

This class wraps the functional pattern for single-tenant usage with state held in local memory.

4.1 Initialization

```

1  class UniversalPredictor:
2      def __init__(self, config: PredictorConfig):
3          """
4              Initialize the JAX compute graph (XLA JIT compilation).
5              Allocate static device buffers (VRAM).
6              Internal state stores persistent rolling buffers updated with
7              functional ops to avoid CPU<->VRAM transfers.
8          """
9          self.config = config
10         self._state = self._initialize_state()
11         self._jit_update = jax.jit(self._core_update_step)
12         self._last_timestamp_ns = 0
13
14     def fit_history(self, history: list[float]) -> bool:
15         """
16             Cold-start bootstrapping. Requires at least N_buf samples.
17             Returns True if Sinkhorn and CUSUM converge.
18         """
19         if len(history) < self.config.wtmm_buffer_size:
20             raise ValueError(f"Insufficient history. Required: {self.config.wtmm_buffer_size}")
21
22         self._state, final_metrics = self._jit_scan_history(self._state, jnp.array(history))
23
24         is_converged = final_metrics.sinkhorn_converged
25         is_stable = final_metrics.cusum_drift < self.config.cusum_h
26         if not (is_converged and is_stable):
27             logger.warning("Cold start finished without stable convergence.")
28             return False
29         return True

```

4.2 Execution Method ($t \rightarrow t + 1$)

```

1  def step(self, obs: ProcessState) -> PredictionResult:
2      """Execute one prediction cycle with domain and TTL validation."""
3      if not obs.validate_domain():
4          logger.error("Catastrophic outlier detected. Ignoring tick.")
5          return self._last_valid_result
6
7      current_time = time.time_ns()
8      latency = current_time - obs.timestamp_ns
9      is_stale = latency > self.config.staleness_ttl_ns
10
11     dt_arrival = obs.timestamp_ns - self._last_timestamp_ns
12     is_sparse = (self._last_timestamp_ns > 0) and (
13         dt_arrival > self.config.besov_nyquist_interval_ns
14     )
15     if is_sparse:
16         logger.warning(
17             f"FrequencyWarning: interval {dt_arrival}ns > Nyquist limit. WTMM may alias."
18         )
19
20     self._last_timestamp_ns = obs.timestamp_ns
21     should_freeze = is_stale or is_sparse
22
23     new_state, result_data = self._jit_update(
24         self._state,

```

```

25         obs.magnitude,
26         obs.reference,
27         freeze_weights=should_freeze,
28     )
29     self._state = new_state
30
31     return PredictionResult(
32         predicted_next=result_data.y_next,
33         holder_exponent=result_data.H_t,
34         sinkhorn_converged=result_data.converged,
35         # map remaining fields
36     )

```

5 Preventing VRAM Fragmentation (JAX Memory Management)

Production problem: JAX preallocates 90% of GPU memory on first access. Long-running systems may fragment VRAM and hit silent OOM after weeks.

Solution: Configure environment variables **before** importing JAX:

```

1 import os
2
3 os.environ['XLA_PYTHON_CLIENT_MEM_FRACTION'] = '0.7'
4 os.environ['XLA_PYTHON_CLIENT_ALLOCATOR'] = 'platform'
5 os.environ['TF_FORCE_GPU_ALLOW_GROWTH'] = 'true'
6
7 import jax
8 import jax.numpy as jnp

```

6 VRAM Monitoring

```

1 import psutil
2 import subprocess
3
4 def monitor_vram_fragmentation(interval_seconds=60):
5     """Background thread for VRAM monitoring."""
6     import time
7     import threading
8
9     def _monitor():
10         while True:
11             try:
12                 result = subprocess.run(
13                     ['nvidia-smi', '--query-gpu=memory.used,memory.total',
14                      '--format=csv,nounits,noheader'],
15                     capture_output=True, text=True, timeout=5
16                 )
17                 if result.returncode == 0:
18                     used, total = map(float, result.stdout.strip().split(','))
19                     utilization = 100.0 * used / total
20                     if utilization > 0.95:
21                         print(f"[WARNING] VRAM near saturation: {utilization:.1f}%")
22                     elif utilization > 0.85:
23                         print(f"[INFO] VRAM utilization: {utilization:.1f}% (elevated)")
24                     time.sleep(interval_seconds)
25                 except Exception as e:
26                     print(f"[ERROR] VRAM monitoring failed: {e}")
27                     break
28
29             thread = threading.Thread(target=_monitor, daemon=True)
30             thread.start()

```

7 Recommended Deployment Configuration

```

1 #!/bin/bash
2 # deployment/run_predictor.sh
3
4 export XLA_PYTHON_CLIENT_MEM_FRACTION=0.7
5 export XLA_PYTHON_CLIENT_ALLOCATOR=platform
6 export TF_FORCE_GPU_ALLOW_GROWTH=true
7
8 echo "[INFO] XLA VRAM Fraction: 0.7 (28/40 GB on A100)"
9 echo "[INFO] Allocator: platform (dynamic)"
10 echo "[INFO] GPU growth: enabled"
11
12 python3 -u predictor_service.py \
13     --config config.yaml \
14     --device gpu \
15     --pool-size 100 \
16     --monitor-interval 300

```

8 Persistence (Atomic Snapshotting)

```

1 import hashlib
2 import msgpack
3
4 def save_snapshot(self, filepath: str):
5     """
6         Export internal state Sigma_t as MessagePack.
7         Append SHA-256 checksum at the end of the file.
8     """
9     state_dict = self._serialize_jax_state(self._state)
10    payload = {
11        "schema_version": self.config.schema_version,
12        "timestamp": time.time_ns(),
13        "config": asdict(self.config),
14        "global": state_dict["global"],
15        "telemetry": {
16            "kurtosis": float(self._state.kurtosis),
17            "dgm_entropy": float(self._state.dgm_entropy),
18            "adaptive_threshold": float(self._state.h_adaptive)
19        },
20        "flags": {
21            "degraded_inference": bool(self._state.degraded_mode),
22            "emergency": bool(self._state.emergency_mode),
23            "regime_change": bool(self._state.regime_changed),
24            "mode_collapse": bool(self._state.mode_collapse_warning)
25        },
26        "kernels": {
27            "A": state_dict["kernel_a"],
28            "B": state_dict["kernel_b"],
29            "C": state_dict["kernel_c"],
30            "D": state_dict["kernel_d"]
31        }
32    }
33    data_bytes = msgpack.packb(payload)
34    checksum = hashlib.sha256(data_bytes).hexdigest()
35
36    with open(filepath, "wb") as f:
37        f.write(data_bytes)
38        f.write(checksum.encode('utf-8'))
39
40
41 def load_snapshot(self, filepath: str):
42     """
43         Load state. Validate SHA-256 and schema_version.
44         Raise ValueError if validation fails.
45     """
46     with open(filepath, "rb") as f:
47         content = f.read()
48
49     data_bytes = content[:-64]
50     stored_checksum = content[-64:].decode('utf-8')
51
52     computed = hashlib.sha256(data_bytes).hexdigest()

```

```
53     if computed != stored_checksum:
54         raise ValueError("Snapshot corrupt: checksum mismatch.")
55
56     payload = msgpack.unpackb(data_bytes)
57     loaded_schema = payload.get('schema_version', 'unknown')
58     if loaded_schema != self.config.schema_version:
59         raise ValueError(
60             f"Schema version mismatch: snapshot={loaded_schema}, current={self.config.
61             schema_version}." )
62
63     self._state = self._deserialize_jax_state(payload)
```

9 Asynchronous I/O for Snapshots (Non-Blocking)

```
1 import concurrent.futures
2 import hashlib
3 import msgpack
4 import threading
5 import time
6
7 class UniversalPredictor_AsyncIO:
8     def __init__(self, n_worker_threads=2):
9         self.io_executor = concurrent.futures.ThreadPoolExecutor(
10             max_workers=n_worker_threads,
11             thread_name_prefix="snapshot_io_"
12         )
13         self.pending_snapshot_future = None
14         self.snapshot_lock = threading.Lock()
15
16     def _compute_and_save_async(self, filepath: str, data_bytes: bytes):
17         checksum = hashlib.sha256(data_bytes).hexdigest()
18         temp_filepath = filepath + ".tmp"
19         try:
20             with open(temp_filepath, "wb") as f:
21                 f.write(data_bytes)
22                 f.write(checksum.encode('utf-8'))
23             import os
24             os.replace(temp_filepath, filepath)
25             return {
26                 'status': 'success',
27                 'filepath': filepath,
28                 'filesize_bytes': len(data_bytes),
29                 'checksum': checksum,
30                 'timestamp': time.time()
31             }
32         except Exception as e:
33             return {
34                 'status': 'error',
35                 'filepath': filepath,
36                 'error': str(e),
37                 'timestamp': time.time()
38             }
39
40     def save_snapshot_nonblocking(self, filepath: str) -> concurrent.futures.Future:
41         state_dict = self._serialize_jax_state(self._state)
42         payload = {
43             "schema_version": self.config.schema_version,
44             "timestamp": time.time_ns(),
45             "config": asdict(self.config),
46             "global": state_dict["global"],
47             "telemetry": {
48                 "kurtosis": float(self._state.kurtosis),
49                 "dgm_entropy": float(self._state.dgm_entropy),
50                 "adaptive_threshold": float(self._state.h_adaptive)
51             },
52             "flags": {
53                 "degraded_inference": bool(self._state.degraded_mode),
54                 "emergency": bool(self._state.emergency_mode),
55                 "regime_change": bool(self._state.regime_changed),
56                 "mode_collapse": bool(self._state.mode_collapse_warning)
57             }
58         }
59
60         return self.io_executor.submit(self._compute_and_save_async, filepath, payload)
```

```

57     },
58     "kernels": {
59         "A": state_dict["kernel_a"],
60         "B": state_dict["kernel_b"],
61         "C": state_dict["kernel_c"],
62         "D": state_dict["kernel_d"]
63     }
64 }
65 data_bytes = msgpack.packb(payload)
66 future = self.io_executor.submit(self._compute_and_save_async, filepath, data_bytes)
67 with self.snapshot_lock:
68     self.pending_snapshot_future = future
69 return future

```

10 Graceful Shutdown for Containers

```

1 import signal
2 import sys
3 import threading
4 import time
5 import logging
6 from typing import Optional
7
8 class UniversalPredictor_GracefulShutdown:
9     def __init__(self, config: PredictorConfig):
10         self.config = config
11         self.predictor = UniversalPredictor_AsyncIO(config)
12         self.shutdown_requested = threading.Event()
13         self.is_accepting_data = True
14         self.input_buffer_lock = threading.Lock()
15         self.residual_buffer = []
16
17         signal.signal(signal.SIGTERM, self._handle_sigterm)
18         signal.signal(signal.SIGINT, self._handle_sigint)
19
20         self.logger = logging.getLogger("predictor.shutdown")
21         self.logger.info("[INIT] Graceful shutdown handler registered")
22
23     def _handle_sigterm(self, signum, frame):
24         self.logger.warning(f"[SIGTERM] Received signal {signum}. Initiating graceful shutdown")
25         self.shutdown_requested.set()
26
27     def _handle_sigint(self, signum, frame):
28         self.logger.warning(f"[SIGINT] Received signal {signum}. Initiating graceful shutdown")
29         self.shutdown_requested.set()
30
31     def accept_observation(self, obs: ProcessState) -> Optional[PredictionResult]:
32         if self.shutdown_requested.is_set() or not self.is_accepting_data:
33             self.logger.warning(f"[REJECT] Observation rejected (shutdown in progress): {obs.timestamp_ns}")
34             return None
35         with self.input_buffer_lock:
36             self.residual_buffer.append(obs)
37         return self._process_observation(obs)
38
39     def _process_observation(self, obs: ProcessState) -> PredictionResult:
40         result = self.predictor.predict_next(obs)
41         with self.input_buffer_lock:
42             if obs in self.residual_buffer:
43                 self.residual_buffer.remove(obs)
44         return result
45
46     def graceful_shutdown(self, timeout_seconds: int = 25):
47         shutdown_start = time.time()
48         self.logger.info("GRACEFUL SHUTDOWN INITIATED")
49         self.is_accepting_data = False
50         time.sleep(0.1)
51
52         with self.input_buffer_lock:

```

```

53     for obs in list(self.residual_buffer):
54         if time.time() - shutdown_start > timeout_seconds - 10:
55             self.logger.warning("Timeout approaching, aborting residual processing")
56             break
57         try:
58             _ = self._process_observation(obs)
59         except Exception as e:
60             self.logger.error(f"Error processing residual: {e}")
61
62     pending_snapshot = self.predictor.pending_snapshot_future
63     if pending_snapshot is not None and not pending_snapshot.done():
64         try:
65             remaining_time = max(1, timeout_seconds - (time.time() - shutdown_start))
66             pending_snapshot.result(timeout=remaining_time)
67         except Exception as e:
68             self.logger.error(f"Async snapshot failed: {e}")
69
70     try:
71         final_snapshot_path = f"snapshots/shutdown_{int(time.time())}.pkl"
72         self.predictor.save_snapshot(final_snapshot_path)
73     except Exception as e:
74         self.logger.error(f"Final snapshot failed: {e}")
75
76     try:
77         if hasattr(self.predictor, 'io_executor'):
78             self.predictor.io_executor.shutdown(wait=True, cancel_futures=False)
79     except Exception as e:
80         self.logger.error(f"Error closing resources: {e}")
81
82     total_time = time.time() - shutdown_start
83     self.logger.info(f"SHUTDOWN COMPLETED ({total_time:.2f}s)")
84     sys.exit(0)

```

11 Prometheus Integration

```

1 from prometheus_client import Counter, Histogram
2
3 class UniversalPredictor_GracefulShutdown_Monitored:
4     def __init__(self, config: PredictorConfig):
5         self.shutdown_counter = Counter(
6             'predictor_graceful_shutdowns_total',
7             'Total number of graceful shutdowns executed'
8         )
9         self.shutdown_duration = Histogram(
10            'predictor_shutdown_duration_seconds',
11            'Time taken to complete graceful shutdown',
12            buckets=[0.5, 1.0, 2.0, 5.0, 10.0, 20.0, 30.0]
13        )
14         self.residual_observations = Histogram(
15            'predictor_shutdown_residual_observations',
16            'Number of observations in buffer during shutdown',
17            buckets=[0, 1, 5, 10, 50, 100, 500, 1000]
18        )

```

12 Adaptive CUSUM Threshold

The system implements the adaptive threshold lemma based on kurtosis:

$$h_t = k \cdot \sigma_t \cdot \left(1 + \ln\left(\frac{\kappa_t}{3}\right)\right)$$

13 Grace Period (Post-Regime Refractory Window)

After a regime change ($G^+ > h_t$), the system resets weights to uniform. A grace period prevents a cascade of false alarms while weights re-converge. The detector continues to compute G^+ but does not emit an alarm until the counter expires.

14 Operational Flags and Recovery

The system exposes explicit flags:

- `degraded_inference_mode`: TTL exceeded; weights frozen.
- `emergency_mode`: $H_t < H_{min}$; force Kernel D and Huber loss.
- `regime_change_detected`: CUSUM alarm; entropy reset.
- `mode_collapse_warning`: DGM entropy below threshold for > 10 steps.

15 Error Handling and Exceptions

Standard alerts:

- `DomainError`: catastrophic outlier $> 20\sigma$
- `StalenessWarning`: TTL exceeded
- `FrequencyWarning`: Nyquist limit violated
- `IntegrityError`: snapshot verification failed

16 Production Logging Example

```
1 import logging
2 import os
3 from datetime import datetime
4
5 def save_emergency_dump(predictor, result, asset_id: str):
6     dump_dir = os.path.expanduser("~/predictor_emergency_dumps")
7     os.makedirs(dump_dir, exist_ok=True)
8
9     timestamp = datetime.now().isoformat()
10    dump_file = f"{dump_dir}/{asset_id}_emergency_{timestamp}.msgpack"
11
12    debug_payload = {
13        "emergency_timestamp": timestamp,
14        "asset_id": asset_id,
15        "holder_exponent": float(result.holder_exponent),
16        "weights": [float(w) for w in result.weights],
17        "signal_buffer": predictor._state.signal_circular_buffer.tolist(),
18        "regime_history": predictor._state.cusum_history.tolist(),
19        "telemetry_snapshot": {
20            "kurtosis": float(result.kurtosis),
21            "dgm_entropy": float(result.dgm_entropy),
22            "adaptive_threshold": float(result.adaptive_threshold),
23            "distance_to_collapse": float(result.distance_to_collapse)
24        },
25        "flags_at_emergency": {
26            "degraded_inference": bool(result.degraded_inference_mode),
27            "regime_change": bool(result.regime_change_detected),
28            "mode_collapse": bool(result.mode_collapse_warning)
29        }
30    }
31
32    with open(dump_file, "wb") as f:
33        msgpack.packb(debug_payload, file=f)
34
35    logging.critical(f"Emergency dump saved to {dump_file} for forensic analysis")
```

17 Deterministic Floating-Point Reproducibility

Configure deterministic reductions and PRNG before importing JAX:

```
1 import os
2 import numpy as np
3 import jax
4
5 os.environ['XLA_FLAGS'] = '--xla_cpu_use_cross_replica_callbacks=false'
6 os.environ['JAX_DETERMINISTIC_REDUCTIONS'] = '1'
7 os.environ['JAX_TRACEBACK_FILTERING'] = 'off'
8
9 np.random.seed(42)
10
11 jax.config.update('jax_default_prng_impl', 'threefry2x32')
12 key = jax.random.PRNGKey(42)
13
14 jax.config.update('jax_enable_x64', True)
```

18 Load Shedding (Adaptive Topological Pruning)

When tick rate spikes, dynamically reduce signature depth M based on EWMA latency and jitter. Precompile multiple JIT graphs for $M \in \{2, 3, 5\}$ and switch by thresholds to prevent backlog.

19 Jitter Telemetry

Measure latency jitter using `time.perf_counter_ns()` and degrade if jitter exceeds 80% of Nyquist limit. Expose P95/P99 in telemetry and Prometheus.

20 Dependency Pinning

Strict version pinning is mandatory. Any update must be tested for bit-exact parity and documented. Use exact versions in `requirements.txt` and `environment.yml`, never open ranges.

21 Meta-Optimization API (Bayesian Hyperparameter Tuning)

To support the autonomous Level 4 operation defined in `Stochastic_Predictor_Implementation.tex` (tiered meta-optimization), the API must expose contracts for persistent, resumable Bayesian optimization.

21.1 BayesianMetaOptimizer Class

The `BayesianMetaOptimizer` encapsulates the Tree-structured Parzen Estimator (TPE) logic for both Fast Tuning and Deep Tuning regimes.

```
1 from typing import Callable, Dict, Any, Optional
2 from dataclasses import dataclass
3 import pickle
4 import hashlib
5 from pathlib import Path
6
7 @dataclass(frozen=True)
8 class SearchSpace:
9     """Defines hyperparameter search space with constraints."""
10    name: str
11    param_type: str # "float", "int", "categorical", "log_uniform"
12    range: tuple[Any, Any] # (min, max) or list of choices
13    locked: bool = False # Immutable parameters (security/io sections)
14    constraint: Optional[str] = None # "must be power of 2", etc.
15
16 class BayesianMetaOptimizer:
17     """
18         Resumable Bayesian optimization using TPE algorithm.
19         Supports checkpointing for long-running Deep Tuning campaigns.
```

```

20     """
21
22     def __init__(self,
23                  search_space: Dict[str, SearchSpace],
24                  objective_fn: Callable[[Dict[str, Any]], float],
25                  study_name: str,
26                  max_iterations: int,
27                  tier: str = "fast" # "fast" or "deep"
28                  ):
29         """
30         Initialize optimizer with search space and objective function.
31
32     Args:
33         search_space: Dictionary of parameter names to SearchSpace definitions
34         objective_fn: Walk-forward validation function returning MAPE
35         study_name: Unique identifier for this optimization campaign
36         max_iterations: Budget (50 for Fast, 500 for Deep)
37         tier: Optimization tier ("fast" or "deep")
38
39     Raises:
40         ValueError: If search_space contains locked parameters
41     """
42
43     self.search_space = self._validate_search_space(search_space)
44     self.objective_fn = objective_fn
45     self.study_name = study_name
46     self.max_iterations = max_iterations
47     self.tier = tier
48
49     # Internal TPE state (optuna.Study or similar)
50     self._study = None
51     self._best_params = None
52     self._best_value = float('inf')
53     self._iteration = 0
54     self._checkpoint_counter = 0
55
56     def _validate_search_space(self, space: Dict[str, SearchSpace]) -> Dict[str, SearchSpace]:
57         """Verify no locked parameters in search space."""
58         locked_params = [name for name, spec in space.items() if spec.locked]
59         if locked_params:
60             raise ValueError(
61                 f"Cannot optimize locked parameters: {locked_params}. "
62                 f"Remove from search space or set locked=False."
63             )
64         return space
65
66     def optimize(
67                 self,
68                 checkpoint_interval: int = 10,
69                 early_stopping_patience: int = 50
70                 ) -> Dict[str, Any]:
71         """
72         Execute Bayesian optimization with automatic checkpointing.
73
74     Args:
75         checkpoint_interval: Emit checkpoint every N trials
76         early_stopping_patience: Stop if no improvement for N trials
77
78     Returns:
79         Best hyperparameter configuration found
80
81     Note:
82         This method blocks until completion or early stopping.
83         For long-running Deep Tuning, consider running in separate process.
84     """
85     no_improvement_count = 0
86
87     for i in range(self._iteration, self.max_iterations):
88         # Sample next candidate from TPE surrogate
89         candidate = self._suggest_next_candidate()
90
91         # Evaluate via walk-forward validation (expensive!)
92         objective_value = self.objective_fn(candidate)

```

```

93     # Update TPE model
94     self._report_trial(candidate, objective_value)
95
96     # Track best result
97     if objective_value < self._best_value:
98         self._best_value = objective_value
99         self._best_params = candidate
100        no_improvement_count = 0
101        # Checkpoint immediately on improvement
102        self.save_study(f"io/snapshots/{self.study_name}_best.pkl")
103    else:
104        no_improvement_count += 1
105
106    self._iteration = i + 1
107
108    # Periodic checkpointing
109    if (i + 1) % checkpoint_interval == 0:
110        checkpoint_path = f"io/snapshots/{self.study_name}_iter{i+1}.pkl"
111        self.save_study(checkpoint_path)
112
113    # Early stopping
114    if no_improvement_count >= early_stopping_patience:
115        print(f"Early stopping: No improvement for {early_stopping_patience} trials")
116        break
117
118
119    return self._best_params
120
121 def save_study(self, path: str) -> None:
122 """
123     Serialize TPE study state to disk for resumability.
124
125     Implementation must guarantee atomic write via temporary file + os.replace().
126     Includes SHA-256 hash for integrity verification on load.
127
128     Args:
129         path: Target checkpoint file path (e.g., "io/snapshots/study.pkl")
130
131     Protocol:
132         1. Serialize study state to temporary file
133         2. Compute SHA-256 hash of serialized data
134         3. Atomically replace target file (POSIX os.replace)
135         4. Store hash in metadata sidecar file
136
137     Note:
138         This operation is I/O-bound and may block for 100-500ms.
139         For production systems running live prediction, execute in
140         separate thread or process to avoid blocking telemetry collection.
141
142     Example:
143         >>> optimizer.save_study("checkpoints/deep_tuning_iter250.pkl")
144     """
145     path_obj = Path(path)
146     path_obj.parent.mkdir(parents=True, exist_ok=True)
147
148     # Prepare checkpoint payload
149     checkpoint_data = {
150         'study_name': self.study_name,
151         'search_space': self.search_space,
152         'tier': self.tier,
153         'iteration': self._iteration,
154         'best_params': self._best_params,
155         'best_value': self._best_value,
156         'trial_history': self._study.trials if self._study else [],
157         'parzen_estimators': self._study._storage if self._study else None,
158         'rng_state': self._get_rng_state(),
159         'timestamp': time.time_ns()
160     }
161
162     # Serialize to temporary file (atomic write protocol)
163     tmp_path = path_obj.with_suffix('.tmp')
164     with open(tmp_path, 'wb') as f:
165         serialized = pickle.dumps(checkpoint_data, protocol=pickle.HIGHEST_PROTOCOL)

```

```

166         f.write(serialized)
167         f.flush()
168         os.fsync(f.fileno()) # Force kernel buffer flush
169
170     # Compute integrity hash
171     with open(tmp_path, 'rb') as f:
172         hash_value = hashlib.sha256(f.read()).hexdigest()
173
174     # Atomic replacement (POSIX guarantee)
175     os.replace(tmp_path, path)
176
177     # Store hash in sidecar
178     hash_path = path_obj.with_suffix('.pkl.sha256')
179     with open(hash_path, 'w') as f:
180         f.write(f"{hash_value} {path_obj.name}\n")
181
182 def load_study(self, path: str) -> None:
183     """
184     Deserialize TPE study state from checkpoint.
185
186     Verifies SHA-256 hash before loading to detect corruption.
187     Reconstructs Parzen estimators and random state for exact resumption.
188
189     Args:
190         path: Checkpoint file path
191
192     Raises:
193         IntegrityError: If SHA-256 verification fails
194         FileNotFoundError: If checkpoint or hash file missing
195         ValueError: If checkpoint schema version incompatible
196
197     Protocol:
198         1. Verify SHA-256 hash matches expected value
199         2. Deserialize checkpoint data
200         3. Reconstruct TPE study object
201         4. Restore RNG state for deterministic sampling
202         5. Validate search space matches current configuration
203
204     Note:
205         After successful load, optimizer continues from iteration N+1.
206         No re-evaluation of previous trials occurs (warm start).
207
208     Example:
209         >>> optimizer = BayesianMetaOptimizer(...)
210         >>> optimizer.load_study("checkpoints/deep_tuning_iter250.pkl")
211         >>> optimizer.optimize() # Resumes from iteration 251
212     """
213     path_obj = Path(path)
214     hash_path = path_obj.with_suffix('.pkl.sha256')
215
216     # Verify integrity
217     if not hash_path.exists():
218         raise FileNotFoundError(f"Hash file missing: {hash_path}")
219
220     with open(hash_path, 'r') as f:
221         expected_hash = f.read().strip().split()[0]
222
223     with open(path, 'rb') as f:
224         actual_hash = hashlib.sha256(f.read()).hexdigest()
225
226     if actual_hash != expected_hash:
227         raise IntegrityError(
228             f"Checkpoint corrupted: hash mismatch. "
229             f"Expected {expected_hash}, got {actual_hash}"
230         )
231
232     # Deserialize
233     with open(path, 'rb') as f:
234         checkpoint_data = pickle.load(f)
235
236     # Validate schema
237     if checkpoint_data['study_name'] != self.study_name:
238         raise ValueError(

```

```

239         f"Study name mismatch: checkpoint is for '{checkpoint_data['study_name']}', "
240         f"but optimizer is '{self.study_name}'"
241     )
242
243     # Restore state
244     self._iteration = checkpoint_data['iteration']
245     self._best_params = checkpoint_data['best_params']
246     self._best_value = checkpoint_data['best_value']
247
248     # Reconstruct TPE study (replay trials)
249     self._study = self._create_study()
250     for trial_data in checkpoint_data['trial_history']:
251         self._study.add_trial(trial_data)
252
253     # Restore RNG state for deterministic continuation
254     self._restore_rng_state(checkpoint_data['rng_state'])
255
256     print(f'Resumed from iteration {self._iteration}, best value: {self._best_value:.6f}')

```

21.2 OptimizationResult Schema (Pydantic)

The result of a meta-optimization campaign must be exportable to `config.toml` using the atomic mutation protocol.

```

1  from pydantic import BaseModel, validator
2  from typing import Dict, Any, Optional
3  import toml
4  import os
5  import time
6
7  class OptimizationResult(BaseModel):
8      """
9          Immutable result of Bayesian optimization campaign.
10         Includes export capability for atomic config mutation.
11      """
12
13     study_name: str
14     tier: str # "fast" or "deep"
15     best_params: Dict[str, Any]
16     best_objective: float
17     total_iterations: int
18     early_stopped: bool
19     convergence_delta: float # Improvement over last N trials
20     timestamp_utc: str
21
22     @validator('tier')
23     def validate_tier(cls, v):
24         if v not in ['fast', 'deep']:
25             raise ValueError(f"Invalid tier: {v}. Must be 'fast' or 'deep'")
26         return v
27
28     def export_to_toml(
29         self,
30         path: str = "config.toml",
31         backup: bool = True,
32         validate: bool = True
33     ) -> None:
34         """
35             Export optimized parameters to config.toml using atomic mutation protocol.
36
37             Implements the Configuration Mutation Protocol specified in
38             Stochastic_Predictor_IO.tex §3.3.
39
40             Args:
41                 path: Target config file path (default: "config.toml")
42                 backup: Create timestamped backup before mutation (default: True)
43                 validate: Validate merged config against schema (default: True)
44
45             Protocol:
46                 1. Validate new parameters against schema (ranges, types, constraints)
47                 2. Create immutable backup (config.toml.bak + timestamped archive)
48                 3. Write to temporary file (config.toml.tmp)
49                 4. Fsync to guarantee durability

```

```

49      5. Atomic replacement via os.replace()
50      6. Log mutation to audit trail (io/mutations.log)
51
52      Locked Subsections (Never Modified):
53          - [io]: snapshot_path, telemetry_buffer_maxlen, credentials_vault_path
54          - [security]: telemetry_hash_interval_steps, snapshot_integrity_hash_algorithm
55          - [core]: float_precision, jax_platform (partial lock)
56          - [meta_optimization]: max_deep_tuning_iterations, checkpoint_path
57
58      Raises:
59          ConfigMutationError: If validation fails or locked parameter modified
60          IOError: If atomic write fails (disk full, permissions, etc.)
61
62      Note:
63          This operation blocks for ~50-200ms due to fsync requirement.
64          For production systems, execute in separate thread/process.
65
66      Example:
67          >>> result = optimizer.optimize()
68          >>> opt_result = OptimizationResult(
69              ...     study_name="deep_tuning_2026",
70              ...     tier="deep",
71              ...     best_params=result,
72              ...     best_objective=0.0234,
73              ...
74              ...
75          )
76          >>> # Non-blocking export in separate thread
77          >>> import threading
78          >>> export_thread = threading.Thread(
79              ...     target=opt_result.export_to_toml,
80              ...     kwargs={'backup': True, 'validate': True}
81              ...
82          )
83          >>> export_thread.start()
84
85      """
86      path_obj = Path(path)
87      if not path_obj.exists():
88          raise FileNotFoundError(f"Config file not found: {path}")
89
90      # Phase 1: Load and merge
91      current_config = toml.load(path)
92      merged_config = self._merge_with_validation(current_config, validate)
93
94      # Phase 2: Backup
95      if backup:
96          timestamp = time.strftime("%Y-%m-%dT%H:%M:%SZ", time.gmtime())
97          backup_timestamped = path_obj.with_suffix(f'.bak.{timestamp}')
98          backup_latest = path_obj.with_suffix('.bak')
99
100     import shutil
101     shutil.copy2(path, backup_timestamped)
102     shutil.copy2(path, backup_latest)
103
104     # Phase 3: Atomic write via temporary file
105     tmp_path = path_obj.with_suffix('.tmp')
106
107     # Prevent concurrent mutations
108     if tmp_path.exists():
109         raise IOError(
110             f"Concurrent mutation detected: {tmp_path} exists. "
111             f"Another optimizer may be writing. Aborting."
112         )
113
114         with open(tmp_path, 'w') as f:
115             toml.dump(merged_config, f)
116             f.flush()
117             os.fsync(f.fileno()) # CRITICAL: Force kernel buffer flush
118
119     # Phase 4: Atomic replacement (POSIX guarantee)
120     os.replace(tmp_path, path)

# Phase 5: Audit logging
delta = self._compute_delta(current_config, merged_config)

```

```

121         self._log_mutation(delta, timestamp if backup else time.strftime("%Y-%m-%dT%H:%M:%SZ"))
122     )
123
124     def _merge_with_validation(
125         self,
126         current_config: Dict[str, Any],
127         validate: bool
128     ) -> Dict[str, Any]:
129         """
130             Merge optimized parameters into current config with validation.
131
132             Prevents modification of locked subsections.
133             Validates ranges and constraints.
134         """
135
136         merged = current_config.copy()
137
138         # Determine target subsection based on tier
139         target_section = "sensitivity" if self.tier == "fast" else "structural"
140
141         if target_section not in merged:
142             merged[target_section] = {}
143
144         for param_name, param_value in self.best_params.items():
145             # Check if parameter is in locked subsection
146             if self._is_locked_parameter(param_name):
147                 raise ConfigMutationError(
148                     f"Attempted to modify locked parameter: {param_name}. "
149                     f"This violates invariant protection rules."
150                 )
151
152             merged[target_section][param_name] = param_value
153
154         if validate:
155             self._validate_config(merged)
156
157     return merged
158
159     def _is_locked_parameter(self, param_name: str) -> bool:
160         """Check if parameter belongs to locked subsection."""
161         locked_params = {
162             'snapshot_path', 'telemetry_buffer maxlen', 'credentials_vault_path',
163             'telemetry_hash_interval_steps', 'snapshot_integrity_hash_algorithm',
164             'float_precision', 'jax_platform',
165             'max_deep_tuning_iterations', 'checkpoint_path'
166         }
167
168         return param_name in locked_params
169
170     def _validate_config(self, config: Dict[str, Any]) -> None:
171         """Validate merged config against schema."""
172         # Implementation: Check ranges, types, constraints
173         # Raise ConfigMutationError if validation fails
174         pass
175
176     def _compute_delta(
177         self,
178         old_config: Dict[str, Any],
179         new_config: Dict[str, Any]
180     ) -> Dict[str, Any]:
181         """
182             Compute parameter delta for audit logging.
183         """
184         delta = {}
185         # Compare configs and extract changes
186         return delta
187
188     def _log_mutation(self, delta: Dict[str, Any], timestamp: str) -> None:
189         """Append mutation record to audit trail."""
190         log_path = Path("io/mutations.log")
191         log_path.parent.mkdir(parents=True, exist_ok=True)
192
193         with open(log_path, 'a') as f:
194             f.write(f"[{timestamp}] MUTATION_SUCCESS\n")
195             f.write(f"  Trigger: {self.tier.capitalize()}Tuning_{self.study_name}\n")
196             f.write(f"  Best_Objective: {self.best_objective:.6f}\n")
197             f.write(f"  Delta:\n")

```

```

193         for param, change in delta.items():
194             f.write(f"    - {param}: {change}\n")
195             f.write("\n")
196
197 class ConfigMutationError(Exception):
198     """Raised when config mutation violates invariant protection rules."""
199     pass
200
201 class IntegrityError(Exception):
202     """Raised when checkpoint integrity verification fails."""
203     pass

```

21.3 Non-Blocking I/O Execution Pattern

Meta-optimization I/O operations (checkpoint save/load, TOML export) are blocking by nature due to `fsync()` requirements. To prevent interference with live prediction and telemetry collection, these operations must execute in separate threads or processes.

```

1 import threading
2 import queue
3 from concurrent.futures import ThreadPoolExecutor
4
5 class AsyncMetaOptimizer:
6     """
7         Wrapper for BayesianMetaOptimizer with non-blocking I/O.
8         Checkpoints and config exports execute in background threads.
9     """
10
11     def __init__(self, optimizer: BayesianMetaOptimizer):
12         self.optimizer = optimizer
13         self._io_executor = ThreadPoolExecutor(max_workers=2, thread_name_prefix="meta_io")
14         self._checkpoint_queue = queue.Queue(maxsize=5)
15
16     def save_study_async(self, path: str) -> None:
17         """
18             Non-blocking checkpoint save.
19             Submits to thread pool and returns immediately.
20
21             Note:
22                 If checkpoint queue is full (5 pending), oldest is dropped (backpressure).
23
24             if self._checkpoint_queue.full():
25                 # Drop oldest pending checkpoint to prevent memory buildup
26                 try:
27                     self._checkpoint_queue.get_nowait()
28                 except queue.Empty:
29                     pass
30
31             future = self._io_executor.submit(self.optimizer.save_study, path)
32             self._checkpoint_queue.put(future)
33
34     def export_config_async(
35         self,
36         result: OptimizationResult,
37         path: str = "config.toml"
38     ) -> None:
39         """
40
41             Non-blocking config export.
42             Returns immediately, actual write happens in background.
43
44             WARNING:
45                 The config file mutation happens asynchronously.
46                 Do not restart the predictor until export completes.
47                 Use wait_for_io_completion() to block until done.
48
49             self._io_executor.submit(result.export_to_toml, path=path, backup=True)
50
51     def wait_for_io_completion(self, timeout_seconds: float = 60.0) -> bool:
52         """
53             Block until all pending I/O operations complete.
54
55             Args:

```

```

55     timeout_seconds: Maximum wait time
56
57     Returns:
58         True if all operations completed, False if timeout
59
60     Use Case:
61         Before restarting predictor after config mutation:
62         >>> async_opt.export_config_async(result)
63         >>> if async_opt.wait_for_io_completion(timeout_seconds=30):
64             >>>     predictor.reload_config() # Safe to reload
65         """
66     deadline = time.time() + timeout_seconds
67
68     while not self._checkpoint_queue.empty():
69         remaining = deadline - time.time()
70         if remaining <= 0:
71             return False
72
73     try:
74         future = self._checkpoint_queue.get(timeout=remaining)
75         future.result(timeout=remaining) # Wait for completion
76     except queue.Empty:
77         break
78     except Exception as e:
79         print(f"I/O operation failed: {e}")
80         return False
81
82     # Wait for executor to finish all tasks
83     self._io_executor.shutdown(wait=True, cancel_futures=False)
84     return True
85
86 # Example usage in production
87 if __name__ == "__main__":
88     # Setup optimizer
89     search_space = {
90         'cusum_k': SearchSpace('cusum_k', 'float', (0.3, 1.5)),
91         'dgm_width_size': SearchSpace('dgm_width_size', 'int', (32, 256)),
92         # ... more parameters
93     }
94
95     optimizer = BayesianMetaOptimizer(
96         search_space=search_space,
97         objective_fn=walk_forward_validation,
98         study_name="deep_tuning_2026_Q1",
99         max_iterations=500,
100        tier="deep"
101    )
102
103 # Wrap for non-blocking I/O
104 async_optimizer = AsyncMetaOptimizer(optimizer)
105
106 # Attempt resume from checkpoint
107 checkpoint_path = "io/snapshots/deep_tuning_2026_Q1_iter250.pkl"
108 if Path(checkpoint_path).exists():
109     optimizer.load_study(checkpoint_path) # Blocking load is OK (one-time startup)
110
111 # Run optimization with non-blocking checkpoints
112 for i in range(optimizer._iteration, optimizer.max_iterations):
113     candidate = optimizer._suggest_next_candidate()
114     objective = walk_forward_validation(candidate)
115     optimizer._report_trial(candidate, objective)
116
117     if (i + 1) % 10 == 0:
118         # Non-blocking checkpoint (doesn't interrupt optimization loop)
119         checkpoint = f"io/snapshots/deep_tuning_2026_Q1_iter{i+1}.pkl"
120         async_optimizer.save_study_async(checkpoint)
121
122 # Export results to config.toml (non-blocking)
123 result = OptimizationResult(
124     study_name="deep_tuning_2026_Q1",
125     tier="deep",
126     best_params=optimizer._best_params,
127     best_objective=optimizer._best_value,

```

```

128     total_iterations=optimizer._iteration,
129     early_stopped=False,
130     convergence_delta=0.0001,
131     timestamp_utc=time.strftime("%Y-%m-%dT%H:%M:%SZ", time.gmtime())
132   )
133   async_optimizer.export_config_async(result)
134
135   # Wait for all I/O to complete before exiting
136   if async_optimizer.wait_for_io_completion(timeout_seconds=60):
137     print("All I/O operations completed successfully")
138   else:
139     print("WARNING: Some I/O operations timed out")

```

21.4 Integration with Prediction Pipeline

Meta-optimization runs offline (batch mode) separate from live prediction. However, config mutations can occur during production operation. The predictor must detect config changes and reload safely.

```

1 class ConfigReloadablePredictor(UniversalPredictor):
2   """
3     UniversalPredictor with hot-reload capability for config mutations.
4   """
5
6   def __init__(self, config_path: str = "config.toml"):
7     self.config_path = config_path
8     self._config_mtime = os.path.getmtime(config_path)
9
10    config = self._load_config(config_path)
11    super().__init__(config)
12
13  def check_and_reload_config(self) -> bool:
14    """
15      Check if config.toml has been modified and reload if necessary.
16
17      Returns:
18          True if config was reloaded, False if unchanged
19
20      Note:
21          Reloading config triggers full state reinitialization.
22          Call this only during safe windows (e.g., market closed, low traffic).
23    """
24    current_mtime = os.path.getmtime(self.config_path)
25
26    if current_mtime > self._config_mtime:
27      print(f"Config file modified. Reloading from {self.config_path}")
28      new_config = self._load_config(self.config_path)
29
30      # Reinitialize with new config
31      self.config = new_config
32      self._state = self._initialize_state()
33      self._jit_update = jax.jit(self._core_update_step)
34
35      self._config_mtime = current_mtime
36      return True
37
38    return False

```

22 Dependency Pinning

Strict version pinning is mandatory. Any update must be tested for bit-exact parity and documented. Use exact versions in `requirements.txt` and `environment.yml`, never open ranges.