

# Input/Output Interface Specification - Universal Predictor System

Systems Architecture

February 21, 2026

## Contents

<b>1 Executive Summary</b>	<b>2</b>
<b>2 Configuration Vector (Hyper-Inputs)</b>	<b>2</b>
<b>3 Input Flow (Data Injection)</b>	<b>2</b>
3.1 1. Calibration Phase (Bootstrapping) . . . . .	2
3.2 2. Operational Phase (Online Stream) . . . . .	2
<b>4 Security Policies in the I/O Layer (Credentials)</b>	<b>3</b>
<b>5 System Outputs</b>	<b>3</b>
5.1 1. Control signal (Prediction signal) . . . . .	3
5.2 2. State telemetry . . . . .	4
<b>6 Abstract I/O Diagram</b>	<b>5</b>
6.1 Internal Process Cycle . . . . .	5
<b>7 Persistence (Snapshotting)</b>	<b>6</b>
7.1 Atomic and Verified Snapshotting Protocol . . . . .	6
<b>8 Telemetry and Deterministic Logging</b>	<b>6</b>
8.1 Non-Blocking Telemetry . . . . .	6
8.2 Parity Audit Hashes . . . . .	6
<b>9 Configuration Mutation Protocol (Autonomous Self-Calibration)</b>	<b>7</b>
9.1 Motivation and Threat Model . . . . .	7
9.2 Atomic TOML Update Algorithm . . . . .	7
9.3 Critical Implementation Details . . . . .	7
9.3.1 Fsync Requirement . . . . .	7
9.3.2 Concurrent Mutation Prevention . . . . .	7
9.3.3 Rollback Procedure . . . . .	9
9.4 Invariant Protection: Locked Configuration Subsections . . . . .	9
9.4.1 Locked Subsections . . . . .	9
9.4.2 Mutable Subsections (Optimizer-Accessible) . . . . .	10
9.5 Validation Schema . . . . .	10
9.6 Mutation Audit Trail . . . . .	10
9.7 Rate Limiting and Safety Guardrails . . . . .	11
9.8 Integration with Meta-Optimization Workflow . . . . .	11

# 1 Executive Summary

This document defines the abstract input/output (I/O) interface for the Universal Predictor System, independent of the concrete implementation (Python/JAX, C++, Rust, FPGA). It describes the configuration vectors needed to instantiate the system, the runtime data flow, and the structure of the output signals and telemetry.

## 2 Configuration Vector (Hyper-Inputs)

The system is initialized with a configuration vector  $\Lambda$  that defines module topology and sensitivity. These parameters are typically static during an operating session or tuned by an external meta-optimizer.

Parameter	Symbol	Functional Description
Entropic regularization	$\epsilon$	Mass transport smoothing in the JKO orchestrator (Sinkhorn).
Learning rate	$\tau$	Adaptation speed of weights $\rho$ under energy gradients.
Signature depth	$L$	Truncation order of the log-signature (Kernel D - topological).
WTMM memory	$N_{buf}$	Sliding window size for singularity estimation.
Besov cone	$C_{besov}$	Influence radius for wavelet maxima tracking.
Holder threshold	$H_{min}$	Critical regularity that triggers the circuit breaker.
CUSUM threshold	$h$	Accumulated deviation level that triggers weight reset.
CUSUM slack	$k$	Drift tolerance ("white noise") allowed without error accumulation.
Volatility memory	$\alpha$	EMA decay rate to estimate error variance.

Table 1: Hyperparameter vector  $\Lambda$

## 3 Input Flow (Data Injection)

### 3.1 1. Calibration Phase (Bootstrapping)

Initial state required before sequential operation.

**Input:** History  $\mathcal{H} = \{y_{-T}, \dots, y_0\}$

- **Structure:** Time series of vectors  $\mathbf{y} \in \mathbb{R}^d$  or scalars  $y \in \mathbb{R}$ .
- **Purpose:**
  - Initialize history-dependent kernels (e.g., Levy parameters).
  - Stabilize initial orchestrator weights  $\rho_0$ .
  - Fill the singularity buffer for the WTMM module.

### 3.2 2. Operational Phase (Online Stream)

Step-by-step update cycle at time  $t$ .

**Input at step  $t$ :** Tuple  $(y_t, y_{target}, \tau_{epoch})$

- **Timestamp ( $\tau_{epoch}$ ):** Absolute timestamp (Unix nanoseconds). Required for synchronization and latency checks in the staleness policy.
- **Current observation ( $y_t$ ):**
  - New data point available at  $t$ .
  - Used to feed kernels ( $K_A, K_B, K_C, K_D$ ) and generate predictions for  $t + 1$ .
  - **Domain error handling:** If  $|y_t| > 20\sigma$  (relative to historical normalization), the point is classified as a catastrophic outlier. The system must discard the input, keep the inertial state, and emit a critical validation alert to protect kernels from numerical divergence.
  - **Frozen signal detection:** If the stream injects the exact same value for  $N_{freeze} \geq 5$  consecutive steps, the system must:

1. Compute the variance of the last  $N_{freeze}$  values:  $\text{Var}([y_{t-4}, y_{t-3}, y_{t-2}, y_{t-1}, y_t]) = 0$
  2. Identify this as sensor failure or data source corruption
  3. Emit `FrozenSignalAlarmEvent` with the event timestamp
  4. **Mathematical impact:** The Holder exponent in Branch D requires variability:  $H_t = \lim_{s \rightarrow 0} \frac{\log |\gamma(t+s) - \gamma(t)|}{\log s}$ . With a frozen signal, the numerator is zero, causing singularities or indeterminate values. This invalidates the multifractal spectrum. The system must:
    - \* Freeze the topological branch (Kernel D) at the last valid value
    - \* NOT update orchestrator weights (keep inertia)
    - \* Activate degraded inference mode
    - \* Continue predictions using Branches A, B, C only
  5. **Recovery:** Once variance  $> 0.1 \times \text{Var}_{historical}$  is detected for 2 consecutive steps, release the Kernel D lock and resume normal operation.
- **Inference grid (anti-aliasing input):**
    - **Sampling frequency vs scales ( $N_{buf}$ ):** To guarantee WTMM stability (Kernel D singularities), the data injection frequency must maintain sufficient density relative to the finest wavelet scales.
    - *Restriction:* A **minimum injection frequency** (Nyquist soft limit) is enforced based on  $C_{besov}$ . If event density falls below this threshold, the multifractal spectrum collapses and the system must freeze the topological branch update.
  - **Validation target ( $y_{target}$ ):**
    - The "real" value corresponding to the prediction generated at the previous step ( $t - 1$ ).
    - Typically  $y_{target} \equiv y_t$  for causal one-step-ahead prediction.
    - Used to compute error  $e_t = y_{target} - \hat{y}_{t|t-1}$  and the energy gradient  $\nabla E$  for JKO transport.
  - **Staleness policy:**
    - **Time-to-live (TTL):** Parameter  $\Delta_{max}$ .
    - **Violation behavior:** If the delay of  $y_{target}$  exceeds  $\Delta_{max}$ , the JKO update is canceled.
    - **Integrity signal:** The system must emit a persistent *degraded inference* flag ("stale weights"). This alerts the executor that, although the prediction  $\hat{y}$  is still produced, the weights  $\rho$  are stale and risk is no longer optimized geometrically.

## 4 Security Policies in the I/O Layer (Credentials)

The Stochastic Predictor design requires high-frequency ingestion against external market infrastructure (e.g., institutional WebSockets, brokers, or REST APIs). Access to these systems introduces critical vulnerability vectors.

### Secure environment injection

Hardcoding tokens, API keys, database secrets, or connection credentials in any future source module (e.g., `iostreams.py`) is strictly forbidden. Every implementation **must** apply an environment injection pattern, reading credentials at runtime from OS variables or local `.env` files.

### Version control exclusion

The resulting implementation repository must include explicit exclusion rules in version control (e.g., enforce `*.env` in `.gitignore`) to ensure no secret is exposed.

## 5 System Outputs

### 5.1 1. Control signal (Prediction signal)

Primary output for decision-making.

**Output:**  $\hat{y}_{t+1}$

- **Description:** Estimate of the expected process value for the next instant.
- **Inference grid (output quantization):**
  - Output  $\hat{y}_{t+1}$  is delivered in normalized space (Z-score) consistent with input  $y_t$ .
  - The actor/executor applies inverse normalization using rolling-window statistics if an absolute price is required.
- **Composition:** Convex combination of base kernels:  $\hat{y}_{t+1} = \sum_{i \in \{A, B, C, D\}} \rho_i^{(t)} \cdot K_i(y_t)$ .

## 5.2 2. State telemetry

Latent variables that describe system "health" and market regime.

- **Risk state ( $\mathbb{S}_{risk}$ ):**
  - **Local Holder exponent ( $H_t$ ):** Pointwise regularity measure.  $H_t < 0.5$  indicates antipersistence/noise;  $H_t < H_{min}$  indicates imminent crash/shock.
  - **Empirical kurtosis ( $\kappa_t$ ):** Fourth standardized moment of prediction residuals over a rolling window:
$$\kappa_t = \frac{E[(e_t - \mu_e)^4]}{(\sigma_e)^4}$$
where  $e_t = y_{target} - \hat{y}_{t|t-1}$  are prediction residuals.
    - \* **Purpose:** Validate the adaptive CUSUM threshold. Values  $\kappa_t > 3$  indicate leptokurtic distributions (heavy tails), activating the logarithmic adjustment  $h_t = k \cdot \sigma \cdot (1 + \ln(\kappa_t/3))$ .
    - \* **Interpretation:**  $\kappa_t \approx 3$  (Gaussian),  $\kappa_t \in [5, 10]$  (standard financial volatility regime),  $\kappa_t > 15$  (crisis regime with frequent extreme events).
    - \* **Alert:** If  $\kappa_t > 20$  persistently, emit a warning of potential residual model failure or undetected systematic outliers.
  - **DGM predictor entropy ( $H_{DGM}$ ):** Differential entropy of the neural value function  $V_\theta(t, x)$  across the spatial domain:

$$H_{DGM} = - \int p_V(v) \log p_V(v) dv$$

where  $p_V(v)$  is the empirical density of  $V_\theta$  values evaluated on a domain grid.

- \* **Purpose:** Monitor the health of Branch B (HJB solution via Deep Galerkin Method) and detect mode collapse when the network predicts a constant or degenerate solution.
- \* **Collapse threshold:** Compare against the terminal condition entropy:  $H_{DGM} \geq \gamma \cdot H[g]$  with  $\gamma \in [0.5, 1.0]$ . If the inequality is violated persistently (more than 10 consecutive steps), emit ModeDegradationAlert.
- \* **Corrective action:** Reduce the weight of Branch B in the orchestrator ( $\rho_B \rightarrow 0$ ) and prioritize alternative branches until the DGM network is retrained or reinitialized.
- \* **Note:** This indicator is only relevant if Branch B is active ( $\rho_B > 0.05$ ). For systems that do not use DGM, this field may be omitted or reported as NaN.
- **CUSUM statistic ( $G^+$ ):** Accumulated structural mismatch level.
- **Distance to collapse ( $h_t - G^+$ ):** Safety margin before a forced model reset. Note:  $h_t$  is now dynamic and depends on  $\sigma_t$  and  $\kappa_t$ .
- **Residual free energy ( $\mathcal{F}$ ):** Instant value of the JKO functional. It monitors whether the model is "stuck" in a stable local minimum or whether entropic regularization  $\epsilon$  is too high, over-smoothing mass transport and diluting predictive power.
- **Orchestrator state ( $\rho$ ):**
  - **Weight vector:**  $[\rho_A, \rho_B, \rho_C, \rho_D]$  such that  $\sum \rho = 1$ .
  - Indicates which "physics" currently dominates the market (jumps vs diffusion vs memory vs topology).

- **Stochastic health-check:**

- **Sinkhorn convergence (bool):** Indicates whether the mass transport algorithm converged within the maximum iteration count.
- *True*: Exact Wasserstein distance. *False*: Sub-optimal approximation (numerical precision alert).

- **Operational flags (mode and circuit breakers):**

- **Base operation mode:**

- \* *Standard (MSE)*: Normal operation under local Gaussian assumptions.
- \* *Robust (Huber)*: Defensive operation triggered by singularities ( $H_t < H_{min}$ ) or high volatility.

- **Degraded inference mode:** Critical boolean flag for temporal quality monitoring:

- \* **Activation condition:** It activates when the time-to-live (TTL) of  $y_{target}$  exceeds the maximum threshold:

$$\text{TTL}(y_{target}) = t_{current} - t_{signal} > \Delta_{max}$$

- \* **Operational implications:**

1. JKO transport update is suspended immediately
2. Weights  $\rho$  freeze at their last valid value (inertial mode)
3. Predictions  $\hat{y}_{t+1}$  continue, but are sub-optimal because they do not reflect the true market state
4. Risk is no longer optimized geometrically

- \* **Executor signaling:** This flag must explicitly warn that:

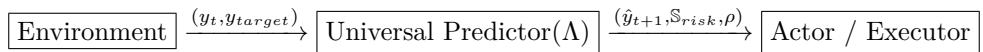
- Current predictions have *degraded confidence*
- Weights are stale
- Exposure should be reduced or a conservative mode should be used
- The system operates in "survival mode" until fresh data flow is restored

- \* **Recovery:** The flag is automatically cleared when a fresh signal arrives with  $\text{TTL}(y_{target}) < 0.8 \cdot \Delta_{max}$  (hysteresis threshold to avoid oscillation). At that moment, JKO transport resumes and `NormalOperationRestoredEvent` is emitted.

- **Emergency mode (singularity fallback):** Flag indicating whether emergency mode was triggered by critical singularity ( $H_t < H_{min}$ ), forcing  $w_D \rightarrow 1.0$  and switching to the Huber metric.

- **Regime change detected:** Flag indicating whether CUSUM detected a regime change at the last step, with entropy reset to a uniform distribution.

## 6 Abstract I/O Diagram



### 6.1 Internal Process Cycle

1. **Ingestion:** Receive  $y_t$ . Update local history.
2. **Singularity analysis:** Compute  $H_t$  using WTMM on the recent window.
3. **Quality control (CUSUM):**
  - Compute error  $e_t$  using  $y_{target}$  and stored prediction  $\hat{y}_{t|t-1}$ .
  - Update drift accumulator  $G^+$ .
  - If  $G^+ > h$  or  $H_t < H_{min} \rightarrow$  emit reset/alert signal.
4. **Transport (JKO):**

- Compute energy gradient  $\nabla E$  based on  $e_t$ .
- Transport probability mass  $\rho_{t-1} \rightarrow \rho_t$  (Sinkhorn).

### 5. Projection:

- Execute kernels  $K_i(y_t)$  to obtain components.
- Aggregate components using new weights  $\rho_t$  to obtain  $\hat{y}_{t+1}$ .

## 7 Persistence (Snapshotting)

To ensure operational continuity, the system must be able to serialize its full internal state  $\Sigma_t$  at any time  $t$ .

$$\Sigma_t = \{\rho_t, G_t^+, \sigma_{ema}^2, \kappa_t, H_{DGM}, \text{Flags}, \text{WTMMBuffer}, \text{KernelsState}\}$$

where:

- $\kappa_t$ : Rolling empirical kurtosis of prediction errors (window size = 252).
- $H_{DGM}$ : Differential entropy of the DGM predictor (for mode collapse detection).
- **Flags**: Boolean flags including `DegradedInferenceMode`, `EmergencyMode`, and `RegimeChangeDetected`.

The `KernelsState` structure must be segmented into independent sub-blocks (K-blocks) to allow modular or partial updates:

$$\text{KernelsState} = \{S_A(\text{Levy}), S_B(\text{PDE}), S_C(\text{Memory}), S_D(\text{Topology})\}$$

The restore operation  $Load(\Sigma_t)$  must allow the flow to resume at  $t + 1$  without recalibration over history  $\mathcal{H}$ . Correct restoration of  $\kappa_t$  and  $H_{DGM}$  is critical to preserve anomaly detection and mode collapse sensitivity after a restart.

### 7.1 Atomic and Verified Snapshotting Protocol

Binary serialization formats (e.g., Protocol Buffers, MessagePack) are required instead of text (JSON/XML) to minimize I/O latency for critical hot-start operations.

- **Mandatory integrity checksum:** Because dense binary formats are used, a single-bit error in kernel matrices or the WTMM buffer could trigger system collapse or undefined behavior. Therefore,  $\Sigma_t$  must include a robust validation hash (e.g., SHA-256 or CRC32c) at the end of the block.
- **Pre-injection validation:** The restore routine  $Load(\Sigma_t)$  must recompute and verify this hash *before* injecting the state into operational memory. If validation fails, the snapshot must be discarded and the system must restart in cold-start mode (history reload).

## 8 Telemetry and Deterministic Logging

### 8.1 Non-Blocking Telemetry

Telemetry emission must be decoupled from JAX/XLA execution. Compute threads must never block on I/O. Telemetry buffers are enqueued in a non-blocking structure and consumed by a separate process or thread.

### 8.2 Parity Audit Hashes

For hardware parity validation, log SHA-256 hashes of the weight vector  $\rho$  and the OT cost at configurable intervals. Hash inputs must use canonical float64 serialization to guarantee deterministic parity across CPU and GPU.

- Hash interval is configuration-driven.
- Logs are append-only and immutable.
- No raw signals or secrets are emitted in telemetry.

## 9 Configuration Mutation Protocol (Autonomous Self-Calibration)

For the Universal Stochastic Predictor to achieve Level 4 autonomy (unsupervised operation), it must be capable of **self-modifying its configuration vector**  $\Lambda$  in response to meta-optimization results without human intervention. This requires a secure, atomic, and auditable protocol for mutating the `config.toml` file on disk while preserving system integrity.

### 9.1 Motivation and Threat Model

The configuration vector  $\Lambda$  governs critical numerical and algorithmic parameters. Unsafe mutation can lead to catastrophic failure modes:

1. **Invalid Parameter Ranges:** Optimizer writes `cusum_k = -1.5` (negative threshold), causing CUSUM detector to enter undefined behavior.
2. **Circular Dependencies:** Optimizer modifies `telemetry_hash_interval` or `snapshot_path`, corrupting its own audit trail.
3. **Race Conditions:** Concurrent write by optimizer and manual operator edit causes partial file corruption.
4. **Loss of Rollback Capability:** Overwriting `config.toml` without backup prevents recovery from pathological configurations.

The mutation protocol must guarantee:

- **Atomicity:** Configuration updates are all-or-nothing (no partial writes visible to readers).
- **Durability:** Successful writes survive process crashes and power failures.
- **Auditability:** All mutations are logged with timestamp, triggering event, and parameter delta.
- **Rollback:** Previous configurations are preserved for manual recovery.
- **Invariant Protection:** Critical subsections are immutable to prevent self-corruption.

### 9.2 Atomic TOML Update Algorithm

The configuration mutation follows a strict POSIX-compliant protocol to ensure filesystem-level atomicity:

### 9.3 Critical Implementation Details

#### 9.3.1 Fsync Requirement

The `fsync()` system call (POSIX) or `FlushFileBuffers()` (Windows) is **mandatory** after writing the temporary file. Without it, the write may reside in kernel page cache and be lost during power failure or system crash before the atomic replacement occurs.

On Linux with ext4/xfs filesystems:

- `fsync(fd)` flushes file data and metadata to disk.
- `os.replace()` is implemented via `renameat2()` with `RENAME_NOREPLACE` flag (kernel 3.15+), guaranteeing atomicity even on NFS.

#### 9.3.2 Concurrent Mutation Prevention

The use of `O_EXCL` flag when creating `config.toml.tmp` ensures that if two processes attempt simultaneous mutations, only one succeeds. The losing process receives `EEXIST` error and must retry or abort.

For production deployments, an additional file-based lock can be used:

```
# Acquire exclusive lock before mutation
with open("config.toml.lock", "x") as lock:
    perform_atomic_mutation()
```

---

**Algorithm 1** Atomic Configuration Mutation

---

1: **Input:** New configuration dictionary  $\Lambda_{\text{new}}$ , current config path  $P_{\text{cfg}} = \text{config.toml}$   
2: **Output:** Success/Failure status, error diagnostics  
3:  
4: **Phase 1: Validation**  
5:  $\Lambda_{\text{current}} \leftarrow \text{Load}(P_{\text{cfg}})$  ▷ Parse existing TOML  
6:  $\Lambda_{\text{merged}} \leftarrow \text{Merge}(\Lambda_{\text{current}}, \Lambda_{\text{new}})$  ▷ Apply updates  
7: **Validate**( $\Lambda_{\text{merged}}$ ) ▷ Check ranges, types, constraints  
8: **if** validation fails **then**  
9:     **Error:** "Invalid parameter values, aborting mutation"  
10:    **return** FAILURE  
11: **end if**  
12:  
13: **Phase 2: Immutable Backup**  
14:  $P_{\text{bak}} \leftarrow P_{\text{cfg}} + .\text{bak}$   
15:  $P_{\text{timestamp}} \leftarrow P_{\text{cfg}} + .\text{bak.} + \text{ISO8601\_timestamp}$   
16: **Copy**( $P_{\text{cfg}}$ ,  $P_{\text{timestamp}}$ ) ▷ Timestamped archive  
17: **Copy**( $P_{\text{cfg}}$ ,  $P_{\text{bak}}$ ) ▷ Latest backup (overwrite)  
18:  
19: **Phase 3: Atomic Write via Temporary File**  
20:  $P_{\text{tmp}} \leftarrow P_{\text{cfg}} + .\text{tmp}$   
21:  $\text{FD}_{\text{tmp}} \leftarrow \text{Open}(P_{\text{tmp}}, \text{O\_WRONLY} | \text{O_CREAT} | \text{O_EXCL})$   
22: **if** open fails (file exists) **then**  
23:     **Error:** "Concurrent mutation detected, aborting"  
24:    **return** FAILURE  
25: **end if**  
26: **Write**( $\text{FD}_{\text{tmp}}$ ,  $\text{Serialize}(\Lambda_{\text{merged}})$ ) ▷ TOML format  
27: **Fsync**( $\text{FD}_{\text{tmp}}$ ) ▷ Force kernel buffer flush to disk  
28: **Close**( $\text{FD}_{\text{tmp}}$ )  
29:  
30: **Phase 4: Atomic Replacement (POSIX os.replace)**  
31: **Replace**( $P_{\text{tmp}}$ ,  $P_{\text{cfg}}$ ) ▷ Atomic inode swap  
32: ▷ On POSIX: `os.replace()` guarantees atomicity  
33: ▷ On Windows: `ReplaceFileW()` with backup flag  
34:  
35: **Phase 5: Audit Logging**  
36:  $\Delta \leftarrow \text{Diff}(\Lambda_{\text{current}}, \Lambda_{\text{merged}})$   
37: **Log**(`io/mutations.log`, {timestamp, trigger,  $\Delta$ })  
38: **return** SUCCESS

---

### 9.3.3 Rollback Procedure

If the system exhibits pathological behavior after a configuration mutation (e.g., NaN detection, CUSUM alarm storm), the operator can manually rollback:

```
# Restore from backup
cp config.toml.bak config.toml

# Or restore from specific timestamp
cp config.toml.bak.2026-02-19T14:32:05Z config.toml
```

The system must be restarted after rollback to reload the configuration.

## 9.4 Invariant Protection: Locked Configuration Subsections

Not all configuration parameters are safe for autonomous mutation. The following subsections of `config.toml` are **strictly immutable** and must be excluded from the optimizer's search space:

### 9.4.1 Locked Subsections

#### 1. [io] Section: I/O paths and serialization parameters

- `snapshot_path`: Changing this path during runtime would orphan existing snapshots, breaking the ability to resume from checkpoints.
- `telemetry_buffer_maxlen`: Modifying ring buffer size requires memory reallocation, violating JIT compilation assumptions.
- `credentials_vault_path`: Altering credential paths could expose secrets or cause authentication failures.

**Justification:** These parameters define the system's I/O contract with external infrastructure. Mutation breaks persistence guarantees.

#### 2. [security] Section: Audit and validation parameters

- `telemetry_hash_interval_steps`: Changing the hash interval would invalidate hardware parity validation protocols.
- `snapshot_integrity_hash_algorithm`: Switching from SHA-256 to CRC32 mid-session would cause all existing snapshots to fail validation.
- `allowed_mutation_rate_per_hour`: Maximum number of config mutations per hour (prevents optimizer runaway).

**Justification:** The optimizer must not modify its own audit trail or disable its safety constraints (analogous to Asimov's Zeroth Law for AI systems).

#### 3. [core] Section (Partial Lock): Fundamental numerical precision

- `float_precision`: Locked to 64. Switching to 32-bit would break Malliavin calculus and signature immutability.
- `jax_platform`: Locked (e.g., "cpu" or "gpu"). Changing platform mid-session invalidates XLA compilation cache.

**Justification:** These parameters define the computational substrate. Mutation requires full system recompilation.

#### 4. [meta\_optimization] Section (Partial Lock):

- `max_deep_tuning_iterations`: Locked to prevent infinite optimization loops.
- `checkpoint_path`: Locked to preserve resumability of Deep Tuning.
- `mutation_protocol_version`: Locked (semantic versioning of this protocol).

**Justification:** Prevents the optimizer from disabling its own termination criteria or corrupting its checkpoint state.

#### 9.4.2 Mutable Subsections (Optimizer-Accessible)

The following subsections are **safe for autonomous mutation** and constitute the search space for Fast Tuning and Deep Tuning:

- **[sensitivity]**: CUSUM thresholds, grace periods, EMA smoothing factors, entropy windows, learning rates.
- **[kernels]**: DGM architecture (width, depth, activation), SDE solver thresholds (stiffness\_low, stiffness\_high), WTMM parameters (num\_scales, sigma, threshold).
- **[orchestrator]**: Weight decay rates, mode collapse thresholds, frozen signal recovery ratios.
- **[numerical]**: Sinkhorn epsilon, max iterations, signature depth (within safe range [3, 5]).

## 9.5 Validation Schema

Before any mutation is committed, the merged configuration  $\Lambda_{\text{merged}}$  must pass validation against a strict schema:

```
# Example validation rules (Python-like pseudocode)
schema = {
    "cusum_k": {"type": float, "range": [0.3, 1.5]},
    "dgm_width_size": {"type": int, "range": [32, 256],
                       "constraint": "must be power of 2"},
    "stiffness_low": {"type": float, "range": [50, 500],
                      "constraint": "must be < stiffness_high"},
    "float_precision": {"type": int, "locked": True, "value": 64},
    "snapshot_path": {"type": str, "locked": True},
    ...
}

def validate(config, schema):
    for param, rules in schema.items():
        if rules.get("locked") and config[param] != rules["value"]:
            raise ConfigMutationError(f"{param} is immutable")
        if not (rules["range"][0] <= config[param] <= rules["range"][1]):
            raise ConfigMutationError(f"{param} out of safe range")
        # Additional constraint checks...
    return True
```

## 9.6 Mutation Audit Trail

All configuration mutations are logged to `io/mutations.log` in append-only mode:

```
[2026-02-19T14:32:05.123456Z] MUTATION_START
Trigger: DeepTuning_Iteration_127
Best_Objective: 0.0234 (MAPE)
Delta:
- cusum_k: 0.5 -> 0.72 (+0.22)
- dgm_width_size: 128 -> 256 (doubled)
- stiffness_low: 100 -> 143 (+43)
Validation: PASSED
Backup: config.toml.bak.2026-02-19T14:32:05Z
Status: SUCCESS
```

```
[2026-02-19T15:45:12.987654Z] MUTATION_REJECTED
Trigger: FastTuning_Iteration_42
Delta:
- float_precision: 64 -> 32 (LOCKED)
Validation: FAILED (immutable parameter)
Status: ABORTED
```

This log enables forensic analysis of configuration evolution and debugging of optimizer behavior.

## 9.7 Rate Limiting and Safety Guardrails

To prevent optimizer pathologies (e.g., thrashing between configurations), the mutation protocol enforces:

- **Maximum Mutation Rate:** No more than  $N_{\max} = 10$  mutations per hour.
- **Minimum Stability Period:** New configuration must run for at least  $T_{\text{stable}} = 1000$  prediction steps before re-mutation allowed.
- **Delta Magnitude Limit:** Parameter changes must satisfy:

$$\left| \frac{\theta_{\text{new}} - \theta_{\text{old}}}{\theta_{\text{old}}} \right| \leq 0.5 \quad (\text{50\% max relative change})$$

Exception: Discrete parameters (e.g., `dgm_depth`) can change by  $\pm 1$  level.

- **Rollback on Degradation:** If prediction RMSE increases by  $> 30\%$  within 500 steps of mutation, automatically rollback to `config.toml.bak`.

## 9.8 Integration with Meta-Optimization Workflow

The mutation protocol integrates with the tiered optimization framework (Deep Tuning + Fast Tuning) defined in `Stochastic_Predictor_Implementation.tex`:

1. **Deep Tuning Phase:** After 500 iterations, the best configuration  $\theta_{\text{deep}}^*$  is committed via atomic mutation to `[structural]` subsections.
2. **Fast Tuning Phase:** After 50 iterations, the best configuration  $\theta_{\text{fast}}^*$  is committed to `[sensitivity]` subsections.
3. **Live Adaptation:** During production operation, if a regime shift is detected (CUSUM alarm + entropy spike), trigger Fast Tuning with current config as initialization. Upon completion, mutate `[sensitivity]` parameters.

This creates a closed-loop autonomous system:

Deployment → Monitor Performance → Detect Degradation → Trigger Optimization → Mutate Config → Reload System

With the Configuration Mutation Protocol, the Universal Stochastic Predictor achieves **Level 4 Autonomy**: it can self-optimize and self-reconfigure without human intervention, while maintaining safety guarantees through locked invariants, atomic updates, and auditable mutation trails.