

# **Validation and Test Protocol for the Universal Stochastic Predictor**

Adaptive Meta-Prediction Development Consortium

February 19, 2026

# Contents

<b>1 Unit Tests: Kernels and Fundamental Algorithms</b>	<b>2</b>
1.1 Entropy and Random Variable Generation . . . . .	2
1.1.1 CMS Algorithm . . . . .	2
1.1.2 Pseudo-Random Generator Integrity . . . . .	2
1.2 Singularity Analysis (SIA) . . . . .	2
1.2.1 Holder Exponent Detection . . . . .	2
1.2.2 Cone of Influence . . . . .	3
1.2.3 Soft Nyquist Limit Validation . . . . .	3
1.3 Algebraic Structures (Branch D) . . . . .	4
1.3.1 Chen Identity . . . . .	4
1.3.2 Scale Invariance . . . . .	4
<b>2 Integration Tests and Stochastic Convergence</b>	<b>5</b>
2.1 Stochastic Differential Equation (SDE) Solvers . . . . .	5
2.1.1 Convergence of Numerical Schemes . . . . .	5
2.1.2 Mixed CFL Condition . . . . .	5
2.2 Transport Optimization (Orchestrator) . . . . .	6
2.2.1 Sinkhorn Algorithm Stability . . . . .	6
2.2.2 Simplex Normalization . . . . .	6
2.3 HJB Solution via DGM (Branch B) . . . . .	6
2.3.1 Gradient Stability . . . . .	6
2.3.2 Viscosity Solution Validation . . . . .	7
2.3.3 Training Entropy Test (Mode Collapse) . . . . .	7
2.3.4 Mesh Refinement Convergence . . . . .	9
2.3.5 Simplified Output Variance Test (Mode Collapse) . . . . .	9
<b>3 Robustness Tests and Circuit Breakers</b>	<b>11</b>
3.1 Outlier and Regime Handling . . . . .	11
3.1.1 Outlier Injection . . . . .	11
3.1.2 CUSUM Trigger . . . . .	11
3.2 Emergency Mode (Robustness Postulate) . . . . .	12
3.2.1 Critical Singularity . . . . .	12
<b>4 I/O and Persistence Tests</b>	<b>13</b>
4.1 Snapshot Protocol . . . . .	13
4.1.1 Hot-Start . . . . .	13
4.1.2 Checksum Validation . . . . .	13
4.2 I/O Failure Recovery . . . . .	14
4.2.1 Write Interruption (Atomicity) . . . . .	14
4.2.2 Silent Disk Corruption . . . . .	14
4.2.3 Disk Space Exhaustion . . . . .	15

<b>5</b>	<b>Hardware Parity and Fidelity Tests (Cross-Platform)</b>	<b>17</b>
5.1	Bit-Consistency Tests . . . . .	17
5.2	Numerical Drift Test . . . . .	17
5.2.1	Branch D: Signatures on FPGA vs CPU . . . . .	17
5.2.2	Deterministic Reproducibility . . . . .	18
5.3	Latency and Throughput Validation . . . . .	19
<b>6</b>	<b>Final Validation Protocol (Causality)</b>	<b>20</b>
6.1	Generalization . . . . .	20
6.2	Meta-Optimization Efficiency . . . . .	20
6.3	Temporal Integrity . . . . .	21
6.3.1	Lag Injection Test . . . . .	21
<b>7</b>	<b>Edge Cases and Operational Limits</b>	<b>23</b>
7.1	CUSUM: Adaptive Dynamic Threshold . . . . .	23
7.2	Orchestrator: Maximum Entropy Convergence . . . . .	24
7.3	Branch D: Time Reparametrization Invariance . . . . .	24
7.4	Edge Case Summary Table . . . . .	25
<b>8</b>	<b>Acceptance Criteria Summary</b>	<b>26</b>
8.1	Final Considerations . . . . .	26

# Chapter 1

# Unit Tests: Kernels and Fundamental Algorithms

These tests verify isolated implementation of critical algorithms without relying on global system state.

## 1.1 Entropy and Random Variable Generation

### 1.1.1 CMS Algorithm

**Test Case 1.1** (Validation of  $\alpha$ -Stable Distributions). *Validate that generating  $\alpha$ -stable variables via the Chambers-Mallows-Stuck method produces distributions with the desired parameters  $(\alpha, \beta, \gamma, \delta)$ .*

**Criterion 1.1.** *For a sample size  $N \geq 10^4$ , empirical moments must match the theoretical properties of the stable distribution within a 95% confidence interval.*

### 1.1.2 Pseudo-Random Generator Integrity

**Test Case 1.2** (Mersenne Twister/PCG64). *Verify the absence of serial correlations and the long period guarantees of the pseudo-random number generator.*

**Criterion 1.2.** *Apply standard statistical test batteries (TestU01, Diehard) and verify that no randomness tests fail. The generator period must be  $\geq 2^{127}$  for Mersenne Twister and  $\geq 2^{128}$  for PCG64.*

## 1.2 Singularity Analysis (SIA)

### 1.2.1 Holder Exponent Detection

**Test Case 1.3** (WTMM Validation). *Use synthetic signals with known Holder exponent  $H$  to validate that the WTMM (Wavelet Transform Modulus Maxima) algorithm recovers the singularity spectrum  $D(h)$  with error < 5%.*

**Criterion 1.3.** *Let  $f(t)$  be a synthetic signal with known  $H = H_0$ . The Holder exponent estimated by WTMM must satisfy:*

$$|\hat{H} - H_0| < 0.05 \cdot H_0$$

*where  $\hat{H}$  is the estimate from multiscale analysis.*

### 1.2.2 Cone of Influence

**Test Case 1.4** (Besov Influence Radius). Verify that maxima linking in the scale space respects the influence radius defined by  $C_{\text{besov}}$ .

**Criterion 1.4.** For two consecutive maxima at scales  $s_1 < s_2$ , their temporal positions  $(t_1, t_2)$  must satisfy:

$$|t_2 - t_1| \leq C_{\text{besov}} \cdot (s_2 - s_1)$$

where  $C_{\text{besov}}$  is the Besov-space constant associated with the chosen wavelet.

### 1.2.3 Soft Nyquist Limit Validation

The I/O specification defines a minimum data injection frequency to preserve multifractal analysis integrity. This test explicitly validates that operational limit.

**Test Case 1.5** (Multifractal Aliasing). Gradually reduce the input sampling frequency until the spectrum  $D(h)$  collapses, validating that the system detects the undersampling condition before irreversible degradation occurs.

**Criterion 1.5.** Consider a reference multifractal signal with known singularity spectrum  $D_0(h)$  and nominal sampling frequency  $f_0$ . Use the following protocol:

1. Reduce sampling frequency by decimation:  $f_k = f_0/2^k$  with  $k = 0, 1, 2, \dots$
  2. For each  $f_k$ , compute the estimated spectrum  $\hat{D}_k(h)$  via WTMM
  3. Compute relative error:
- $$\varepsilon_k = \frac{\|D_0(h) - \hat{D}_k(h)\|_{L^2}}{\|D_0(h)\|_{L^2}}$$
4. Monitor the dominant Holder exponent:  $\hat{H}_k = \arg \max_h D_k(h)$

The system must satisfy:

- **Preventive detection:** Emit `FreezingTopologicalBranchEvent` when  $\varepsilon_k > 0.05$  (error > 5%)
- **Acceptance criterion:** The signal must trigger before Holder error exceeds 10%:

$$\frac{|\hat{H}_k - H_0|}{H_0} < 0.10 \quad \text{at the moment of freezing}$$

- **Corrective action:** After the signal, freeze the topological branch weight at its last reliable value:  $w_C \rightarrow w_C^{\text{frozen}}$  with no dynamic updates

**Note 1.1.** The Nyquist limit in multifractal analysis is a soft boundary. Self-similarity allows some tolerance to undersampling, but below a threshold fine-scale structures collapse and  $D(h)$  degrades. This test ensures the system operates in a safe resolution regime.

**Note 1.2.** The critical frequency  $f_{\text{critical}}$  depends on:

1. The scale range  $[s_{\min}, s_{\max}]$  of the wavelet transform
2. The temporal support of the mother wavelet  $\psi(t)$
3. The expected minimum Holder regularity  $H_{\min}$

A practical heuristic:

$$f_{\min} \geq \frac{10}{s_{\min}} \cdot (1 + H_{\min}^{-1})$$

This ensures at least 10 points per minimum scale, adjusted by roughness.

## 1.3 Algebraic Structures (Branch D)

### 1.3.1 Chen Identity

**Test Case 1.6** (Signature Concatenation). *Validate that concatenating signatures of two path segments via the tensor product ( $\otimes$ ) equals the signature of the full path.*

**Criterion 1.6.** *Let  $\gamma_1 : [0, T_1] \rightarrow \mathbb{R}^d$  and  $\gamma_2 : [0, T_2] \rightarrow \mathbb{R}^d$  be two paths. Chen's identity states:*

$$\text{Sig}(\gamma_1 \star \gamma_2) = \text{Sig}(\gamma_1) \otimes \text{Sig}(\gamma_2)$$

*where  $\gamma_1 \star \gamma_2$  denotes concatenation. Numerical error must be  $< 10^{-6}$  in Euclidean norm.*

### 1.3.2 Scale Invariance

**Test Case 1.7** (Temporal Reparametrization). *Check that the level- $M$  truncated signature is invariant under strictly increasing time reparametrizations of the input path.*

**Criterion 1.7.** *For a strictly increasing reparametrization  $\phi : [0, 1] \rightarrow [0, 1]$  with  $\phi(0) = 0$  and  $\phi(1) = 1$ , we must have:*

$$\text{Sig}^{(M)}(\gamma) = \text{Sig}^{(M)}(\gamma \circ \phi)$$

*where  $\text{Sig}^{(M)}$  is the level- $M$  truncated signature.*

## Chapter 2

# Integration Tests and Stochastic Convergence

Validate that component interactions respect continuous probability laws and numerical stability conditions.

## 2.1 Stochastic Differential Equation (SDE) Solvers

### 2.1.1 Convergence of Numerical Schemes

**Test Case 2.1** (Euler-Maruyama vs Milstein). *For diffusion with non-constant volatility, verify that Milstein achieves strong convergence order 1.0 versus order 0.5 for Euler-Maruyama.*

**Criterion 2.1.** *Consider the SDE:*

$$dX_t = \mu(X_t, t) dt + \sigma(X_t, t) dW_t$$

*with non-constant  $\sigma(x, t)$ . For a sequence of time steps  $\Delta t_n = 2^{-n} \Delta t_0$ , the strong error must satisfy:*

$$\begin{aligned} E[|X_T - X_T^{EM}|^2] &= O(\Delta t^{0.5}) \quad (\text{Euler-Maruyama}) \\ E[|X_T - X_T^M|^2] &= O(\Delta t^{1.0}) \quad (\text{Milstein}) \end{aligned}$$

*where  $X_T^{EM}$  and  $X_T^M$  are the numerical approximations.*

### 2.1.2 Mixed CFL Condition

**Test Case 2.2** (Stability Violation). *Force a time step  $\Delta t$  that violates the Courant-Friedrichs-Lowy restriction and confirm numerical instability or NaNs as expected behavior to calibrate the safety monitor.*

**Criterion 2.2.** *For an explicit scheme, the CFL condition requires:*

$$\Delta t \leq \frac{C}{\sup_x |\mu'(x)| + \sigma^2(x)/2}$$

*Force  $\Delta t > 10C$  and verify divergence detection via:*

- *Nan or Inf emergence in simulated trajectories*
- *Instability alert emitted by the safety module*

## 2.2 Transport Optimization (Orchestrator)

### 2.2.1 Sinkhorn Algorithm Stability

**Test Case 2.3** (Log-Domain Convergence). *Evaluate Sinkhorn convergence in the log domain with decreasing regularization  $\varepsilon$ , ensuring no underflow down to  $\varepsilon \geq 10^{-4}$ .*

**Criterion 2.3.** *The log-domain Sinkhorn-Knopp algorithm must converge when:*

$$\varepsilon \geq 10^{-4}$$

*For  $\varepsilon < 10^{-4}$ , the system must detect underflow risk and emit a warning. Convergence is measured by:*

$$\|K \text{diag}(u) K^T \text{diag}(v) - \mu\|_1 < 10^{-6}$$

*where  $K_{ij} = \exp(-C_{ij}/\varepsilon)$  is the Gibbs kernel.*

### 2.2.2 Simplex Normalization

**Test Case 2.4** (Probabilistic Mass Conservation). *Confirm that after any JKO update the sum of kernel weights is strictly  $\sum_i \rho_i = 1.0$ .*

**Criterion 2.4.** *After each JKO iteration:*

$$\rho^{n+1} = \arg \min_{\rho \in \mathcal{P}(\mathcal{X})} \{W_2^2(\rho, \rho^n) + \tau \mathcal{F}[\rho]\}$$

*verify that  $\rho^{n+1}$  belongs to the probability simplex:*

$$\sum_{i=1}^N \rho_i^{n+1} = 1.0, \quad \rho_i^{n+1} \geq 0 \quad \forall i$$

*with numerical tolerance  $|\sum_i \rho_i^{n+1} - 1.0| < 10^{-10}$ .*

## 2.3 HJB Solution via DGM (Branch B)

Branch B solves the Hamilton-Jacobi-Bellman equation with Deep Galerkin Method (DGM). These tests validate convergence and stability.

### 2.3.1 Gradient Stability

**Test Case 2.5** (Gradient Explosion Under High Volatility). *Monitor gradient norms during PDE training to detect gradient explosions in high-volatility regimes.*

**Criterion 2.5.** *The HJB equation:*

$$\frac{\partial V}{\partial t} + \sup_{u \in \mathcal{U}} \{\mathcal{L}^u V(x, t) + f(x, u, t)\} = 0$$

*is approximated by a neural network  $V_\theta(x, t)$  trained with DGM. During training, monitor:*

1. *Loss gradient norm:*

$$\|\nabla_\theta \mathcal{L}_{DGM}(\theta)\|_2 = \left\| \frac{\partial}{\partial \theta} \mathbb{E}[|PDE[V_\theta]|^2 + |BC[V_\theta]|^2] \right\|_2$$

2. *Apply gradient clipping when  $\|\nabla_\theta \mathcal{L}\|_2 > C_{clip}$  with  $C_{clip} = 10.0$*

3. In high-volatility scenarios ( $\sigma(x, t) > 2\sigma_0$ ), the system must:

- Detect if  $\|\nabla_\theta \mathcal{L}\|_2$  exceeds  $C_{clip}$  for more than 5 consecutive iterations
- Emit `GradientInstabilityEvent`
- Adaptively reduce learning rate:  $\eta \rightarrow 0.5\eta$
- Verify stabilization within the next 20 epochs

**Note 2.1.** Gradient explosions in DGM often arise from interactions between second-order derivatives and diffusion coefficients. Early detection and adaptive clipping are essential for convergence in high-volatility regimes.

### 2.3.2 Viscosity Solution Validation

**Test Case 2.6** (Crandall-Lions Comparison Principle). Validate that the neural solution  $V_\theta(x, t)$  respects the comparison principle for viscosity solutions, ensuring uniqueness and consistency for non-linear PDEs.

**Criterion 2.6.** The comparison principle states: if  $u$  is a viscosity supersolution and  $v$  is a viscosity subsolution of the same HJB equation with  $u \geq v$  on the boundary, then  $u \geq v$  on the full domain.

Validation procedure:

1. Construct a reference solution  $V_{ref}(x, t)$  using a standard method (upwind finite differences, method of lines)
2. Train the DGM network to obtain  $V_\theta(x, t)$
3. Verify time monotonicity (finite horizon):

$$V_\theta(x, t_1) \geq V_\theta(x, t_2) \quad \forall t_1 < t_2, \forall x \in \mathcal{D}$$

(assuming non-negative costs)

4. Validate sub/supersolution constraints on a test grid  $\{(x_i, t_j)\}_{i,j}$ :

$$\begin{aligned} \text{Subsolution: } & PDE[V_\theta](x_i, t_j) \leq \epsilon_{tol} \\ \text{Supersolution: } & PDE[V_\theta](x_i, t_j) \geq -\epsilon_{tol} \end{aligned}$$

with  $\epsilon_{tol} = 10^{-3}$

5. Compare with reference solution:

$$\|V_\theta - V_{ref}\|_{L^\infty(\mathcal{D})} < \delta_{viscosity}$$

where  $\delta_{viscosity} = 0.05 \cdot \|V_{ref}\|_{L^\infty}$  (relative error < 5%)

**Note 2.2.** Crandall-Lions viscosity theory is the rigorous framework for HJB equations. DGM networks, being smooth functions, must approximate these generalized solutions. The comparison principle test prevents spurious oscillations and causality violations.

### 2.3.3 Training Entropy Test (Mode Collapse)

During neural PDE training, the model can collapse to trivial constant solutions. This test detects and prevents that behavior.

**Test Case 2.7** (Mode Collapse Detection). Verify that the neural network does not collapse to a trivial solution that satisfies the PDE but fails to capture structure.

**Criterion 2.7.** *Mode collapse occurs when the network learns a minimum-variance solution that minimizes PDE loss without solving the boundary-value problem.*

*Detection protocol:*

1. Compute spatial variance at time  $t < T$ :

$$\text{Var}_x[V_\theta(x, t)] = \mathbb{E}_x[(V_\theta(x, t) - \bar{V}_t)^2]$$

where  $\bar{V}_t = \mathbb{E}_x[V_\theta(x, t)]$

2. Compute variance of the terminal condition:

$$\text{Var}[g(\xi)] = \mathbb{E}_\xi[(g(\xi) - \bar{g})^2]$$

3. Verify proportionality:

$$\kappa_{low} \leq \frac{\text{Var}_x[V_\theta(x, t)]}{\text{Var}[g(\xi)]} \leq \kappa_{high}$$

with typical thresholds  $\kappa_{low} = 0.3$  and  $\kappa_{high} = 1.2$

4. If the ratio falls below  $\kappa_{low}$ : detect mode collapse

5. Monitor differential entropy:

$$H[V_\theta] = - \int p(v) \log p(v) dv$$

A collapsed solution yields  $H[V_\theta] \rightarrow -\infty$  (delta distribution)

Acceptance criteria:

- Variance ratio in  $[\kappa_{low}, \kappa_{high}]$  for at least 90% of  $t \in [0, T]$
- Differential entropy satisfies  $H[V_\theta] > H_{min}$
- Spatial gradient norm does not collapse:

$$\mathbb{E}_x[\|\nabla_x V_\theta(x, t)\|_2] > \epsilon_{grad} > 0$$

**Note 2.3.** Mode collapse is common in:

1. Uniform or symmetric boundary conditions
2. Excessive learning rates leading to minimum-norm solutions
3. Under-parameterized networks
4. Biased weight initialization

Early variance monitoring allows reinitialization or hyperparameter tuning before collapse is permanent.

**Note 2.4.** From a control theory perspective, a collapsed solution  $V_\theta(x, t) \approx c$  yields a degenerate control policy  $u^*(x, t)$  that does not respond to state changes, which is operationally useless.

### 2.3.4 Mesh Refinement Convergence

**Test Case 2.8** (Consistency under Refinement). Verify that the DGM solution converges to the exact (or reference) solution as the density of collocation points increases.

**Criterion 2.8.** For nested meshes  $\mathcal{M}_1 \subset \mathcal{M}_2 \subset \mathcal{M}_3$  with densities  $N_1 < N_2 < N_3$ :

1. Train DGM on each mesh to loss convergence:  $\mathcal{L}_{DGM}^{(k)} < 10^{-4}$
2. Compute error on a fixed fine evaluation mesh:

$$e_k = \|V_{\theta_k} - V_{ref}\|_{L^2(\mathcal{D}_{eval})}$$

3. Verify monotone convergence:

$$e_1 > e_2 > e_3$$

4. Estimate empirical rate:

$$r = \frac{\log(e_1/e_2)}{\log(N_2/N_1)}$$

Expected  $r \geq 0.5$  (at least order  $N^{-0.5}$  convergence)

### 2.3.5 Simplified Output Variance Test (Mode Collapse)

This test complements entropy monitoring with a direct variance ratio check.

**Test Case 2.9** (Minimum Variance Threshold). Compare output variance of the neural solution  $V_{\theta}(x, t)$  against a reference solution  $V_{ref}(x, t)$  and verify that the network captures at least 10% of the true variability.

**Criterion 2.9.** Let  $V_{ref}(x, t)$  be a reference solution obtained with a standard method. For the trained DGM network  $V_{\theta}(x, t)$ :

1. Compute reference variance at time  $t \in [0, T]$ :

$$Var_{ref}(t) = \frac{1}{|\mathcal{X}|} \sum_{x_i \in \mathcal{X}} (V_{ref}(x_i, t) - \bar{V}_{ref}(t))^2$$

2. Compute neural variance:

$$Var_{\theta}(t) = \frac{1}{|\mathcal{X}|} \sum_{x_i \in \mathcal{X}} (V_{\theta}(x_i, t) - \bar{V}_{\theta}(t))^2$$

3. Variance ratio:

$$R_{var}(t) = \frac{Var_{\theta}(t)}{Var_{ref}(t)}$$

4. Failure criterion:

$$R_{var}(t) < 0.10 \quad \text{for any } t \in [0, 0.9T]$$

5. Acceptance criterion:

$$R_{var}(t) \geq 0.10 \quad \forall t \in [0, 0.9T]$$

and

$$\text{median}_{t \in [0, T]} R_{var}(t) \geq 0.50$$

**Note 2.5.** If  $R_{var}(t) < 0.10$  before convergence, interrupt training and adjust hyperparameters (learning rate, architecture, initialization). The 10% threshold is conservative; any solution below it is effectively constant. The  $t \in [0, 0.9T]$  restriction excludes the terminal region where variance naturally contracts.

**Note 2.6.** This criterion complements the entropy test. Practical interpretation:

- $R_{var} < 0.10$ : critical collapse (fail)
- $0.10 \leq R_{var} < 0.30$ : low-variance warning
- $R_{var} \geq 0.50$ : normal operation

# Chapter 3

# Robustness Tests and Circuit Breakers

These tests verify system protection against market anomalies or data failures.

## 3.1 Outlier and Regime Handling

### 3.1.1 Outlier Injection

**Test Case 3.1** (Extreme Outlier Values). *Inject values  $> 20\sigma$  in the input stream and verify that the system rejects the point, emits a validation alert, and preserves inertial weights.*

**Criterion 3.1.** *Given observations  $\{y_t\}$  with rolling mean  $\mu_t$  and standard deviation  $\sigma_t$ , inject:*

$$\tilde{y}_t = \mu_t + 20\sigma_t$$

*The system must:*

1. Detect  $|\tilde{y}_t - \mu_t| > \theta\sigma_t$  with  $\theta = 10$
2. Reject the observation and NOT update the meta-state  $\Xi_t$
3. Emit `OutlierDetectedEvent` with metadata for the rejected value
4. Keep weights  $\{w_i\}_{i=A}^D$  unchanged

### 3.1.2 CUSUM Trigger

**Test Case 3.2** (Structural Regime Change). *Simulate a regime change (structural drift) and validate that the change event is emitted exactly when  $G_t^+$  exceeds the dynamic threshold  $h$ .*

**Criterion 3.2.** *CUSUM detects mean shifts via:*

$$G_t^+ = \max(0, G_{t-1}^+ + (y_t - \mu_0) - k)$$

*where  $k$  is the slack and  $h$  is the alarm threshold. Verify:*

1.  $G_t^+ > h \Rightarrow$  emit `RegimeChangedEvent`
2. Reset  $G_t^+ = 0$  after detection
3. Detection delay is at most  $\tau$  observations after the true change point

## 3.2 Emergency Mode (Robustness Postulate)

### 3.2.1 Critical Singularity

**Test Case 3.3** (Extreme Roughness Regime). *Artificially reduce the Holder exponent below  $H_{min}$  and verify that the orchestrator forces  $w_D \rightarrow 1.0$  and switches the cost to the Huber metric.*

**Criterion 3.3.** Define a critical threshold  $H_{min}$  (typically  $H_{min} = 0.25$ ). When WTMM detects:

$$\hat{H}_t < H_{min}$$

The system must:

1. Activate emergency mode:  $w_A = w_B = w_C = 0$ ,  $w_D = 1.0$

2. Switch orchestrator cost from Wasserstein to Huber:

$$C(x, y) = \begin{cases} \frac{1}{2}|x - y|^2 & \text{if } |x - y| \leq \delta \\ \delta(|x - y| - \frac{\delta}{2}) & \text{if } |x - y| > \delta \end{cases}$$

3. Emit *CriticalSingularityEvent*

4. Maintain this state until  $\hat{H}_t > H_{min} + \epsilon_{hysteresis}$

# Chapter 4

## I/O and Persistence Tests

These tests guarantee operational continuity and latent state integrity.

### 4.1 Snapshot Protocol

#### 4.1.1 Hot-Start

**Test Case 4.1** (State Continuity). *Serialize the meta-state  $\Xi_t$ , restart the system, and load it. The first prediction after restart must match the prediction without interruption.*

**Criterion 4.1.** Define the full meta-state:

$$\Xi_t = \{\{w_i\}_{i=A}^D, \{\theta_i^*\}_{i=A}^D, \mathcal{H}_t, \text{Sig}_t, G_t^\pm, \mu_t, \sigma_t^2\}$$

Validation procedure:

1. At time  $t_0$ , serialize  $\Xi_{t_0}$  to a binary file
2. Generate prediction  $\hat{y}_{t_0+1}^{\text{original}}$
3. Restart the system (release memory)
4. Load  $\Xi_{t_0}$  from file
5. Generate prediction  $\hat{y}_{t_0+1}^{\text{restored}}$
6. Verify:  $|\hat{y}_{t_0+1}^{\text{original}} - \hat{y}_{t_0+1}^{\text{restored}}| < 10^{-12}$

#### 4.1.2 Checksum Validation

**Test Case 4.2** (Cryptographic Integrity). *Corrupt a single bit in the snapshot file and verify that the system rejects the load via SHA-256, forcing a cold start.*

**Criterion 4.2.** Each snapshot must include a SHA-256 hash. Load must:

1. Read the binary file
2. Compute  $H' = \text{SHA256}(\text{content})$
3. Compare to stored hash  $H$
4. If  $H' \neq H$ : reject load, emit `CorruptedSnapshotEvent`, initialize cold start
5. If  $H' = H$ : proceed with deserialization

Validate by flipping one random bit and verifying rejection.

## 4.2 I/O Failure Recovery

Persisting  $\Xi_t$  is critical for continuity. These tests validate robustness against write, read, and storage failures.

### 4.2.1 Write Interruption (Atomicity)

**Test Case 4.3** (Power Loss Simulation). *Simulate a sudden power loss during snapshot serialization and verify that partially written files are handled safely.*

**Criterion 4.3.** *Snapshotting must guarantee atomicity via write-then-rename:*

1. *Serialize  $\Xi_t$  to temporary file:  $\text{snapshot}_{\{\text{timestamp}\}}.\text{tmp}$*
2. *Compute  $H = \text{SHA256}(\Xi_t)$  and append to file*
3. *Perform  $fsync()$  to flush to disk*
4. *Atomically rename:  $\text{snapshot}_{\{\text{timestamp}\}}.\text{tmp} \rightarrow \text{snapshot}_{\{\text{timestamp}\}}.\text{bin}$*

*Validation:*

1. *Start snapshot at time  $t_0$*
2. *Interrupt at random progress  $p \in \{0.1, 0.3, 0.5, 0.7, 0.9\}$*
3. *Simulate abrupt termination ( $SIGKILL$  or  $I/O$  cut)*
4. *Restart system*
5. *The system must:*
  - *Detect absence of main  $.bin$  file*
  - *Detect presence of corrupted  $.tmp$*
  - *Ignore the temp file*
  - *Load the latest valid snapshot before  $t_0$*
  - *If none exist: execute cold start*
  - *Avoid infinite restart loops*

*Success criterion:*

$$\text{Recovery time} < T_{\text{recovery}} = 30 \text{ seconds}$$

*No functional degradation after recovery.*

**Note 4.1.** *Atomic snapshot writes avoid the torn write problem. Most modern file systems (ext4, XFS, NTFS, APFS) guarantee atomic rename, which this protocol relies on.*

### 4.2.2 Silent Disk Corruption

**Test Case 4.4** (Bit Rot and Storage Errors). *Detect and handle silent data corruption that can occur between snapshot write and read.*

**Criterion 4.4.** *The protocol must:*

1. *Store verification metadata with each snapshot:*
  - *SHA-256 hash*
  - *Creation timestamp*

- *Serialization format version*
- *Optional CRC32 pre-check*

2. *On load:*

- *Verify CRC32 if available*
- *Verify full SHA-256*
- *On failure, mark snapshot corrupt and search for previous valid snapshot*

3. *Retention policy:*

- *Keep last  $N = 5$  valid snapshots*
- *Allow recovery from  $t-k$  if  $t_0$  is corrupt*

*Validation:*

1. *Create valid snapshot at  $t_0$*
2. *Corrupt random bytes*
3. *Attempt load*
4. *Verify system:*
  - *Detects corruption*
  - *Emits CorruptedSnapshotEvent*
  - *Falls back to previous valid snapshot*
  - *Executes cold start if no valid snapshot exists*
  - *Avoids infinite retry loops*

#### 4.2.3 Disk Space Exhaustion

**Test Case 4.5** (Insufficient Capacity Handling). *Validate behavior when storage is full during snapshot write.*

**Criterion 4.5.** *During snapshotting, the system must:*

1. *Before writing, verify free space:*

$$\text{FreeSpace} \geq 2 \times \text{EstimatedSize}(\Xi_t)$$

*(factor 2 for temp + rename)*

2. *If insufficient space:*

- *Emit InsufficientStorageEvent*
- *Do not attempt write*
- *Keep last valid snapshot*
- *Continue operation in memory until space is available*

3. *If failure occurs mid-write:*

- *Catch I/O exception*
- *Delete corrupted temp file*
- *Emit SnapshotWriteFailedEvent*

- *Preserve last valid snapshot*

*Validation:*

1. *Create a limited storage volume*
2. *Fill it, leaving insufficient space*
3. *Attempt snapshot*
4. *Verify graceful handling without crash*

**Note 4.2.** *Disk exhaustion is common in production. Design priorities: (1) do not corrupt valid snapshots, (2) degrade gracefully in memory, (3) provide clear telemetry, (4) recover automatically when space frees.*

# Chapter 5

## Hardware Parity and Fidelity Tests (Cross-Platform)

Given heterogeneous targets (CPU, GPU, FPGA), add bit-consistency tests to ensure equivalent outputs across architectures.

### 5.1 Bit-Consistency Tests

**Test Case 5.1** (Multi-Architecture Equivalence). *Verify that critical algorithms produce consistent numerical results across hardware platforms (CPU, GPU, FPGA), within precision limits of each architecture.*

**Criterion 5.1.** *For each critical component (random generation, signatures, SDE integration), execute identical inputs on:*

1. *CPU with IEEE 754 floating-point (64-bit)*
2. *GPU with floating-point (32 or 64-bit)*
3. *FPGA with fixed-point arithmetic (configurable precision)*

*Relative difference must satisfy:*

$$\frac{\|x^{CPU} - x^{GPU}\|_2}{\|x^{CPU}\|_2} < \epsilon_{GPU}, \quad \frac{\|x^{CPU} - x^{FPGA}\|_2}{\|x^{CPU}\|_2} < \epsilon_{FPGA}$$

*where  $\epsilon_{GPU} = 10^{-6}$  for 32-bit arithmetic and  $\epsilon_{FPGA}$  is the quantization error of the lower-precision hardware.*

### 5.2 Numerical Drift Test

#### 5.2.1 Branch D: Signatures on FPGA vs CPU

**Test Case 5.2** (Fixed-Point Error Accumulation). *Compare Branch D signatures on FPGA (fixed-point) against CPU (64-bit floating-point).*

**Criterion 5.2** (Drift Acceptance Criteria). *For a path  $\gamma : [0, T] \rightarrow \mathbb{R}^d$ :*

$$Sig^{(M)}(\gamma) = \left( 1, \int_0^T d\gamma_{t_1}, \int_0^T \int_0^{t_1} d\gamma_{t_2} \otimes d\gamma_{t_1}, \dots \right)$$

*For 10,000 iterations:*

1. Compute CPU signature:  $Sig_{10000}^{CPU}$
2. Compute FPGA signature:  $Sig_{10000}^{FPGA}$
3. Compute accumulated divergence:

$$\Delta_{acum} = \|Sig_{10000}^{CPU} - Sig_{10000}^{FPGA}\|_\infty$$

4. Primary criterion:

$$\Delta_{acum} \leq N \cdot \epsilon_{quant}$$

where  $N = 10000$  and  $\epsilon_{quant} = 2^{-n+1}$  for  $n$ -bit fixed point

5. Secondary topological preservation:

- Norm preservation:

$$\frac{|\|Sig_{10000}^{FPGA}\|_2 - \|Sig_{10000}^{CPU}\|_2|}{\|Sig_{10000}^{CPU}\|_2} < \tau_{norm}$$

with  $\tau_{norm} = 0.01$

- Sign preservation:

$$sgn(s_i^{(k),CPU}) = sgn(s_i^{(k),FPGA}) \quad \forall i, k$$

- Angular distance:

$$\cos(\theta) = \frac{\langle Sig_{10000}^{CPU}, Sig_{10000}^{FPGA} \rangle}{\|Sig_{10000}^{CPU}\|_2 \cdot \|Sig_{10000}^{FPGA}\|_2}$$

Require  $\cos(\theta) > 0.9999$  (angular deviation  $< 0.81^\circ$ )

- Relative error per level:

$$\frac{\|Sig_{10000}^{CPU,(k)} - Sig_{10000}^{FPGA,(k)}\|_2}{\|Sig_{10000}^{CPU,(k)}\|_2} < \tau_k$$

where  $\tau_k = 0.05 \cdot k$

**Note 5.1.** Fixed-point drift is inevitable due to repeated truncation. The primary criterion bounds cumulative error by the worst-case independent error accumulation. For Branch D, qualitative properties must also be preserved: orientation, angular similarity, and low-level tensor accuracy.

**Note 5.2.** For FPGA with  $n = 32$  bits (Q16.16),  $\epsilon_{quant} = 2^{-15} \approx 3.05 \times 10^{-5}$ . After 10,000 iterations, the primary criterion allows  $\Delta_{acum} \leq 0.305$ , while topological criteria enforce tighter constraints:

- Relative norm error  $< 1\%$
- Angular deviation  $< 1^\circ$
- Level-1 error  $< 5\%$

### 5.2.2 Deterministic Reproducibility

**Test Case 5.3** (Controlled Seed Initialization). Guarantee that, given the same pseudo-random seed, all platforms (CPU, GPU, FPGA) produce the same state sequence.

**Criterion 5.3.** Set deterministic seed  $s_0$  and run 1,000 simulation steps on each platform. Verify:

$$\{X_t^{CPU}\}_{t=1}^{1000} = \{X_t^{GPU}\}_{t=1}^{1000} = \{X_t^{FPGA}\}_{t=1}^{1000}$$

Equality must be bit-for-bit for platforms with the same representation (CPU and GPU floating-point). For FPGA, compare after converting fixed-point to floating-point.

### 5.3 Latency and Throughput Validation

**Test Case 5.4** (Cross-Platform Performance Benchmark). *Measure execution time and throughput for each predictor branch on all architectures to identify bottlenecks and validate heterogeneous acceleration.*

**Criterion 5.4.** *For a batch of  $N = 1000$  predictions:*

$$T_{CPU} = \text{total time on } CPU$$

$$T_{GPU} = \text{total time on } GPU$$

$$T_{FPGA} = \text{total time on } FPGA$$

*Expected:*

- *GPU outperforms CPU for massively parallel operations:  $T_{GPU} < 0.3 \cdot T_{CPU}$*
- *FPGA outperforms CPU for deterministic low-latency operations:  $T_{FPGA} < 0.1 \cdot T_{CPU}$*

*Throughput:*

$$\text{Throughput} = \frac{N}{T} \quad [\text{predictions/second}]$$

# Chapter 6

## Final Validation Protocol (Causality)

This is the definitive predictive test before any deployment.

### 6.1 Generalization

**Test Case 6.1** (Rolling Walk-Forward). *Zero look-ahead bias. Training uses only data strictly prior to the test horizon.*

**Criterion 6.1.** *Split the dataset into rolling windows:*

$$\mathcal{D} = \{(t_1, y_1), \dots, (t_N, y_N)\}$$

For each test window  $\mathcal{T}_k = \{t_{n_k}, \dots, t_{n_k+w}\}$ :

1. Train only with  $\mathcal{D}_{train}^k = \{(t_i, y_i) : t_i < t_{n_k}\}$
2. Predict on  $\mathcal{T}_k$
3. Advance window:  $k \rightarrow k + 1$
4. Aggregate out-of-sample metrics (RMSE, MAE, Sharpe ratio)

The system must ensure that no future data is used to train the model at time  $t$ .

### 6.2 Meta-Optimization Efficiency

**Test Case 6.2** (Bayesian Optimization). *Iterative improvement of Expected Improvement on generalization error compared to random search.*

**Criterion 6.2.** *Use a Gaussian Process surrogate model for the hyperparameter space  $\Theta = \{\alpha, \beta, \gamma, \dots\}$ . For  $n$  optimization iterations:*

1. Random search:  $\theta_i \sim \text{Uniform}(\Theta)$
2. Bayesian optimization:  $\theta_i = \arg \max_{\theta} EI(\theta | \mathcal{D}_{1:i-1})$

Acceptance criterion:

$$\min_{i \leq n} \mathcal{L}(\theta_i^{BO}) < \min_{i \leq n} \mathcal{L}(\theta_i^{random})$$

where  $\mathcal{L}$  is validation loss. Improvement must be significant ( $p\text{-value} < 0.05$  in Mann-Whitney test).

## 6.3 Temporal Integrity

**Test Case 6.3** (Staleness Metric (TTL)). *Cancel JKO update if target signal delay  $y_{target}$  exceeds  $\Delta_{max}$ .*

**Criterion 6.3.** *Define time-to-live:*

$$TTL(y_t) = t_{current} - t$$

*If:*

$$TTL(y_t) > \Delta_{max}$$

*The system must:*

1. *Discard the signal*
2. *NOT perform JKO update*
3. *Emit StaleDataEvent*
4. *Log the rejected timestamp*

*Typical value:  $\Delta_{max} = 5$  seconds for high-frequency systems.*

### 6.3.1 Lag Injection Test

The staleness policy prevents operating with obsolete weights. This test validates behavior under extreme latency.

**Test Case 6.4** (Artificial Signal Delay). *Artificially delay  $y_{target}$  beyond  $\Delta_{max}$  and validate that degraded inference mode is activated and optimal transport is suspended.*

**Criterion 6.4.** *Validation procedure:*

1. *Configure system with  $\Delta_{max} = 5$  seconds*
2. *Generate real-time data stream with correct timestamps*
3. *Inject delayed signal  $\tilde{y}_t$ :*

$$TTL(\tilde{y}_t) = t_{current} - t = \Delta_{max} + \delta$$

*where  $\delta > 0$  (typically  $\delta = 1$  second)*

4. *Verify the sequence:*
  - *Detect  $TTL(\tilde{y}_t) > \Delta_{max}$*
  - *Activate `DegradedInferenceMode = True`*
  - *Suspend JKO transport:*

$$\text{JKO\_update}(\rho^{n+1}) \rightarrow \text{SUSPENDED}$$

- *Freeze orchestrator weights at last valid value:  $\{w_i\}_{frozen}$*
- *Emit events:*
  - *StaleDataEvent with metadata  $(t, t_{current}, TTL)$*
  - *DegradedInferenceModeActivated*
- *Log lag duration for post-mortem analysis*
- *Continue predictions using frozen weights only*

5. Verify recovery when fresh signals satisfy  $TTL(y_t) < 0.8 \cdot \Delta_{max}$ :

- Deactivate `DegradedInferenceMode`
- Resume JKO transport
- Emit `NormalOperationRestoredEvent`

Acceptance criteria:

- Detection time  $< 100$  ms from receiving  $\tilde{y}_t$
- No JKO iteration executed with stale data
- No crash or undefined state
- Predictions continue (degraded) with last valid configuration

**Note 6.1.** Degraded inference mode is critical in high-frequency systems. Operating with stale weights is equivalent to optimizing the wrong objective. It is better to operate with consistent static weights than with weights optimized on obsolete data.

# Chapter 7

# Edge Cases and Operational Limits

This chapter documents boundary conditions for theoretical and operational behavior.

## 7.1 CUSUM: Adaptive Dynamic Threshold

**Test Case 7.1** (Volatility Regime Adaptation). *Validate that the CUSUM threshold adapts correctly to low and high volatility via  $h = k \cdot \sigma_{\text{resid}}$ .*

**Criterion 7.1.** *CUSUM uses dynamic threshold:*

$$h_t = k \cdot \sigma_{\text{resid},t}$$

where  $k \in [3, 5]$  and  $\sigma_{\text{resid},t}$  is rolling residual standard deviation.

*Validation:*

**1. Low volatility:**

- Generate signal with  $\sigma_{\text{true}} = 0.01$
- Estimate  $\hat{\sigma}_{\text{resid}} \approx 0.01$
- Verify  $h_t = k \cdot 0.01$
- Inject small drift  $\Delta\mu = 0.05$
- Detector should trigger when  $G_t^+ > h_t$

**2. High volatility:**

- Generate signal with  $\sigma_{\text{true}} = 0.50$
- Estimate  $\hat{\sigma}_{\text{resid}} \approx 0.50$
- Verify  $h_t = k \cdot 0.50$
- Inject the same drift  $\Delta\mu = 0.05$
- Detector should NOT trigger
- Inject larger drift  $\Delta\mu = 2.0$
- Detector should trigger

**3. Transition:**

- Simulate transition from low to high volatility
- Verify  $h_t$  updates smoothly (rolling window)
- No spurious activations during transition

*Acceptance criterion:*

$$\frac{h_{\text{high}}}{h_{\text{low}}} = \frac{\sigma_{\text{high}}}{\sigma_{\text{low}}} \pm 0.1$$

*Threshold ratio must match volatility ratio within 10% tolerance.*

## 7.2 Orchestrator: Maximum Entropy Convergence

**Test Case 7.2** (Uniform Weights under Total Uncertainty). *Confirm that the system converges to uniform weights  $w = [0.25, 0.25, 0.25, 0.25]$  when Sinkhorn regularization tends to infinity ( $\varepsilon \rightarrow \infty$ ), representing maximum uncertainty.*

**Criterion 7.2.** *In Sinkhorn, entropy regularization  $\varepsilon$  controls smoothing:*

$$\min_{\pi \in \Pi(\mu, \nu)} \{ \langle C, \pi \rangle - \varepsilon H(\pi) \}$$

with  $H(\pi) = - \sum_{ij} \pi_{ij} \log \pi_{ij}$ .

*Limit behavior:*

1.  $\varepsilon \rightarrow 0$ : deterministic optimal transport
2.  $\varepsilon \rightarrow \infty$ : maximum entropy dispersion

*Validation:*

1. Configure four branches ( $A, B, C, D$ )
2. Run with  $\varepsilon_k = 10^k$  for  $k \in \{0, 1, 2, 3, 4\}$
3. Record weights  $\{w_A^k, w_B^k, w_C^k, w_D^k\}$
4. Verify convergence:

$$\lim_{\varepsilon \rightarrow \infty} \{w_i\} = \left\{ \frac{1}{4}, \frac{1}{4}, \frac{1}{4}, \frac{1}{4} \right\}$$

5. Numerical criterion for  $\varepsilon = 10^4$ :

$$\max_{i \in \{A, B, C, D\}} |w_i - 0.25| < 0.01$$

*Interpretation:*

- High  $\varepsilon$  implies indistinguishable branches, thus uniform weights
- Maximum entropy principle (Jaynes)

**Note 7.1.** *This test validates that the orchestrator respects fundamental information theory. Under total uncertainty it must not favor any branch.*

## 7.3 Branch D: Time Reparametrization Invariance

**Test Case 7.3** (Signature Invariance to Time Warping). *Test a signal and its time-stretched variants; the rough path signature must be identical under strictly increasing reparametrizations.*

**Criterion 7.3.** *Signature invariance:*

$$Sig(\gamma) = Sig(\gamma \circ \phi)$$

for any strictly increasing  $\phi : [0, 1] \rightarrow [0, 1]$  with  $\phi(0) = 0$ ,  $\phi(1) = 1$ .

*Validation:*

1. Generate reference signal  $\gamma(t)$  for  $t \in [0, 1]$
2. Compute truncated signature  $S_0 = Sig^{(M)}(\gamma)$

3. Apply nonlinear reparametrizations:

$$\phi_1(t) = t^2$$

$$\phi_2(t) = \sqrt{t}$$

$$\phi_3(t) = \frac{1}{2}(1 - \cos(\pi t))$$

4. For each, compute:

$$\gamma_i(t) = \gamma(\phi_i(t)), \quad S_i = \text{Sig}^{(M)}(\gamma_i)$$

5. Verify:

$$\|S_i - S_0\|_2 < \epsilon_{inv} \quad \forall i \in \{1, 2, 3\}$$

with  $\epsilon_{inv} = 10^{-8}$

Negative control:

1. Apply a non-monotone reparametrization  $\psi(t) = t^2 - 0.5t$

2. Verify  $\text{Sig}(\gamma \circ \psi) \neq \text{Sig}(\gamma)$

**Note 7.2.** Signature invariance to time warping captures intrinsic path geometry regardless of execution speed. In finance, this means the signature captures the shape of a price move whether it occurred in 1 minute or 1 hour.

## 7.4 Edge Case Summary Table

Table 7.1: Limit Scenarios and Edge Cases

Module	Test Scenario	Purpose
CUSUM	Dynamic threshold	Validate $h = k \cdot \sigma_{resid}$ adapts to low/high volatility
Orchestrator	Maximum entropy ( $\varepsilon \rightarrow \infty$ )	Confirm convergence to uniform weights [0.25, 0.25, 0.25, 0.25]
Branch D	Time reparametrization invariance	Signature remains identical under time warping ( $< 10^{-8}$ error)

# Chapter 8

## Acceptance Criteria Summary

Table 8.1: Global Validation Table

Test	Method	Acceptance Criterion
Generalization	Rolling walk-forward	No look-ahead bias; training uses $t < t_{test}$
Meta-optimization efficiency	Bayesian optimization (GP)	Expected Improvement beats random search
Temporal integrity	TTL staleness metric	Cancel JKO update if $\text{TTL}(y) > \Delta_{max}$
Lag injection	Artificial delay $> \Delta_{max}$	Immediate degraded mode; JKO suspended
Nyquist soft limit	Multifractal aliasing (SIA)	Freeze topological branch before $H$ error exceeds 10%
HJB-DGM stability	Gradient monitoring	Clip and reduce $\eta$ under sustained explosion
Viscosity solutions	Comparison principle	Error $< 5\%$ vs reference; sub/supersolution verified
Mode collapse	Training entropy	Variance ratio in $[0.3, 1.2]$ and proportional to term
I/O atomicity	Write interruption	Safe cold start; recovery $< 30$ s
Silent corruption	Bit rot detection	SHA-256 check with fallback to previous snapshot
Dynamic CUSUM	Volatility adaptation	Threshold ratio follows volatility ratio within 10%
Maximum entropy	Sinkhorn $\varepsilon \rightarrow \infty$	Uniform weights within 1%
Time invariance	Path reparametrization	Signature error $< 10^{-8}$
Bit consistency	Cross-platform tests	CPU/GPU/FPGA equivalence within quantization
Numerical drift	Branch D fixed-point	Cumulative divergence $\leq N \cdot \epsilon_{quant}$ after 10k iterations

### 8.1 Final Considerations

All protocols here are language-agnostic and based solely on the mathematical and algorithmic foundations of the Universal Stochastic Predictor.

**Note 8.1.** *The full test suite must run before any production deployment. Results must be documented in a validation report including:*

- *Executive summary of all tests executed*
- *Numeric metrics for each criterion*
- *Convergence plots and distributions*

- *Failure case analysis (if any)*
- *Parameter calibration recommendations*

**Note 8.2.** *This protocol is a living document that must evolve with the system. Any architectural change in theory must be reflected in new test cases or updates to existing ones.*