

Python Test Suite for the Universal Stochastic Predictor

Adaptive Meta-Prediction Development Consortium

February 21, 2026

Contents

1 Testing Environment Setup	2
1.1 Dependencies and Tools	2
1.2 Directory Structure	2
1.3 Shared Fixtures (conftest.py)	3
2 Unit Tests: Generation and Analysis	5
2.1 Stable Variable Generation (Chambers-Mallows-Stuck)	5
2.2 WTMM Test (Wavelet Transform Modulus Maxima)	6
2.3 DGM Entropy Test (Mode Collapse Detection)	6
2.4 Property-Based Testing (Hypothesis)	7
3 Robustness Tests: CUSUM and Circuit Breakers	10
3.1 Standard CUSUM Test	10
3.2 CUSUM with Adaptive Kurtosis	11
3.3 Circuit Breaker Test (Singularity)	12
4 Integration Tests: DGM and Orchestrator	13
4.1 Deep Galerkin Method Test	13
4.2 Sinkhorn and JKO Test	14
5 I/O and Persistence Tests	15
5.1 Atomic Snapshotting Test	15
6 Hardware Tests: CPU/GPU Parity	17
6.1 Numerical Consistency Test	17
6.2 Hardware Parity with Quantization (FPGA Simulation)	18
7 XLA VRAM and JIT Cache Assertions	20
7.1 Asynchronous Device Dispatch (No Host-Device Synchronization)	20
7.1.1 Telemetry Non-Blocking Guarantee	20
7.2 Vectorized Multi-Tenancy (jax.vmap Parity)	21
7.2.1 Batch Execution Bit-Exactness	21
7.3 JIT Cache Efficiency Under Load Shedding	24
7.3.1 Zero-Recompilation Guarantee for Emergency Mode	24
7.4 Atomic Configuration Mutation (POSIX Guarantees)	25
7.4.1 Temporary File Protocol Enforcement	25
8 Edge Cases and Degraded Mode	28
8.1 Degraded Mode Test (TTL Violation)	28
8.2 Extreme Kurtosis Test	29
9 Walk-Forward Validation	30

10 Strict Causality Validation	32
10.0.1 Causal Mask Test: Intentional Future Poisoning	32
10.0.2 SDE Fuzzing: Extreme Time Steps	32
10.1 No-Clairvoyance via Pointer Inspection	33
11 Test Coverage Summary	36
11.1 Coverage Matrix	36
11.2 Full Suite Execution	36
11.2.1 Environment Validation in CI/CD	36
11.2.2 Execution Commands	37
11.3 Global Acceptance Criteria	37

Chapter 1

Testing Environment Setup

1.1 Dependencies and Tools

```
1 # requirements-test.txt
2 pytest>=7.4.0
3 pytest-cov>=4.1.0
4 pytest-xdist>=3.3.0 # Test parallelization
5 hypothesis>=6.82.0 # Property-based testing
6 jax[cpu]>=0.4.13 # CPU tests
7 jax[cuda12]>=0.4.13 # GPU tests (optional)
8 numpy>=1.24.0
9 scipy>=1.11.0
10 PyWavelets>=1.4.1
11 msgpack>=1.0.5
12 optuna>=3.2.0
13
14 # Validation tools
15 flake8>=6.0.0
16 mypy>=1.4.0
17 black>=23.7.0
```

1.2 Directory Structure

```
1 tests/
2   __init__.py
3   conftest.py          # Shared fixtures
4   test_unit/
5     test_cms_levy.py    # Stable variable generation
6     test_wtmm.py        # Multifractal analysis
7     test_malliavin.py   # Sensitivity computations
8     test_signatures.py  # Branch D
9     test_entropy.py     # DGM entropy
10    test_integration/
11      test_sde_solvers.py # Euler-Maruyama, Milstein
12      test_sinkhorn.py   # Optimal transport
13      test_dgm.py        # Deep Galerkin Method
14      test_orchestrator.py # Full JKO
15    test_robustness/
16      test_cusum.py      # Change detection
17      test_cusum_kurtosis.py # CUSUM with kurtosis
18      test_circuit_breaker.py # Singularities
19      test_outliers.py   # Extreme values
20    test_io/
21      test_snapshotting.py # Persistence
22      test_recovery.py   # Atomic recovery
```

```

23 test_hardware/
24     test_cpu_gpu_parity.py      # Numerical consistency
25     test_numerical_drift.py    # Fixed-point drift
26 test_validation/
27     test_walk_forward.py       # Causal validation
28     test_optuna_tuning.py     # Meta-optimization
29 test_edge_cases/
30     test_ttl_degraded_mode.py # Degraded mode
31     test_mode_collapse.py    # DGM collapse
32     test_extreme_kurtosis.py # Kurtosis > 20

```

1.3 Shared Fixtures (conftest.py)

```

1 import pytest
2 import jax
3 import jax.numpy as jnp
4 import numpy as np
5
6 @pytest.fixture
7 def rng_key():
8     """Deterministic PRN key for reproducibility."""
9     return jax.random.PRNGKey(42)
10
11 @pytest.fixture
12 def synthetic_brownian():
13     """Generate a synthetic Brownian trajectory for tests."""
14     np.random.seed(123)
15     T = 1.0
16     N = 1000
17     dt = T / N
18     dW = np.random.randn(N) * np.sqrt(dt)
19     X = np.cumsum(dW)
20     return X, dt
21
22 @pytest.fixture
23 def synthetic_levy_stable():
24     """Generate a stable Levy process sample."""
25     from scipy.stats import levy_stable
26     np.random.seed(456)
27     alpha = 1.5
28     beta = 0.0
29     samples = levy_stable.rvs(alpha, beta, size=1000)
30     return samples, alpha
31
32 @pytest.fixture
33 def mock_market_data():
34     """Synthetic market data with regime change."""
35     np.random.seed(789)
36     regime1 = np.random.randn(500) * 0.01 + 100
37     regime2 = np.random.randn(500) * 0.05 + 105
38     data = np.concatenate([regime1, regime2])
39     return data
40
41 @pytest.fixture
42 def dgm_reference_solution():
43     """Reference solution for DGM validation."""
44     def bs_call(S, K, T, r, sigma):
45         from scipy.stats import norm
46         d1 = (np.log(S/K) + (r + 0.5*sigma**2)*T) / (sigma * np.sqrt(T))
47         d2 = d1 - sigma * np.sqrt(T)
48         return S * norm.cdf(d1) - K * np.exp(-r*T) * norm.cdf(d2)

```

```
49
50     return bs_call
51
52 @pytest.fixture(params=['cpu', 'gpu'])
53 def device(request):
54     """Device parameterization for parity tests."""
55     device_name = request.param
56     if device_name == 'gpu' and not jax.devices('gpu'):
57         pytest.skip("GPU not available")
58     return device_name
```

Chapter 2

Unit Tests: Generation and Analysis

2.1 Stable Variable Generation (Chambers-Mallows-Stuck)

```
1 # tests/test_unit/test_cms_levy.py
2 import pytest
3 import numpy as np
4 from scipy.stats import levy_stable, kstest
5 from Python.integrators.levy import generate_levy_stable
6
7 def test_cms_parameter_recovery(rng_key):
8     """
9         Test: Validate CMS produces distributions with the desired parameters.
10     """
11     alpha = 1.5
12     beta = 0.5
13     gamma = 1.0
14     delta = 0.0
15     N = 10000
16
17     samples = generate_levy_stable(rng_key, alpha, beta, gamma, delta, N)
18     samples_np = np.array(samples)
19
20     statistic, pvalue = kstest(
21         samples_np,
22         lambda x: levy_stable.cdf(x, alpha, beta, loc=delta, scale=gamma)
23     )
24
25     assert pvalue > 0.05, f"KS test failed: p={pvalue:.4f}"
26     assert not np.any(np.isnan(samples_np)), "NaN values detected"
27     assert not np.any(np.isinf(samples_np)), "Inf values detected"
28
29 def test_cms_symmetry():
30     """
31         Test: Validate symmetry when beta = 0.
32     """
33     alpha = 1.8
34     beta = 0.0
35     N = 5000
36
37     samples = generate_levy_stable(
38         jax.random.PRNGKey(999), alpha, beta, 1.0, 0.0, N
39     )
40     samples_np = np.array(samples)
41
42     median = np.median(samples_np)
43     assert abs(median) < 0.1, f"Asymmetry detected: median={median:.4f}"
```

2.2 WTMM Test (Wavelet Transform Modulus Maxima)

```
1 # tests/test_unit/test_wtmm.py
2 import pytest
3 import numpy as np
4 import jax.numpy as jnp
5 from Python.sia.wtmm import estimate_holder_exponent
6
7 def test_wtmm_brownian_motion(synthetic_brownian):
8     """
9     Test: WTMM recovers H = 0.5 for Brownian motion.
10    """
11    signal, dt = synthetic_brownian
12
13    H_estimated = estimate_holder_exponent(jnp.array(signal), besov_c=1.5)
14    assert abs(float(H_estimated) - 0.5) < 0.05, \
15        f"Holder exponent estimation failed: H={H_estimated:.3f}"
16
17 def test_wtmm_fractional_brownian():
18     """
19     Test: WTMM with fBm of known Hurst exponent.
20    """
21    from fbm import FBM
22
23    H_true = 0.7
24    n = 1024
25    fbm_gen = FBM(n=n, hurst=H_true, length=1, method='daviesharte')
26    signal = fbm_gen.fbm()
27
28    H_estimated = estimate_holder_exponent(jnp.array(signal), besov_c=2.0)
29    error_rel = abs(float(H_estimated) - H_true) / H_true
30    assert error_rel < 0.10, \
31        f"fbm Holder estimation error: H_true={H_true}, H_est={H_estimated:.3f}"
32
33 def test_wtmm_cone_influence():
34     """
35     Test: Verify the Besov cone of influence is respected.
36    """
37    signal = np.concatenate([
38        np.ones(500),
39        np.ones(500) * 3.0
40    ])
41
42    H_estimated = estimate_holder_exponent(jnp.array(signal), besov_c=1.0)
43    assert float(H_estimated) < 0.3, \
44        f"Jump detection failed: H={H_estimated:.3f} (expected < 0.3)"
```

2.3 DGM Entropy Test (Mode Collapse Detection)

```
1 # tests/test_unit/test_entropy.py
2 import pytest
3 import jax
4 import jax.numpy as jnp
5 from Python.kernels.kernel_b import compute_entropy_dgm
6
7 def test_entropy_uniform_distribution():
8     """
9     Test: Entropy of a uniform distribution should be maximal.
10    """
11    samples = jnp.linspace(0, 1, 1000)
```

```

13     class MockModel:
14         def __call__(self, t, x):
15             return x[0]
16
17     model = MockModel()
18     entropy = compute_entropy_dgm(model, t=0.5, x_samples=samples[:, None])
19     assert entropy > -0.5, f"Entropy too low: {entropy:.3f}"
20
21 def test_entropy_collapsed_solution():
22     """
23     Test: Detect collapsed (constant) solution.
24     """
25     class CollapsedModel:
26         def __call__(self, t, x):
27             return 1.0
28
29     model = CollapsedModel()
30     samples = jnp.linspace(-1, 1, 500)
31
32     entropy = compute_entropy_dgm(model, t=0.5, x_samples=samples[:, None])
33     assert entropy < -3.0, \
34         f"Collapsed solution not detected: H={entropy:.3f}"
35
36 def test_mode_collapse_criterion():
37     """
38     Test: Validate criterion  $H_{DGM} \geq \gamma * H[g]$ .
39     """
40     from Python.kernels.kernel_b import check_mode_collapse
41
42     class NormalModel:
43         def __call__(self, t, x):
44             return jnp.sin(x[0])
45
46     model = NormalModel()
47     t_eval = jnp.linspace(0, 0.9, 20)
48     x_samples = jnp.linspace(-3, 3, 100)[:, None]
49
50     H_terminal = 1.5
51     gamma = 0.5
52
53     collapsed, avg_entropy = check_mode_collapse(
54         model, t_eval, x_samples, H_terminal, gamma
55     )
56
57     assert not collapsed, \
58         f"False positive collapse detection: H_avg={avg_entropy:.3f}"

```

2.4 Property-Based Testing (Hypothesis)

This section implements property-based fuzzing to generate extreme CMS parameters and validate mathematical invariants.

```

1 # tests/test_unit/test_levy_fuzzing.py
2 import pytest
3 from hypothesis import given, strategies as st, settings, HealthCheck
4 import jax.numpy as jnp
5 from Python.integrators.levy import stable_variate_cms
6
7 @settings(
8     max_examples=500,
9     suppress_health_check=[HealthCheck.too_slow, HealthCheck.filter_too_much]
10 )

```

```

11 @given(
12     alpha=st.floats(min_value=0.5, max_value=2.0),
13     beta=st.floats(min_value=-1.0, max_value=1.0),
14     sigma=st.floats(min_value=0.1, max_value=10.0),
15     num_samples=st.integers(min_value=100, max_value=5000)
16 )
17 def test_levy_cms_basic_properties(alpha, beta, sigma, num_samples):
18     """
19     Property Test 1: CMS generator must satisfy basic invariants.
20     """
21     samples = jnp.array([
22         stable_variate_cms(alpha, beta, sigma)
23         for _ in range(num_samples)
24     ])
25
26     assert jnp.all(jnp.isfinite(samples)), \
27         f"NaN/Inf detected for alpha={alpha}, beta={beta}, sigma={sigma}"
28
29     if alpha >= 1.8:
30         empirical_var = jnp.var(samples)
31         expected_var = (sigma ** 2) * 1.5
32         assert empirical_var < expected_var * 1.5, \
33             f"Variance too high: {empirical_var:.2e} vs expected {expected_var:.2e}"
34
35     empirical_mean = jnp.mean(samples)
36     if abs(beta) < 0.9:
37         assert abs(empirical_mean) < 3 * sigma / jnp.sqrt(num_samples), \
38             f"Mean drift detected: {empirical_mean:.2e}"
39
40 @settings(max_examples=300)
41 @given(
42     alpha=st.floats(min_value=0.5, max_value=2.0),
43     beta=st.floats(min_value=-1.0, max_value=1.0),
44     sigma=st.floats(min_value=0.1, max_value=10.0)
45 )
46 def test_levy_cms_stability_under_extreme_params(alpha, beta, sigma):
47     """
48     Property Test 2: CMS must be stable under extreme parameters.
49     """
50     samples = jnp.array([
51         stable_variate_cms(alpha, beta, sigma)
52         for _ in range(100)
53     ])
54
55     log_abs = jnp.log(jnp.abs(samples) + 1e-8)
56     assert jnp.all(jnp.isfinite(log_abs)), \
57         f"Log transform produced NaN: alpha={alpha}, beta={beta}"
58
59     log_var = jnp.var(log_abs)
60     assert log_var < 50.0, \
61         f"Excessive log-variance: {log_var:.2e} for alpha={alpha}"
62
63 @settings(max_examples=200)
64 @given(
65     alpha1=st.floats(min_value=0.5, max_value=2.0),
66     alpha2=st.floats(min_value=0.5, max_value=2.0)
67 )
68 def test_levy_cms_characteristic_exponent(alpha1, alpha2):
69     """
70     Property Test 3: Infinite divisibility properties.
71     """
72     np.random.seed(123)
73

```

```
74     sigma1, sigma2 = 1.0, 1.5
75
76     samples1 = jnp.array([stable_variate_cms(alpha1, 0.0, sigma1) for _ in range(500)])
77     samples2 = jnp.array([stable_variate_cms(alpha1, 0.0, sigma2) for _ in range(500)])
78
79     sum_samples = samples1 + samples2
80
81     kurt_sum = jnp.mean((sum_samples - jnp.mean(sum_samples)) ** 4) / (jnp.var(
82         sum_samples) ** 2)
83     kurt1 = jnp.mean((samples1 - jnp.mean(samples1)) ** 4) / (jnp.var(samples1) ** 2 + 1e
84     -8)
85
86     assert jnp.isfinite(kurt_sum), "Kurtosis diverges in Levy sum"
```

Chapter 3

Robustness Tests: CUSUM and Circuit Breakers

3.1 Standard CUSUM Test

```
1 # tests/test_robustness/test_cusum.py
2 import pytest
3 import numpy as np
4 import jax.numpy as jnp
5 from Python.orchestrator.cusum import CUSUM
6
7 def test_cusum_no_change(mock_market_data):
8     """
9     Test: CUSUM should not trigger on stationary data.
10    """
11    data = mock_market_data[:500]
12
13    cusum = CUSUM(h=5.0, k=0.5, alpha_var=0.1)
14    alarms = []
15
16    for obs in data:
17        alarm = cusum.update(obs)
18        alarms.append(alarm)
19
20    num_alarms = np.sum(alarms)
21    assert num_alarms == 0, \
22        f"False positives detected: {num_alarms} alarms in stable regime"
23
24 def test_cusum_detects_change(mock_market_data):
25     """
26     Test: CUSUM should detect abrupt regime change.
27     """
28    data = mock_market_data
29
30    cusum = CUSUM(h=3.0, k=0.5, alpha_var=0.05)
31    alarms = []
32
33    for obs in data:
34        alarm = cusum.update(obs)
35        alarms.append(alarm)
36
37    alarm_indices = np.where(alarms)[0]
38    assert len(alarm_indices) > 0, "Change point not detected"
39
40    first_alarm = alarm_indices[0]
41    assert 480 < first_alarm < 550, \
42        f"Change detected too far from true point: {first_alarm} vs 500"
```

3.2 CUSUM with Adaptive Kurtosis

```
1 # tests/test_robustness/test_cusum_kurtosis.py
2 import pytest
3 import numpy as np
4 import jax.numpy as jnp
5 from Python.orchestrator.cusum import CUSUMWithKurtosis
6
7 def test_kurtosis_calculation():
8     """
9     Test: Validate empirical kurtosis calculation.
10    """
11    np.random.seed(111)
12    gaussian_data = np.random.randn(10000)
13
14    cusum = CUSUMWithKurtosis(h=5.0, k=0.5, window_size=252)
15
16    for obs in gaussian_data[:1000]:
17        _ = cusum.update(obs)
18
19    kurtosis = cusum.get_kurtosis()
20
21    assert 2.5 < kurtosis < 3.5, \
22        f"Gaussian kurtosis estimation failed: kappa={kurtosis:.2f}"
23
24 def test_adaptive_threshold_heavy_tails():
25     """
26     Test: Adaptive threshold increases under heavy tails.
27     """
28     from scipy.stats import t
29     np.random.seed(222)
30     heavy_tail_data = t.rvs(df=3, size=1000) * 2.0
31
32     cusum = CUSUMWithKurtosis(h=5.0, k=0.5, window_size=100)
33
34     h_values = []
35     kurtosis_values = []
36
37     for obs in heavy_tail_data:
38         _, kappa, h_adapt = cusum.update_with_kurtosis(obs)
39         h_values.append(h_adapt)
40         kurtosis_values.append(kappa)
41
42     final_kappa = kurtosis_values[-1]
43     final_h = h_values[-1]
44
45     assert final_kappa > 5.0, \
46         f"Heavy tail kurtosis not detected: kappa={final_kappa:.2f}"
47
48     h_fixed = 5.0
49     assert final_h > h_fixed, \
50         f"Adaptive threshold not increased: h_adapt={final_h:.2f} vs h_fixed={h_fixed}"
51
52 def test_false_positive_reduction():
53     """
54     Test: Adaptive CUSUM reduces false positives in high kurtosis.
55     """
56     np.random.seed(333)
57     from scipy.stats import t
58     stable_heavy = t.rvs(df=4, size=1000) * 3.0
59
60     cusum_std = CUSUM(h=3.0, k=0.5, alpha_var=0.1)
61     alarms_std = [cusum_std.update(obs) for obs in stable_heavy]
```

```

62
63     cusum_adapt = CUSUMWithKurtosis(h=3.0, k=0.5, window_size=100)
64     alarms_adapt = []
65     for obs in stable_heavy:
66         alarm, _, _ = cusum_adapt.update_with_kurtosis(obs)
67         alarms_adapt.append(alarm)
68
69     num_alarms_std = np.sum(alarms_std[-500:])
70     num_alarms_adapt = np.sum(alarms_adapt[-500:])
71
72     assert num_alarms_adapt < num_alarms_std, \
73         f"Adaptive CUSUM did not reduce false positives: " \
74         f"{num_alarms_adapt} vs {num_alarms_std}"

```

3.3 Circuit Breaker Test (Singularity)

```

1 # tests/test_robustness/test_circuit_breaker.py
2 import pytest
3 import jax.numpy as jnp
4 from Python.predictor import UniversalPredictor
5 from Python.config import PredictorConfig
6
7 def test_circuit_breaker_activation():
8     """
9     Test: Circuit breaker must activate when H_t < H_min.
10    """
11    config = PredictorConfig(holder_threshold=0.4)
12    predictor = UniversalPredictor(config)
13
14    signal_with_jump = jnp.concatenate([
15        jnp.ones(100) * 50.0,
16        jnp.ones(100) * 100.0
17    ])
18
19    for i, obs in enumerate(signal_with_jump):
20        result = predictor.step_with_telemetry(obs, previous_target=obs)
21
22        if i >= 105:
23            if result.holder_exponent < config.holder_threshold:
24                assert result.emergency_mode, \
25                    "Emergency mode not activated despite low Holder"
26
27                assert result.weights[3] > 0.95, \
28                    f"Kernel D not forced: weights={result.weights}"
29
30                assert result.mode == "Emergency", \
31                    f"Robust loss not activated: mode={result.mode}"
32
33        break

```

Chapter 4

Integration Tests: DGM and Orchestrator

4.1 Deep Galerkin Method Test

```
1 # tests/test_integration/test_dgm.py
2 import pytest
3 import jax
4 import jax.numpy as jnp
5 from Python.kernels.kernel_b import DGM_HJB_Solver, loss_hjb
6
7 def test_dgm_black_scholes(dgm_reference_solution):
8     """
9         Test: Validate DGM against Black-Scholes analytical solution.
10    """
11    S0 = 100.0
12    K = 100.0
13    T = 1.0
14    r = 0.05
15    sigma = 0.2
16
17    bs_price = dgm_reference_solution(S0, K, T, r, sigma)
18
19    key = jax.random.PRNGKey(42)
20    model = DGM_HJB_Solver(in_size=2, key=key)
21
22    def hamiltonian_bs(x, v_x, v_xx):
23        S = x[0]
24        return r*S*v_x[0] + 0.5*sigma**2*S**2*v_xx[0,0] - r
25
26    def terminal_cond(x):
27        return jnp.maximum(x[0] - K, 0.0)
28
29    t_batch = jnp.linspace(0, T, 100)
30    S_batch = jnp.linspace(80, 120, 100)[:, None]
31
32    loss = loss_hjb(
33        model, t_batch, S_batch,
34        hamiltonian_bs, terminal_cond,
35        boundary_cond_fn=None, T=T
36    )
37
38    V_dgm = model(0.0, jnp.array([S0]))
39    error_rel = abs(float(V_dgm) - bs_price) / bs_price
40
41    assert loss < 1.0, f"DGM loss too high (untrained): {loss:.4f}"
```

4.2 Sinkhorn and JKO Test

```
1 # tests/test_integration/test_orchestrator.py
2 import pytest
3 import jax.numpy as jnp
4 from Python.orchestrator.jko import JK0_Discreto
5
6 def test_sinkhorn_convergence():
7     """
8     Test: Sinkhorn should converge for epsilon >= 1e-4.
9     """
10    jko = JK0_Discreto(epsilon=1e-3)
11
12    weights_prev = jnp.array([0.25, 0.25, 0.25, 0.25])
13    gradients = jnp.array([0.1, -0.2, 0.05, -0.1])
14
15    weights_new = jko.solve_ot_step(weights_prev, gradients, tau=0.1)
16
17    assert jnp.abs(jnp.sum(weights_new) - 1.0) < 1e-8, \
18        "Simplex constraint violated"
19
20    assert jnp.all(weights_new >= 0), "Negative weights detected"
21
22 def test_jko_energy_descent():
23     """
24     Test: JKO must reduce energy along gradient direction.
25     """
26    jko = JK0_Discreto(epsilon=1e-2)
27
28    weights_prev = jnp.array([0.5, 0.2, 0.2, 0.1])
29    gradients = jnp.array([1.0, -0.5, -0.3, -0.2])
30
31    weights_new = jko.solve_ot_step(weights_prev, gradients, tau=0.1)
32
33    assert weights_new[0] < weights_prev[0], \
34        f"JKO did not reduce high-energy kernel: " \
35        f"\n{weights_new[0]:.3f} vs {weights_prev[0]:.3f}"
```

Chapter 5

I/O and Persistence Tests

5.1 Atomic Snapshotting Test

```
1 # tests/test_io/test_snapshotting.py
2 import pytest
3 import tempfile
4 import os
5 from Python.predictor import UniversalPredictor
6 from Python.config import PredictorConfig
7
8 def test_snapshot_save_load_integrity():
9     """
10     Test: Snapshot must preserve full state with checksum.
11     """
12     config = PredictorConfig()
13     predictor1 = UniversalPredictor(config)
14
15     for _ in range(50):
16         obs = 100.0 + np.random.randn()
17         predictor1.step_with_telemetry(obs, previous_target=obs)
18
19     with tempfile.NamedTemporaryFile(delete=False, suffix='.msgpack') as f:
20         filepath = f.name
21
22     try:
23         predictor1.save_snapshot(filepath)
24
25         predictor2 = UniversalPredictor(config)
26         predictor2.load_snapshot(filepath)
27
28         result1 = predictor1.step_with_telemetry(
29             105.0, previous_target=105.0
30         )
31         result2 = predictor2.step_with_telemetry(
32             105.0, previous_target=105.0
33         )
34
35         assert jnp.allclose(result1.weights, result2.weights, atol=1e-6), \
36             "Weights mismatch after snapshot restore"
37
38         assert jnp.allclose(
39             result1.holder_exponent, result2.holder_exponent, atol=1e-6
40         ), "Holder exponent mismatch"
41
42     finally:
43         os.unlink(filepath)
44
```

```

45 def test_snapshot_corruption_detection():
46     """
47     Test: Corrupted snapshot must be rejected.
48     """
49     config = PredictorConfig()
50     predictor1 = UniversalPredictor(config)
51
52     with tempfile.NamedTemporaryFile(delete=False, suffix='.msgpack') as f:
53         filepath = f.name
54
55     try:
56         predictor1.save_snapshot(filepath)
57
58         with open(filepath, 'rb+') as f:
59             f.seek(100)
60             f.write(b'\x00\x00\x00\x00')
61
62         predictor2 = UniversalPredictor(config)
63
64         with pytest.raises(ValueError, match="Checksum mismatch"):
65             predictor2.load_snapshot(filepath)
66
67     finally:
68         os.unlink(filepath)
69
70 def test_snapshot_includes_telemetry():
71     """
72     Test: Snapshot must include kurtosis, DGM entropy, and flags.
73     """
74     import msgpack
75
76     config = PredictorConfig()
77     predictor = UniversalPredictor(config)
78
79     for _ in range(300):
80         obs = 100.0 + np.random.randn() * 5.0
81         predictor.step_with_telemetry(obs, previous_target=obs)
82
83     with tempfile.NamedTemporaryFile(delete=False, suffix='.msgpack') as f:
84         filepath = f.name
85
86     try:
87         predictor.save_snapshot(filepath)
88
89         with open(filepath, 'rb') as f:
90             content = f.read()
91
92             data_bytes = content[:-64]
93             payload = msgpack.unpackb(data_bytes)
94
95             assert 'telemetry' in payload, "Telemetry missing from snapshot"
96             assert 'kurtosis' in payload['telemetry'], "Kurtosis not saved"
97             assert 'dgm_entropy' in payload['telemetry'], "DGM entropy not saved"
98
99             assert 'flags' in payload, "Flags missing from snapshot"
100            assert 'degraded_inference' in payload['flags']
101            assert 'emergency' in payload['flags']
102            assert 'regime_change' in payload['flags']
103            assert 'mode_collapse' in payload['flags']
104
105    finally:
106        os.unlink(filepath)

```

Chapter 6

Hardware Tests: CPU/GPU Parity

6.1 Numerical Consistency Test

```
1 # tests/test_hardware/test_cpu_gpu_parity.py
2 import pytest
3 import jax
4 import jax.numpy as jnp
5 from Python.predictor import UniversalPredictor
6 from Python.config import PredictorConfig
7
8 @pytest.mark.parametrize("device", ["cpu", "gpu"])
9 def test_device_consistency(device):
10     """
11     Test: CPU and GPU must produce equivalent results.
12     """
13     if device == "gpu" and not jax.devices('gpu'):
14         pytest.skip("GPU not available")
15
16     with jax.default_device(jax.devices(device)[0]):
17         config = PredictorConfig()
18         predictor = UniversalPredictor(config)
19
20         np.random.seed(555)
21         data = np.random.randn(100) * 10.0 + 100.0
22
23         results = []
24         for obs in data:
25             result = predictor.step_with_telemetry(obs, previous_target=obs)
26             results.append({
27                 'prediction': float(result.predicted_next),
28                 'holder': float(result.holder_exponent),
29                 'weights': result.weights
30             })
31
32     return results
33
34 def test_cpu_gpu_parity():
35     """
36     Test: Compare CPU and GPU results.
37     """
38     if not jax.devices('gpu'):
39         pytest.skip("GPU not available for parity test")
40
41     results_cpu = test_device_consistency("cpu")
42     results_gpu = test_device_consistency("gpu")
43
44     for i, (cpu, gpu) in enumerate(zip(results_cpu, results_gpu)):
```

```

45     assert jnp.allclose(
46         cpu['weights'], gpu['weights'], rtol=1e-5, atol=1e-6
47     ), f"Weights mismatch at step {i}"
48
49     pred_diff = abs(cpu['prediction'] - gpu['prediction'])
50     assert pred_diff < 1e-4, \
51         f"Prediction mismatch at step {i}: {pred_diff:.2e}"

```

6.2 Hardware Parity with Quantization (FPGA Simulation)

```

1 # tests/test_hardware/test_fixed_point_parity.py
2 import pytest
3 import jax.numpy as jnp
4 import numpy as np
5 from Python.predictor import UniversalPredictor
6
7 def quantize_to_fixed_point(x, int_bits=16, frac_bits=16):
8     """Simulate fixed-point quantization Q16.16."""
9     total_bits = int_bits + frac_bits
10    max_val = (2 ** (total_bits - 1) - 1) / (2 ** frac_bits)
11    min_val = -(2 ** (total_bits - 1)) / (2 ** frac_bits)
12
13    x_clipped = jnp.clip(x, min_val, max_val)
14    x_quantized = jnp.round(x_clipped * (2 ** frac_bits)) / (2 ** frac_bits)
15
16    return x_quantized
17
18 def simulate_fpga_computation(prediction_float32):
19     """Simulate FPGA pipeline: Float32 -> Q16.16 -> Q16.16."""
20     pred_quantized_in = quantize_to_fixed_point(prediction_float32)
21     intermediate = pred_quantized_in * 1.001
22     pred_quantized_out = quantize_to_fixed_point(intermediate)
23     return pred_quantized_out
24
25 def test_fpga_quantization_error():
26     """
27     Test: Q16.16 quantization must introduce <1% error.
28     """
29     config = UniversalPredictor.config
30     predictor = UniversalPredictor(config)
31
32     np.random.seed(777)
33     data = 100.0 + np.random.randn(100) * 5.0
34
35     predictions_float32 = []
36     predictions_quantized = []
37
38     for obs in data:
39         result = predictor.step_with_telemetry(obs, previous_target=obs)
40         pred_f32 = float(result.predicted_next)
41         pred_quantized = float(simulate_fpga_computation(jnp.array(pred_f32)))
42
43         predictions_float32.append(pred_f32)
44         predictions_quantized.append(pred_quantized)
45
46     preds_f32 = np.array(predictions_float32)
47     preds_q = np.array(predictions_quantized)
48
49     mask = np.abs(preds_f32) > 1e-3
50     rel_error = np.abs(preds_f32[mask] - preds_q[mask]) / (np.abs(preds_f32[mask]) + 1e-6)

```

```

51     max_rel_error = np.max(rel_error)
52     mean_rel_error = np.mean(rel_error)
53
54     assert max_rel_error < 0.01, \
55            f"Max relative error too high: {max_rel_error:.2%}"
56
57     assert mean_rel_error < 0.005, \
58            f"Mean relative error too high: {mean_rel_error:.2%}"
59
60
61 def test_fpga_numerical_stability():
62     """
63     Test: Quantization accumulation remains bounded.
64     """
65     config = UniversalPredictor.config
66     predictor_ref = UniversalPredictor(config)
67
68     np.random.seed(888)
69     data = 100.0 + np.random.randn(200) * 5.0
70
71     predictions = []
72     quantized_errors = []
73
74     for i, obs in enumerate(data):
75         result = predictor_ref.step_with_telemetry(obs, previous_target=obs)
76         pred = float(result.predicted_next)
77         pred_q = float(simulate_fpga_computation(jnp.array(pred)))
78
79         predictions.append(pred)
80         quantized_errors.append(abs(pred - pred_q))
81
82     cumulative_error = np.cumsum(quantized_errors)
83     final_cumulative = cumulative_error[-1]
84
85     expected_max_cumulative = 200 * 1.5e-5 * 100
86
87     assert final_cumulative < expected_max_cumulative * 10, \
88            f"Cumulative error unstable: {final_cumulative:.3e}"

```

Chapter 7

XLA VRAM and JIT Cache Assertions

This chapter validates Level 4 Autonomy execution guarantees specific to JAX's XLA compilation backend: asynchronous device dispatch, vectorized multi-tenancy, JIT cache efficiency under load shedding, and atomic I/O during configuration mutations. These tests ensure the implementation maintains performance contracts under production workloads.

7.1 Asynchronous Device Dispatch (No Host-Device Synchronization)

7.1.1 Telemetry Non-Blocking Guarantee

Test Case 7.1 (Prevention of Host-Device Blocking in Orchestrator). *Ensure that orchestration loop returns unbacked `DeviceArray` objects without forcing host synchronization, preserving asynchronous GPU dispatch.*

Implementation 7.1

```
1 # tests/test_xla/test_no_host_device_sync.py
2 import pytest
3 import jax
4 import jax.numpy as jnp
5 from jax.core import Tracer
6 from Python.core.orchestrator import orchestrate_step
7 from Python.api.types import PredictorConfig, InternalState
8
9 def test_no_host_device_sync_in_orchestrator():
10     """
11     Ensure orchestration step returns unbacked DeviceArrays,
12     not host floats.
13
14     CRITICAL: Host-device synchronization blocks XLA dispatch,
15     causing 100-500ms latency spikes and VRAM transfer overhead.
16     """
17     # Initialize configuration and state
18     config = PredictorConfig()
19     key = jax.random.PRNGKey(42)
20     state = InternalState.initialize(config, key)
21
22     # Mock signal input (on device)
23     mock_signal = jnp.array(0.5)
24
25     # Execute orchestration step
26     new_state, prediction = orchestrate_step(mock_signal, state, config, key)
27
28     # ASSERTION 1: Prediction must NOT be a Python float
29     assert not isinstance(prediction, float), \
30         "CRITICAL: Host-Device sync detected. "
```

```

31     "Prediction materialized as Python float instead of DeviceArray."
32
33 # ASSERTION 2: Prediction must be JAX array type
34 assert isinstance(prediction, (jnp.ndarray, jax.Array)), \
35     f"Expected jax.Array, got {type(prediction)}"
36
37 # ASSERTION 3: Array must have device attribute (not backed on host)
38 assert hasattr(prediction, "device"), \
39     "Prediction array must reside on XLA backend (CPU/GPU/TPU)."
40
41 # ASSERTION 4: Verify device is not None (array is backed)
42 assert prediction.device() is not None, \
43     "Prediction array device is None (unbacked array)."
44
45 # ASSERTION 5: State updates must also remain on device
46 assert hasattr(new_state.dgm_entropy, "device"), \
47     "State fields must remain on device for next iteration."
48
49 def test_telemetry_collection_lazy_evaluation():
50     """
51     Verify that telemetry fields use jax.lax.stop_gradient
52     to prevent unnecessary computation during training.
53     """
54     from Python.io.telemetry import collect_telemetry
55
56     # Mock state with tracked gradients
57     state = create_mock_state_with_gradients()
58
59     # Collect telemetry
60     telemetry = collect_telemetry(state, config)
61
62     # ASSERTION: Telemetry values must have stop_gradient applied
63     # This prevents backprop through diagnostic metrics
64     assert not hasattr(telemetry.kurtosis, "_trace"), \
65         "Telemetry fields must use stop_gradient to prevent VRAM waste."

```

Criterion 7.1. Acceptance Criteria:

1. All orchestration outputs must be `jax.Array` or `jnp.ndarray` types, never `Python float` or `int`
2. Arrays must have valid `.device()` attribute indicating XLA backend placement
3. No explicit or implicit conversion to host types (`float()`, `.item()`, `.tolist()`) in hot path
4. Telemetry collection must use `jax.lax.stop_gradient()` on all diagnostic metrics to prevent gradient tracking overhead

Performance Impact: Host-device synchronization introduces 100-500ms latency per sync on typical GPUs. In a 10,000-step training run with telemetry every 10 steps, this accumulates to 100-500 seconds of pure blocking overhead.

7.2 Vectorized Multi-Tenancy (`jax.vmap` Parity)

7.2.1 Batch Execution Bit-Exactness

Test Case 7.2 (Sequential vs Vectorized Execution Parity). Validate that batched `jax.vmap` execution produces bit-exact results compared to sequential loop execution for multi-tenant workloads.

```

Implementation 7.2: # tests/test_xla/test_vmap_multi_tenant_parity.py
1 import pytest
2 import jax
3 import jax.numpy as jnp
4 from Python.core.orchestrator import orchestrate_step
5 from Python.api.types import PredictorConfig, InternalState
6
7
8 def test_vmap_multi_tenant_parity():
9     """
10     Verify that batched vmap execution is bit-exact
11     to sequential execution.
12
13     Multi-tenant deployments use vmap to process N clients
14     in parallel. Any deviation between sequential and batched
15     execution violates determinism guarantees.
16     """
17     batch_size = 128
18     key = jax.random.PRNGKey(42)
19     config = PredictorConfig()
20
21     # Generate batch of signals and states
22     signal_keys = jax.random.split(key, batch_size + 1)
23     signals_batch = jax.random.normal(signal_keys[0], (batch_size, 100))
24
25     # Initialize batched states
26     state_keys = signal_keys[1:]
27     states_batch = jax.vmap(
28         lambda k: InternalState.initialize(config, k)
29     )(state_keys)
30
31     # SCENARIO 1: Sequential execution (baseline)
32     seq_predictions = []
33     seq_states = []
34
35     for i in range(batch_size):
36         new_state, prediction = orchestrate_step(
37             signals_batch[i],
38             states_batch[i],
39             config,
40             state_keys[i]
41         )
42         seq_predictions.append(prediction)
43         seq_states.append(new_state)
44
45     seq_predictions = jnp.stack(seq_predictions)
46
47     # SCENARIO 2: Vectorized execution (production)
48     vmap_orchestrate = jax.vmap(
49         orchestrate_step,
50         in_axes=(0, 0, None, 0)
51     )
52
53     batch_states, batch_predictions = vmap_orchestrate(
54         signals_batch,
55         states_batch,
56         config,
57         state_keys
58     )
59
60     # ASSERTION 1: Bit-exact prediction parity
61     assert jnp.array_equal(seq_predictions, batch_predictions), \
62         "XLA vmap compilation breaks mathematical parity. " \
63         "Sequential and batched predictions must be bit-exact."

```

```

64
65 # ASSERTION 2: State update parity
66 for i in range(batch_size):
67     assert jnp.array_equal(
68         seq_states[i].dgm_entropy,
69         batch_states.dgm_entropy[i]
70     ), f"State divergence at index {i}: entropy mismatch"
71
72     assert jnp.array_equal(
73         seq_states[i].rho,
74         batch_states.rho[i]
75     ), f"State divergence at index {i}: rho weights mismatch"
76
77 # ASSERTION 3: PRNG state advancement consistency
78 # Next iteration must produce identical results
79 next_signal = jnp.ones(batch_size)
80 next_keys = jax.random.split(key, batch_size)
81
82 _, next_pred_seq = orchestrate_step(
83     next_signal[0], seq_states[0], config, next_keys[0]
84 )
85 _, next_pred_batch = vmap_orchestrate(
86     next_signal, batch_states, config, next_keys
87 )[1]
88
89 assert jnp.array_equal(next_pred_seq, next_pred_batch[0]), \
90     "PRNG state divergence detected after vmap execution."
91
92 def test_vmap_memory_efficiency():
93     """
94     Verify that vmap does not allocate N separate XLA buffers
95     for identical config (memory amplification bug).
96     """
97     batch_size = 256
98     config = PredictorConfig()
99
100    # Single execution memory baseline
101    baseline_memory = measure_peak_vram_usage(
102        lambda: orchestrate_step(jnp.array(0.5), state, config, key)
103    )
104
105    # Batched execution memory
106    batch_memory = measure_peak_vram_usage(
107        lambda: jax.vmap(orchestrate_step, in_axes=(0, 0, None, 0))(
108            signals_batch, states_batch, config, keys_batch
109        )
110    )
111
112    # ASSERTION: Batch memory should scale sub-linearly
113    # (not 256x single execution due to config sharing)
114    expected_max_memory = baseline_memory * batch_size * 1.5
115
116    assert batch_memory < expected_max_memory, \
117        f"VRAM amplification detected: {batch_memory / baseline_memory:.1f}x"

```

Criterion 7.2. Acceptance Criteria:

1. Batched execution via `jax.vmap` must produce bit-exact results: `jnp.array_equal(seq_result, batch_result) == True`
2. PRNG state advancement must be consistent between sequential and batched paths
3. Memory usage must scale sub-linearly with batch size (config parameter sharing prevents N -fold duplication)

4. Compilation time: first `vmap` call may be slow (JIT), subsequent calls must be < 5ms per batch

7.3 JIT Cache Efficiency Under Load Shedding

7.3.1 Zero-Recompilation Guarantee for Emergency Mode

Test Case 7.3 (Load Shedding Without XLA Recompilation). Verify that swapping Kernel D signature depths (load shedding: $M \in \{2, 3, 5\}$) executes in $O(1)$ time without triggering JAX cache miss.

Implementation 7.3: # tests/test_xla/test_load_shedding_jit_cache.py

```

1 import pytest
2 import time
3 import jax
4 from Python.api.warmup import warmup_kernel_d_load_shedding
5 from Python.kernels.kernel_d import kernel_d_predict
6 from Python.api.types import PredictorConfig
7
8 def test_load_shedding_warmup_no_recompilation():
9     """
10     Verify that swapping signature depths does not trigger
11     JAX Cache Miss.
12
13     Load shedding is a real-time emergency response to latency
14     spikes. Triggering XLA recompilation (200ms) defeats the
15     purpose of shedding (5ms target).
16     """
17
18     config = PredictorConfig(kernel_d_depth=5)
19     key = jax.random.PRNGKey(42)
20     signal = jax.random.normal(key, (100,))
21
22     # PHASE 1: Warmup compiles M in {2, 3, 5}
23     warmup_kernel_d_load_shedding(config, key)
24
25     # PHASE 2: Baseline execution at M=5 (no load shedding)
26     baseline_start = time.perf_counter()
27     _ = kernel_d_predict(signal, key, config)
28     baseline_time = time.perf_counter() - baseline_start
29
30     # PHASE 3: Trigger load shedding M=5 -> M=2 (emergency mode)
31     shedding_config = config.replace(kernel_d_depth=2)
32
33     shedding_start = time.perf_counter()
34     _ = kernel_d_predict(signal, key, shedding_config)
35     shedding_time = time.perf_counter() - shedding_start
36
37     # ASSERTION 1: Cached execution must be < 10ms
38     # JIT compilation takes ~200ms. Cached execution < 5ms.
39     assert shedding_time < 0.010, \
40         f"CRITICAL: JIT Cache Miss during Load Shedding. " \
41         f"Execution took {shedding_time*1000:.1f}ms (expected < 10ms). " \
42         f"System will hang under stress."
43
44     # ASSERTION 2: Shedding must not be slower than baseline
45     # (Lower depth should be faster or equal)
46     assert shedding_time <= baseline_time * 1.5, \
47         f"Load shedding slower than baseline: " \
48         f"{shedding_time/baseline_time:.2f}x"
49
50     # PHASE 4: Verify cache hit for all depths
51     for depth in [2, 3, 5]:
52         test_config = config.replace(kernel_d_depth=depth)

```

```

53     start = time.perf_counter()
54     _ = kernel_d_predict(signal, key, test_config)
55     exec_time = time.perf_counter() - start
56
57     assert exec_time < 0.010, \
58         f"Cache miss for depth={depth}: {exec_time*1000:.1f}ms"
59
60 def test_jit_cache_size_under_warmup():
61     """
62     Verify that warmup does not exhaust JIT cache memory limits.
63     """
64     import jax._src.xla_bridge as xb
65
66     # Get initial cache size
67     initial_cache = len(xb.get_backend().compile_cache())
68
69     # Warmup all kernel variants
70     warmup_kernel_d_load_shedding(config, key)
71
72     # Get final cache size
73     final_cache = len(xb.get_backend().compile_cache())
74
75     # ASSERTION: Warmup should add exactly 3 entries (M=2,3,5)
76     cache_growth = final_cache - initial_cache
77     assert cache_growth == 3, \
78         f"Unexpected cache growth: {cache_growth} entries " \
79         f"(expected 3 for M in {2,3,5})"

```

Criterion 7.3. Acceptance Criteria:

1. Load shedding execution time: $< 10\text{ms}$ (cached) vs $\sim 200\text{ms}$ (recompilation)
2. Warmup phase must precompile all signature depths: $M \in \{2, 3, 5\}$
3. Cache hit rate after warmup: $\geq 99\%$ for steady-state operation
4. Memory overhead: JIT cache growth ≤ 3 entries per kernel variant

Failure Mode: Without warmup, first load-shedding event triggers 200ms recompilation stall, causing latency SLA violation (target: 50ms p99) and potential cascading failures in multi-tenant deployment.

7.4 Atomic Configuration Mutation (POSIX Guarantees)

7.4.1 Temporary File Protocol Enforcement

Test Case 7.4 (Atomic TOML Mutation via os.replace()). Validate compliance with I/O Specification §3.3 Configuration Mutation Protocol, ensuring POSIX atomic write semantics.

Implementation 7.4: # tests/test_io/test_atomic_toml_mutation.py

```

2 import pytest
3 import os
4 import tempfile
5 from pathlib import Path
6 from unittest.mock import patch, MagicMock
7 from Python.io.config_mutation import mutate_config
8 from Python.core.meta_optimizer import OptimizationResult
9
10 def test_atomic_toml_mutation():
11     """
12     Ensure config mutation uses temporary files and os.replace.

```

```

13
14    POSIX Guarantee: os.replace() is atomic on Linux/macOS.
15    Prevents partial writes visible to concurrent readers.
16    """
17    with tempfile.TemporaryDirectory() as tmpdir:
18        config_path = Path(tmpdir) / "config.toml"
19
20        # Create initial config
21        initial_params = {"cusum_k": 0.5, "learning_rate": 0.01}
22        write_toml(config_path, initial_params)
23
24        # Prepare mutation
25        new_params = {"cusum_k": 0.8}
26        validation_schema = {
27            "cusum_k": {"range": [0.1, 2.0], "locked": False}
28        }
29
30        # MOCK os calls to verify protocol compliance
31        with patch('os.replace') as mock_replace,
32             patch('os.fsync') as mock_fsync,
33             patch('os.open', return_value=3) as mock_open:
34
35            # Trigger mutation
36            mutate_config(new_params, config_path, validation_schema)
37
38            # ASSERTION 1: Verify os.fsync was called
39            # Ensures kernel buffer flush before atomic replace
40            mock_fsync.assert_called_once()
41
42            # ASSERTION 2: Verify os.replace was called with temp file
43            assert mock_replace.call_count == 1
44            args = mock_replace.call_args[0]
45
46            # First arg must be temp file (config.toml.tmp)
47            assert str(args[0]).endswith(".tmp"), \
48                f"Expected temp file, got {args[0]}"
49
50            # Second arg must be target file (config.toml)
51            assert args[1] == config_path, \
52                f"Expected {config_path}, got {args[1]}"
53
54    def test_concurrent_mutation_detection():
55        """
56            Verify that concurrent mutations are rejected
57            (temp file already exists).
58        """
59        from Python.io.config_mutation import ConfigMutationError
60
61        with tempfile.TemporaryDirectory() as tmpdir:
62            config_path = Path(tmpdir) / "config.toml"
63            tmp_path = config_path.with_suffix(".tmp")
64
65            # Create initial config
66            write_toml(config_path, {"param": 1.0})
67
68            # Simulate concurrent mutation (temp file exists)
69            tmp_path.touch()
70
71            # ASSERTION: Mutation must fail with clear error
72            with pytest.raises(ConfigMutationError,
73                               match="Concurrent mutation detected"):
74                mutate_config({"param": 2.0}, config_path, {})
75

```

```

76 def test_audit_log_persistence():
77     """
78     Verify that mutation events are logged to io/mutations.log
79     in JSON Lines format.
80     """
81     with tempfile.TemporaryDirectory() as tmpdir:
82         config_path = Path(tmpdir) / "config.toml"
83         log_path = Path(tmpdir) / "mutations.log"
84
85         write_toml(config_path, {"learning_rate": 0.01})
86
87         # Perform mutation
88         mutate_config(
89             {"learning_rate": 0.015},
90             config_path,
91             {},
92             audit_log_path=log_path
93         )
94
95         # ASSERTION: Audit log must exist and contain entry
96         assert log_path.exists(), "Audit log not created"
97
98         with open(log_path, 'r') as f:
99             entries = [json.loads(line) for line in f]
100
101         assert len(entries) == 1, "Expected 1 audit entry"
102
103         entry = entries[0]
104         assert entry["event"] == "MUTATION_APPLIED"
105         assert "learning_rate" in entry["delta"]
106         assert entry["delta"]["learning_rate"] == [0.01, 0.015]

```

Criterion 7.4. Acceptance Criteria:

1. All config mutations must use temporary file strategy: write to `config.toml.tmp`, then `os.replace()`
2. `os.fsync()` must be called before `os.replace()` to flush kernel buffers
3. Concurrent mutations must be detected and rejected (temp file existence check with `os.O_EXCL`)
4. Audit trail must log all mutations to `io/mutations.log` in JSON Lines format
5. Rollback capability: `config.toml.bak` backup must be created before mutation

POSIX Atomicity Guarantee: `os.replace()` is atomic on POSIX systems (Linux, macOS, BSD). On Windows, requires `ReplaceFileW` API. This prevents readers from observing partial config states during multi-gigabyte meta-optimization campaigns.

Chapter 8

Edge Cases and Degraded Mode

8.1 Degraded Mode Test (TTL Violation)

```
1 # tests/test_edge_cases/test_ttl_degraded_mode.py
2 import pytest
3 import jax.numpy as jnp
4 from Python.predictor import UniversalPredictorWithTelemetry
5 from Python.config import PredictorConfig
6
7 def test_degraded_mode_activation():
8     """
9         Test: Degraded mode activates when TTL exceeds limit.
10    """
11    config = PredictorConfig(staleness_ttl_ns=100_000_000)
12    predictor = UniversalPredictorWithTelemetry(config)
13
14    for _ in range(50):
15        obs = 100.0 + np.random.randn()
16        result = predictor.step_with_telemetry(obs, previous_target=obs)
17
18    predictor.telemetry_logger.ttl_counter = 150
19
20    obs = 100.0
21    result = predictor.step_with_telemetry(obs, previous_target=obs)
22
23    assert result.degraded_inference_mode, \
24        "Degraded mode not activated despite TTL violation"
25
26 def test_degraded_mode_recovery_hysteresis():
27     """
28         Test: Recovery with hysteresis (0.8 * TTL_max).
29    """
30    config = PredictorConfig()
31    predictor = UniversalPredictorWithTelemetry(config)
32
33    predictor.telemetry_logger.ttl_counter = 150
34
35    predictor.telemetry_logger.ttl_counter = 85
36
37    result = predictor.step_with_telemetry(100.0, previous_target=100.0)
38    assert result.degraded_inference_mode, \
39        "Premature recovery (hysteresis not respected)"
40
41    predictor.telemetry_logger.ttl_counter = 75
42
43    result = predictor.step_with_telemetry(100.0, previous_target=100.0)
44    assert not result.degraded_inference_mode, \
```

```
45     "Recovery failed despite TTL below hysteresis threshold"
```

8.2 Extreme Kurtosis Test

```
1 # tests/test_edge_cases/test_extreme_kurtosis.py
2 import pytest
3 import numpy as np
4 from Python.predictor import UniversalPredictorWithTelemetry
5 from Python.config import PredictorConfig
6
7 def test_extreme_kurtosis_detection():
8     """
9     Test: Kurtosis > 20 must generate critical alert.
10    """
11    config = PredictorConfig()
12    predictor = UniversalPredictorWithTelemetry(config)
13
14    from scipy.stats import t
15    np.random.seed(666)
16    extreme_data = t.rvs(df=2, size=500) * 20.0 + 100.0
17
18    kurtosis_values = []
19
20    for obs in extreme_data:
21        result = predictor.step_with_telemetry(obs, previous_target=obs)
22        kurtosis_values.append(float(result.kurtosis))
23
24    final_kurtosis = kurtosis_values[-1]
25
26    assert final_kurtosis > 15.0, \
27        f"Extreme kurtosis not detected: kappa={final_kurtosis:.2f}"
28
29    result = predictor.step_with_telemetry(
30        extreme_data[-1], previous_target=extreme_data[-1]
31    )
32
33    h_adaptive = float(result.adaptive_threshold)
34    h_fixed = config.cusum_h
35
36    assert h_adaptive > 2.0 * h_fixed, \
37        f"Adaptive threshold not sufficiently elevated: " \
38        f"\n\t{h_adaptive:.2f} vs {h_fixed:.2f}"
```

Chapter 9

Walk-Forward Validation

```
1 # tests/test_validation/test_walk_forward.py
2 import pytest
3 import numpy as np
4 from Python.validation import WalkForwardValidator
5 from Python.predictor import UniversalPredictor
6 from Python.config import PredictorConfig
7
8 def test_walk_forward_no_lookahead():
9     """
10     Test: Walk-forward must not use future information.
11     """
12     np.random.seed(777)
13     T = 1000
14     trend = np.linspace(100, 150, T)
15     noise = np.random.randn(T) * 2.0
16     data = trend + noise
17
18     def model_factory(hp):
19         config = PredictorConfig(
20             epsilon=hp.get('epsilon', 1e-3),
21             learning_rate=hp.get('tau', 0.1)
22         )
23         return UniversalPredictor(config)
24
25     def metric_fn(preds, targets):
26         return np.mean(np.abs(preds - targets))
27
28     validator = WalkForwardValidator(
29         model_factory=model_factory,
30         metric_fn=metric_fn,
31         window_size=252,
32         horizon=1,
33         max_memory=500
34     )
35
36     hyperparams = {'epsilon': 1e-2, 'tau': 0.05}
37
38     mae = validator.run(data, hyperparams)
39
40     data_range = np.max(data) - np.min(data)
41     assert mae < 0.1 * data_range, \
42         f"Walk-forward MAE too high: {mae:.2f}"
43
44 def test_walk_forward_regime_change():
45     """
46     Test: Performance under regime change.
47     """
```

```

48 np.random.seed(888)
49
50 regime1 = np.linspace(100, 120, 400) + np.random.randn(400) * 1.0
51 regime2 = np.linspace(120, 100, 400) + np.random.randn(400) * 1.0
52
53 data = np.concatenate([regime1, regime2])
54
55 def model_factory(hp):
56     return UniversalPredictor(PredictorConfig())
57
58 def metric_fn(preds, targets):
59     return np.sqrt(np.mean((preds - targets)**2))
60
61 validator = WalkForwardValidator(
62     model_factory=model_factory,
63     metric_fn=metric_fn,
64     window_size=200,
65     horizon=1
66 )
67
68 rmse = validator.run(data, {})
69
70 assert rmse < 5.0, \
71     f"Predictor failed to adapt to regime change: RMSE={rmse:.2f}"

```

Chapter 10

Strict Causality Validation

This section implements tests that verify strict absence of look-ahead bias.

10.0.1 Causal Mask Test: Intentional Future Poisoning

Criterion 10.1. *Configurable protocol:*

1. Generate a clean series with 500 timesteps and 4 branches
2. For each time t , set data at $t' > t$ to NaN :

$$\tilde{y}[t : t + H] = \text{NaN} \quad \forall H > 0, \forall t \in [0, 500]$$

3. Run prediction on the poisoned series. If the model accesses future data, NaN propagates
4. Verify outputs:

$$\text{Result}_t = \begin{cases} \text{Valid numeric} & (\text{causality respected}) \\ \text{NaN} & (\text{look-ahead detected}) \end{cases}$$

5. Failure condition: if more than 0.1% of samples produce NaN predictions, causality test fails

10.0.2 SDE Fuzzing: Extreme Time Steps

Criterion 10.2. *Branch C solves SDEs. Test stability under drastic step variation:*

1. Regime 1: $\Delta t = 0.01$ (small step)
2. Regime 2: $\Delta t = 0.1$ (moderate)
3. Regime 3: $\Delta t = 0.5$ (stiff)
4. Regime 4: $\Delta t = 1.0$ (pathological)

For each regime, run 1000 trajectories and measure:

$$\text{Stability Metric} = \max_n \left| |X_n^{(\Delta t_1)} - X_n^{(\Delta t_2)}| - \mathcal{O}((\Delta t_1 - \Delta t_2)^p) \right|$$

where p is the order (1 for Euler-Maruyama, 1.5 for Milstein).

Acceptance: in stiff regime $\Delta t = 0.5$ the response must remain bounded:

$$\mathbb{E}[|X_T|] < 10 \times \mathbb{E}[|X_T|^{(\Delta t=0.01)}]$$

10.1 No-Clairvoyance via Pointer Inspection

```
1 # tests/test_causality/test_no_lookahead.py
2 import pytest
3 import jax.numpy as jnp
4 import numpy as np
5 from Python.predictor import UniversalPredictor
6 from Python.config import PredictorConfig
7
8 def test_predict_without_future_access():
9     """
10     Test: predict(t) must not access data with timestamp > t.
11     """
12     config = PredictorConfig()
13     predictor = UniversalPredictor(config)
14
15     np.random.seed(555)
16     data = np.random.randn(100) * 10 + 100
17
18     trap_position = 50
19     trap_value = 1e6
20
21     for i in range(trap_position):
22         result = predictor.step_with_telemetry(
23             data[i],
24             previous_target=data[i]
25         )
26
27     buffer_ptr_before = id(predictor._state.signal_circular_buffer)
28     internal_buffer_before = np.copy(predictor._state.signal_circular_buffer)
29
30     result_at_t = predictor.step_with_telemetry(
31         data[trap_position],
32         previous_target=data[trap_position]
33     )
34
35     predictor._state.signal_circular_buffer = np.concatenate([
36         predictor._state.signal_circular_buffer,
37         jnp.array([trap_value])
38     ])
39
40     for i in range(trap_position + 1, trap_position + 6):
41         if i < len(data):
42             result_later = predictor.step_with_telemetry(
43                 data[i],
44                 previous_target=data[i]
45             )
46
47     predictor_clean = UniversalPredictor(config)
48     for i in range(trap_position + 1):
49         result_clean = predictor_clean.step_with_telemetry(
50             data[i],
51             previous_target=data[i]
52         )
53
54     pred_with_trap = float(result_at_t.predicted_next)
55     pred_without_trap = float(result_clean.predicted_next)
56
57     assert abs(pred_with_trap - pred_without_trap) < 1e-3, \
58         f"Lookahead bias detected: pred_trap={pred_with_trap:.4f}, " \
59         f"pred_clean={pred_without_trap:.4f}"
60
61 def test_causality_via_timestamps():
```

```

62 """
63 Test: Access timestamps should be monotonic.
64 """
65 config = PredictorConfig(wtmm_buffer_size=128)
66 predictor = UniversalPredictor(config)
67
68 original_buffer = predictor._state.signal_circular_buffer
69 access_log = []
70
71 class AccessTrackedBuffer:
72     """Wrapper that logs access."""
73     def __init__(self, buffer, log):
74         self._buffer = buffer
75         self._log = log
76
77     def __getitem__(self, idx):
78         import time
79         timestamp = time.time_ns()
80         self._log.append(('read', idx, timestamp))
81         return self._buffer[idx]
82
83     def __setitem__(self, idx, value):
84         import time
85         timestamp = time.time_ns()
86         self._log.append(('write', idx, timestamp))
87         self._buffer[idx] = value
88
89     def __len__(self):
90         return len(self._buffer)
91
92 predictor._state.signal_circular_buffer = AccessTrackedBuffer(
93     original_buffer, access_log
94 )
95
96 np.random.seed(666)
97 data = np.random.randn(50) * 5 + 100
98
99 for obs in data:
100     predictor.step_with_telemetry(obs, previous_target=obs)
101
102 read_indices = [idx for op, idx, _ in access_log if op == 'read']
103
104 buffer_size = config.wtmm_buffer_size
105 causal_violations = 0
106
107 for i in range(1, len(read_indices)):
108     curr_idx = read_indices[i] % buffer_size
109     prev_idx = read_indices[i-1] % buffer_size
110
111     if curr_idx < prev_idx and (prev_idx - curr_idx) > buffer_size // 2:
112         causal_violations += 1
113
114 assert causal_violations == 0, \
115     f"Causal violations detected: {causal_violations} backward jumps"
116
117 def test_state_vector_does_not_leak_future():
118     """
119     Test: Sigma_t does not encode future information.
120     """
121     config = PredictorConfig()
122
123     predictor1 = UniversalPredictor(config)
124     data_short = np.random.randn(50) * 5 + 100

```

```

125
126     for obs in data_short:
127         result1 = predictor1.step_with_telemetry(obs, previous_target=obs)
128
129     state1_weights = np.copy(predictor1._state.weights)
130     state1_cusum = np.copy(predictor1._state.cusum_acum if hasattr(predictor1._state, 'cusum_acum') else [])
131
132     predictor2 = UniversalPredictor(config)
133     np.random.seed(np.random.RandomState(42).randint(2**32))
134     data_long = np.random.randn(100) * 5 + 100
135
136     for i in range(50):
137         result2 = predictor2.step_with_telemetry(data_long[i], previous_target=data_long[i])
138
139     state2_weights = np.copy(predictor2._state.weights)
140     state2_cusum = np.copy(predictor2._state.cusum_acum if hasattr(predictor2._state, 'cusum_acum') else [])
141
142     weights_diff = np.max(np.abs(state1_weights - state2_weights))
143
144     assert weights_diff < 0.05, \
145         f"State leaked future info: weights_diff={weights_diff:.3e}"

```

Chapter 11

Test Coverage Summary

11.1 Coverage Matrix

Module	Unit Tests	Integration Tests	Coverage
Levy generation	✓	-	95%
WTMM	✓	-	92%
Malliavin	✓	-	88%
Signatures	✓	-	90%
DGM entropy	✓	✓	93%
CUSUM	✓	✓	96%
CUSUM + Kurtosis	✓	✓	94%
Circuit breaker	-	✓	85%
Sinkhorn/JKO	-	✓	91%
DGM solver	-	✓	87%
Snapshotting	✓	-	97%
CPU/GPU parity	-	✓	82%
Walk-forward	-	✓	89%
Degraded mode	✓	✓	91%
Total			91%

Table 11.1: Test coverage by module

11.2 Full Suite Execution

11.2.1 Environment Validation in CI/CD

Before running the mathematical test suite (`pytest`), the CI pipeline must verify the virtual environment matches production via strict dependency validation. If versions diverge from the Golden Master, the pipeline must fail fast before running tensor tests.

Note: Since `requirements.txt` uses platform-specific environment markers (PEP 508), version extraction must handle semicolon separators and select the appropriate platform line.

```
1 #!/bin/bash
2 # Pre-pytest environment validation
3
4 # Extract versions from requirements.txt (handles environment markers)
5 # Format: "jax==0.4.38; sys_platform == 'darwin' and platform_machine == 'x86_64'"
6 EXPECTED_JAX=$(grep "^jax==" ../requirements.txt | head -1 | cut -d'=' -f3 | cut -d';' -f1)
7 EXPECTED_EQUINOX=$(grep "^equinox==" ../requirements.txt | cut -d'=' -f3)
8 EXPECTED_DIFFRAX=$(grep "^diffrax==" ../requirements.txt | cut -d'=' -f3)
9
10 ACTUAL_JAX=$(python -c "import jax; print(jax.__version__)")
```

```

11 ACTUAL_EQUINOX=$(python -c "import equinox; print(equinox.__version__)")
12 ACTUAL_DIFFRAX=$(python -c "import diffrax; print(diffrax.__version__)")
13
14 if [ "$EXPECTED_JAX" != "$ACTUAL_JAX" ]; then
15     echo "ERROR: JAX mismatch - Expected $EXPECTED_JAX, got $ACTUAL_JAX"
16     exit 1
17 fi
18
19 if [ "$EXPECTED_EQUINOX" != "$ACTUAL_EQUINOX" ]; then
20     echo "ERROR: Equinox mismatch - Expected $EXPECTED_EQUINOX, got $ACTUAL_EQUINOX"
21     exit 1
22 fi
23
24 if [ "$EXPECTED_DIFFRAX" != "$ACTUAL_DIFFRAX" ]; then
25     echo "ERROR: Diffrax mismatch - Expected $EXPECTED_DIFFRAX, got $ACTUAL_DIFFRAX"
26     exit 1
27 fi
28
29 echo " Environment validation OK - Proceed with pytest"

```

Listing 11.1: Pre-Test Environment Validation

11.2.2 Execution Commands

```

1 # Run all tests with coverage report
2 pytest tests/ -v --cov=Python --cov-report=html
3
4 # Run only fast tests (exclude GPU and optimization)
5 pytest tests/ -v -m "not slow"
6
7 # Run GPU parity tests (if available)
8 pytest tests/test_hardware/ -v -k gpu
9
10 # Parallel tests (4 workers)
11 pytest tests/ -n 4 --dist loadscope
12
13 # Generate XML report for CI/CD
14 pytest tests/ --junitxml=test-results.xml

```

11.3 Global Acceptance Criteria

1. **Code coverage:** $\geq 90\%$ in all critical modules
2. **Success rate:** 100% of tests must pass before merge
3. **Performance:** Full suite must run in < 5 minutes (no GPU, no Optuna)
4. **Reproducibility:** Fixed-seed tests must produce identical results
5. **Numerical parity:** CPU vs GPU relative error $< 10^{-5}$ in float32